

Computational Mathematics for Learning and Data Analysis

Computational Foundations, Numerical Linear Algebra, and
Optimization

Academic Year 2026/2027

Contents

Section 0: General Introduction to Computational Models	1
0 Mathematical Models in Machine Learning and Computational Foundations	2
0.1 The Two Pillars of Computational Mathematics in Machine Learning	2
0.2 The Philosophy of Mathematical Modeling: Brunelleschi and Galileo	3
0.2.1 The Genesis of Modeling: From Perspective to Computation	3
0.3 Characterization and Taxonomy of Models	4
0.3.1 The Triangle of Conflicting Requirements	4
0.3.2 Analytical Models vs Data-Driven (Synthetic) Models	5
0.3.3 Fitting vs Machine Learning: Empirical Risk and Generalization	6
0.4 Four Core Mathematical Models in Machine Learning	7
0.4.1 Example 1: Linear Least Squares Estimation	7
0.4.2 Example 2: Low-Rank Approximation and Recommender Systems	7
0.4.3 Example 3: Support Vector Machines (SVM) — Margin, Duality, and Kernels	9
0.4.4 Example 4: Clustering and the K-Means Algorithm	10
0.5 Methodological Principles of Computational Learning	12
0.5.1 The “Magic Academy” Metaphor	12
0.5.2 The Three Guiding Principles of Machine Learning	12
0.6 Synoptic Formula Reference of Introductory Models	13

I	Numerical Linear Algebra and Data Analysis	15
1	Linear Algebra Review: Vectors, Norms, Bases, and Transformations	16
1.1	Introduction to Numerical Computation and Applied Motivations	16
1.2	Vectors, Vector Spaces, and Geometric Representation	17
1.2.1	Notation and Conventions	17
1.2.2	Geometric Interpretation: Points, Arrows, and the Parallelogram Rule	18
1.2.3	Measuring the “Size” of a Vector	19
1.3	Vector Norms and Normed Spaces	19
1.3.1	The ℓ_p -Norm Family on $\mathbb{R}^n$	20
1.3.2	Rigorous Comparison of Vectors $\mathbf{v}$ and $\mathbf{w}$	20
1.4	Standard Inner Product	21
1.4.1	Matrix Notation and Output Dimensions	22
1.5	Bases of Vector Spaces and Linear Independence	22
1.5.1	Prominent Examples of Bases	23
1.5.2	Linear Independence and Cardinality Constraints	23
1.6	Classroom Exercises: Unit Spheres and Norm Geometry . . .	24
1.6.1	Exercise 1: Comparative Norm Calculation in $\mathbb{R}^3$. . .	24
1.6.2	Exercise 2: Plotting Unit Spheres in $\mathbb{R}^2$	25
1.6.3	Exercise 3: Basis Verification in $\mathbb{R}^3$	26
1.7	Three Perspectives on Matrix-Vector Products	27
1.7.1	Definition and Dimensional Compatibility	27
1.7.2	The Three Equivalent Perspectives	27
1.8	Matrices as Linear Transformations in the Plane	29
1.8.1	Core Geometric Examples	29
1.9	Introduction to Computational Complexity and Flops	30
1.9.1	Operation Count for the Inner Product	30
1.9.2	Linear Time Scaling Model	30
1.10	Summary of Core Mathematical Concepts	31
2	Matrix Multiplication, Computational Complexity, and Orthogonal Matrices	32
2.1	Matrix-Matrix Multiplication	32
2.1.1	Dimensional Compatibility Condition	32
2.1.2	Component-Wise Definition	33
2.1.3	Step-by-Step Numerical Example	33
2.2	Computational Complexity and Scaling Law	34
2.2.1	Rigorous FLOP Counting	34
2.2.2	Cubic Time Scaling Law	35

2.2.3	Comparative Overview of Core Linear Algebra Operations	35
2.3	Fundamental Properties: Rank, Subspaces, and Invertibility	35
2.3.1	Subspaces Associated with a Linear Map	35
2.3.2	Invertibility of Square Matrices	37
2.3.3	Guided Example: Rank, Kernel, and Image Calculation	37
2.4	Vector Orthogonality and Orthonormality	38
2.4.1	Core Definitions	38
2.4.2	Constructive Normalization Exercises in $\mathbb{R}^2$	39
2.5	Orthogonal Matrices	39
2.5.1	Definition and Transpose Properties	39
2.5.2	Column and Row Characterization	40
2.5.3	Prototypical Example: Planar Rotation Matrix	41
2.5.4	Closure under Matrix Products	41
2.6	The Four Noteworthy Questions on Orthogonal Matrices	41
2.6.1	Question 1: Upper Bound on Individual Matrix Entries	41
2.6.2	Question 2: Orthogonal Diagonal Matrices	42
2.6.3	Question 3: Orthogonal Upper Triangular Matrices	42
2.6.4	Question 4: Euclidean Isometry and Norm Preservation	43
2.7	Chapter Summary Reference	45
3	Orthonormal Matrices, Eigenvalues, and Quadratic Forms	46
3.1	Review of Square Orthogonal Matrices and Isometry	46
3.2	Rectangular Matrices with Orthonormal Columns	47
3.2.1	Definition and Dimensional Constraint	48
3.2.2	The Inner Product $V^T V$ (Column Gramian)	48
3.2.3	The Reversed Product $V V^T$ (Orthogonal Projector)	49
3.3	Partial Isometry for Orthonormal Matrices	50
3.3.1	Question (a): Norm Preservation for Vectors in $\mathbb{R}^k$	50
3.3.2	Question (b): Norm Preservation for Vectors in $\mathbb{R}^n$ under V^T	51
3.4	Eigenvalues and Eigenvectors: Foundations and Step-by-Step Example	51
3.4.1	Fundamental Definitions	51
3.4.2	Step-by-Step Worked Example	52
3.5	Diagonalization and Dyadic Expansion	53
3.5.1	Diagonalizable Matrices	53
3.5.2	Dyadic (Outer Product) Expansion	53
3.5.3	Pathologies: When Diagonalization Fails	54
3.6	Guided Exercises on Eigenvalues (Q1 – Q4)	55
3.6.1	Exercise Q1: Eigenpairs of a Diagonal Matrix	55
3.6.2	Exercise Q2: Eigenpairs of the Identity Matrix	56

3.6.3	Exercise Q3: Sum and Product of Matrices Sharing an Eigenvalue	56
3.6.4	Exercise Q4: Eigenpairs of a Matrix Power	57
3.7	Symmetric Matrices and the Spectral Theorem	58
3.8	Quadratic Forms and the Rayleigh Quotient Inequality	59
3.8.1	Definition of a Quadratic Form	59
3.8.2	Spectral Bounds Theorem (Rayleigh Inequality)	60
3.8.3	Equality Conditions and the Rayleigh Quotient	61
3.9	Symmetric Positive Definite and Semidefinite Matrices	61
3.9.1	Axiomatic Definitions	61
3.9.2	Equivalent Spectral Characterization	62
3.9.3	Matrix Class Hierarchy and Nested Inclusions	62
3.9.4	Illustrative Examples in $\mathbb{R}^{2 \times 2}$	63
3.10	Natural Generation of SPSD Matrices	63
3.11	Chapter Summary Formulary	66

II Continuous Optimization and Machine Learning 67

4	Elementary Univariate Optimization Problems	68
4.1	Introduction, Motivation, and Learning Objectives	68
4.2	Fundamental Definitions: Spaces, Functions, Graphs, and Level Sets	69
4.2.1	Input and Output Spaces	69
4.2.2	Graph of a Function	69
4.2.3	Image (Effective Codomain)	70
4.2.4	Level Sets and Sublevel Sets	70
4.3	The Unconstrained Univariate Optimization Problem	70
4.3.1	Formal Problem Statement	70
4.3.2	Optimal Value vs Optimal Solution	71
4.3.3	Indifference Between Equivalent Minimizers	71
4.4	Simple Reformulations and Equivalence	72
4.4.1	Min–Max Duality	72
4.4.2	Invariance Under Vertical Shifts (Additive Constant)	72
4.4.3	Invariance Under Positive Scaling (Multiplicative Constant)	72
4.5	Constrained Optimization and the Feasible Region	73
4.5.1	General Formulation of Constrained Optimization	73
4.5.2	Monotone Non-Improvement Under Constraint Restriction	73
4.5.3	Taxonomy of Redundant Constraints	74

4.5.4	Functional Representation of Feasible Sets	74
4.6	Why Optimization is Difficult: Theory vs Computational Practice	75
4.6.1	Modes of Non-Existence of the Minimum	75
4.6.2	Formalization via Extended Real Numbers	76
4.6.3	Optimization in Floating-Point Arithmetic	76
4.6.4	Error Metrics: Absolute Gap vs Relative Gap	77
4.6.5	The Impossibility of Generic Black-Box Optimization .	77
4.7	Univariate Linear Optimization	78
4.7.1	Geometric and Analytical Properties	78
4.7.2	Unconstrained Optimization on $\mathbb{R}$	78
4.7.3	Optimization Over a Bounded Box	79
4.8	The Illusion of Open Sets in Numerical Computation	79
4.8.1	The Analytical Failure of Open Intervals	79
4.8.2	The Practical Numerical Resolution	79
4.9	Univariate Optimization: The Homogeneous Quadratic Case .	80
4.9.1	Definition and Curvature	80
4.9.2	Unconstrained Optimization on $\mathbb{R}$	80
4.9.3	Box-Constrained Optimization	81
4.10	Univariate Optimization: Non-Homogeneous Quadratic Case .	82
4.10.1	Definition and Non-Separability of the Minimum Operator	82
4.10.2	The Coordinate Translation Principle (Origin Shift) . .	82
4.10.3	Step-by-Step Algebraic Derivation	82
4.10.4	Solution in Shifted Coordinates and Inverse Mapping .	83
4.10.5	Geometric Interpretation: The Midpoint of the Roots .	84
4.10.6	Box-Constrained General Quadratic Optimization . . .	85
4.11	Synoptic Formula Reference of Univariate Optimization	86
5	Multivariate Optimization: Geometry, Linear Functions, and Quadratic Forms	87
5.1	From Univariate to Multivariate: Dimensionality of $\mathbb{R}^n$	87
5.1.1	Multivariate Input Space	87
5.1.2	The Curse of Dimensionality	88
5.2	Vector and Multi-Objective Optimization	89
5.2.1	The Limitation of a Single Scalar Objective	89
5.2.2	Illustrative Example: Markowitz Portfolio Selection . .	89
5.2.3	Partial Ordering and the Pareto Frontier	90
5.2.4	Operational Reduction Techniques: Scalarization and Budgeting	91
5.3	Euclidean Geometry in $\mathbb{R}^n$ and Function Tomography	91
5.3.1	Standard Inner Product and Induced Norm	91

5.3.2	Angular Representation and the Cauchy-Schwarz Inequality	92
5.3.3	The Concept of Function Tomography	92
5.4	Multivariate Linear Functions	93
5.4.1	Definition and Separability	93
5.4.2	Geometry of Graph and Level Sets	93
5.5	Linear Tomography and Steepest Descent Direction	94
5.5.1	Tomographic Cross-Section Computation	94
5.5.2	Directional Analysis: The Radar Screen Analogy	94
5.5.3	Formal Derivation of the Steepest Descent Direction	95
5.6	From Linear to Quadratic Functions: The Separable Case	96
5.6.1	Separable Quadratic Functions	96
5.6.2	Anisotropic Curvature and Eccentricity	96
5.7	General Homogeneous Quadratic Forms and Symmetrization	97
5.7.1	Definition and Cross-Terms	97
5.7.2	Symmetrization Theorem	98
5.8	Spectral Decomposition and Matrix Taxonomy	98
5.8.1	Central Symmetry and Tomography	98
5.8.2	The Spectral Theorem for Real Symmetric Matrices	99
5.8.3	Variational Characterization: The Rayleigh Quotient	99
5.8.4	Taxonomy of Matrix Definiteness	100
5.9	Spectral Tomography Across Four Regimes	100
5.9.1	Case I: Positive Definite Matrix ($Q \succ 0$)	101
5.9.2	Case II: Positive Semidefinite Matrix ($Q \succeq 0$)	101
5.9.3	Case III: Indefinite Matrix ($Q \succ \prec 0$)	102
5.9.4	Case IV: Negative Definite Matrix ($Q \prec 0$)	102
5.10	Geometry of Level Sets	102
5.10.1	Diagonalization and Orthonormal Coordinate Rotation	102
5.10.2	Semi-Axis Theorem for Ellipsoids	103
5.10.3	Degeneration into Open Elliptic Cylinders ($\lambda_i \rightarrow 0$)	104
5.10.4	Summary of Level Geometries	105
5.11	Quadratic Optimization: Unconstrained vs Box-Constrained	105
5.11.1	Unconstrained Optimization: Immediate Analytical Solution	105
5.11.2	Box-Constrained Optimization: Complexity Shift	105

Section 0: General Introduction to Computational Models

Chapter 0

Mathematical Models in Machine Learning and Computational Foundations

0.1 The Two Pillars of Computational Mathematics in Machine Learning

The computational treatment of machine learning and large-scale data analysis rests upon two complementary, deeply intertwined mathematical disciplines [1]:

- **Numerical Linear Algebra:** provides the foundational language to represent geometric transformations, project data onto subspaces, compute matrix factorizations (SVD, QR, Cholesky), and quantify floating-point error propagation via condition numbers [2, 3].
- **Mathematical Optimization:** formalizes the decision-making process through constrained and unconstrained formulations, studies analytical optimality conditions (Fermat’s stationary rules and Karush-Kuhn-Tucker conditions), investigates Lagrangian duality, and designs scalable iterative descent algorithms [4, 5, 6].

In real-world Data Science, optimization algorithms cannot function without efficient, stable linear algebra routines, as determining descent directions at each step invariably requires solving a linear algebraic system.




Figure 1: Synergistic conceptual framework: complementary integration of Numerical Linear Algebra and Mathematical Optimization in Machine Learning.

0.2 The Philosophy of Mathematical Modeling: Brunelleschi and Galileo

0.2.1 The Genesis of Modeling: From Perspective to Computation

Extracting actionable insight from massive datasets addresses a fundamental epistemological challenge: raw data is unstructured, cognitively overwhelming, and incapable of making autonomous decisions.

Definition 0.2.1 (Computational Model). **Learning (machine learning)** consists of extracting structural regularities from empirical observations and synthesizing them into an abstract mathematical artifact (**model**), computa-

tionally tractable and designed to drive predictions or operational decisions.

This philosophy has deep historical roots:

- **Filippo Brunelleschi and Perspective:** To erect the free-standing dome of Santa Maria del Fiore in Florence without ground-based wooden falsework, Brunelleschi did not proceed by trial and error on site; he built rigorous mathematical models (geometric projections on an optical cone) capable of predicting radial thrust forces.
- **The Economic Advantage of Mathematics:** Constructing physical scale models (such as wind tunnels in aeronautics) incurs enormous material costs. **Mathematics costs far less than wood or steel:** equations and algorithms allow simulating billions of operational scenarios at negligible cost.
- **Galileo Galilei and Scientific Abstraction:** Modern physics was born when scientists ceased cataloging every accidental empirical detail and instead abstracted parsimonious mathematical laws expressed in algebraic language.

Definition of a Computational Mathematical Model

A model is not an identical replica of physical reality. It is an **abstract mathematical formalization** designed to capture essential causal or correlative patterns while filtering out noise, making the problem tractable on digital computers.

0.3 Characterization and Taxonomy of Models

0.3.1 The Triangle of Conflicting Requirements

Designing any computational model involves balancing three inherently competing objectives:

1. **Accuracy (Accurate):** The fidelity with which the model reflects the actual physical dynamics or behavior of the system.
2. **Computational Inexpensiveness (Inexpensive):** Algorithmic simplicity, translating to minimal memory footprints and fast execution times ($\mathcal{O}(n)$ or quasi-linear complexity).

3. **Generality (General):** The versatility to apply across diverse operational environments and domain distributions.

Property 0.3.1 (Model Incompatibility Principle). *It is theoretically and practically impossible to maximize all three objectives simultaneously: hyper-accurate models require vast parameter counts that destroy computational efficiency; parsimonious, rapid models compromise fine-grained accuracy or generality. Model design is inevitably an **engineering trade-off**.*




Figure 2: The triangle of conflicting requirements in mathematical modeling: accuracy, computational inexpensiveness, and generality.

0.3.2 Analytical Models vs Data-Driven (Synthetic) Models

Table 1 outlines the core differences between the two primary modeling paradigms.

Property	Analytical Approach (White-Box)	Data-Driven Approach (Black-Box)
Principio Base	Models each individual physical component and its differential interactions.	Does not model internal mechanics; reproduces observed input-output behavior directly.
Struttura Matematica	Flexible, non-linear, derived from first physical principles and conservation laws.	Standardized, rigid, parametric (e.g., deep neural networks, SVMs, linear hyperplanes).
Parameters	Few, calibrated parameters with precise, interpretable physical meaning.	Massive parameter counts (millions to billions), estimated automatically via optimization.
Development Cost	High: requires deep domain expertise and invasive internal system access.	Low/scalable: requires only historical observation logs and standardized optimization solvers.

Table 1: Comparison between analytical white-box modeling and empirical data-driven black-box modeling.

0.3.3 Fitting vs Machine Learning: Empirical Risk and Generalization

Every model is an approximation (*“the map is not the territory”*):

- **Fitting (Empirical Risk Minimization):** Identifies parameter vectors that minimize discrepancy purely on known observations (*training set*):

$$\min_{\mathbf{w}} \mathcal{R}_{\text{emp}}(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \ell(f(\mathbf{x}^i; \mathbf{w}), y^i). \quad (1)$$

- **Learning (Generalization):** Requires $f(\mathbf{x}; \mathbf{w})$ to maintain low error on unseen data drawn from the same underlying distribution (*test set*). A hyper-parametrized model that drives training error to zero often suffers from **overfitting**, fitting spurious noise rather than true underlying relationships.

0.4 Four Core Mathematical Models in Machine Learning

0.4.1 Example 1: Linear Least Squares Estimation

Given m observation pairs $(\mathbf{x}^i, y^i) \in \mathbb{R}^n \times \mathbb{R}$, we postulate a linear input-output mapping:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \sum_{j=1}^n w_j x_j + w_0. \quad (2)$$

Absorbing the bias w_0 into the augmented design matrix X :

$$\mathbf{w}_{\text{ext}} = \begin{bmatrix} w_0 \\ \mathbf{w} \end{bmatrix} \in \mathbb{R}^{n+1}, \quad X = \begin{bmatrix} 1 & (\mathbf{x}^1)^T \\ \vdots & \vdots \\ 1 & (\mathbf{x}^m)^T \end{bmatrix} \in \mathbb{R}^{m \times (n+1)}, \quad \mathbf{y} = \begin{bmatrix} y^1 \\ \vdots \\ y^m \end{bmatrix} \in \mathbb{R}^m. \quad (3)$$

The linear least squares problem minimizes the Euclidean residual:

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \|\mathbf{y} - X\mathbf{w}\|_2^2 = \mathbf{y}^T \mathbf{y} - 2\mathbf{w}^T X^T \mathbf{y} + \mathbf{w}^T X^T X \mathbf{w}. \quad (4)$$

Setting the gradient $\nabla \mathcal{L}(\mathbf{w})$ to zero yields the stationary condition, known as the **Gauss Normal Equations**:

$$\nabla \mathcal{L}(\mathbf{w}) = -2X^T \mathbf{y} + 2X^T X \mathbf{w} = \mathbf{0} \implies X^T X \mathbf{w} = X^T \mathbf{y}. \quad (5)$$

When X has full column rank ($\text{rank}(X) = n+1$), $X^T X$ is symmetric positive definite and invertible:

$$\mathbf{w}^* = (X^T X)^{-1} X^T \mathbf{y}. \quad (6)$$

Golden Numerical Rule: Never Compute $(X^T X)^{-1}$ Explicitly

In scientific computing, **one must never explicitly form $X^T X$ nor compute its inverse**. Multiplying $X^T X$ squares the condition number ($\kappa(X^T X) = \kappa(X)^2$), drastically magnifying floating-point roundoff errors. The system is solved stably using QR factorization or truncated SVD (backslash operator in GNU Octave: $\mathbf{w} = \mathbf{X} \backslash \mathbf{y}$).

0.4.2 Example 2: Low-Rank Approximation and Recommender Systems

Given a partially observed data matrix $M \in \mathbb{R}^{n \times m}$ (e.g., n users and m movies, with M_{ij} indicating ratings), we seek a low-rank factorization $k \ll \min(n, m)$

such that $M \approx AB$:

$$A \in \mathbb{R}^{n \times k}, \quad B \in \mathbb{R}^{k \times m} \implies \min_{A, B} \|M - AB\|_F^2 = \sum_{i=1}^n \sum_{j=1}^m (M_{ij} - (AB)_{ij})^2. \quad (7)$$




Figure 3: Low-rank matrix factorization: compressing storage from nm entries to $k(n + m)$.

Memory Compression and Spectral Properties

Storing M explicitly requires $n \cdot m$ memory entries; storing A and B requires $k(n + m)$. For $n = 10^6$, $m = 10^5$, and rank $k = 10$:

$$\text{Raw storage} = 10^6 \times 10^5 = 10^{11} \text{ cells}, \quad (8)$$

$$\text{Factorized storage} = 10 \times (10^6 + 10^5) \approx 1.1 \times 10^7 \text{ cells}, \quad (9)$$

achieving a **99.99% storage reduction**.

By the Eckart-Young-Mirsky Theorem, the optimal low-rank matrix under the Frobenius norm is given by the truncated Singular Value Decomposition (SVD).

Computational Warning on Sparse Matrices

When M is sparse, **one must never form $M^T M$ or MM^T** . Matrix multiplication destroys sparsity, producing a dense matrix that triggers immediate Out-Of-Memory (OOM) failures. One must rely exclusively on iterative methods using matrix-vector products (Lanczos bidiagonalization).

0.4.3 Example 3: Support Vector Machines (SVM) — Margin, Duality, and Kernels

For binary classification with labels $y^h \in \{-1, +1\}$, a linear classifier makes predictions via $\hat{y} = \text{sign}(\mathbf{w}_+^T \mathbf{x} - w_0)$.




Figure 4: Margin geometry in linear Support Vector Machines: maximizing the separation distance between supporting hyperplanes.

Maximizing the geometric margin $2/\|\mathbf{w}_+\|_2$ is equivalent to minimizing $\|\mathbf{w}_+\|_2^2$:

$$\min_{\mathbf{w}_+, w_0} \frac{1}{2} \|\mathbf{w}_+\|_2^2 \quad \text{s.t.} \quad y^h (\mathbf{w}_+^T \mathbf{x}^h - w_0) \geq 1 \quad \forall h = 1, \dots, m. \quad (10)$$

For non-linearly separable data, slack variables $\xi_h \geq 0$ yield the **Soft Margin** problem:

$$\begin{aligned} \min_{\mathbf{w}_+, w_0, \xi} \quad & \frac{1}{2} \|\mathbf{w}_+\|_2^2 + C \sum_{h=1}^m \xi_h \\ \text{s.t.} \quad & y^h (\mathbf{w}_+^T \mathbf{x}^h - w_0) \geq 1 - \xi_h, \quad \xi_h \geq 0 \quad \forall h = 1, \dots, m. \end{aligned} \quad (11)$$

Parameter $C > 0$ controls the trade-off between margin width and classification violations: $C \rightarrow \infty$ forces hard separation, while moderate values allow controlled margin errors to protect test generalization.

Non-Linear Lifting, Duality, and the Kernel Trick

When points are not linearly separable in $\mathbb{R}^n$, a non-linear mapping $\phi : \mathbb{R}^n \rightarrow \mathcal{H}$ lifts data into a higher-dimensional feature space:

$$\mathbf{x} \mapsto \phi(\mathbf{x}). \quad (12)$$

The Wolfe dual problem is given by:

$$\begin{aligned} \max_{\alpha} \quad & \sum_{h=1}^m \alpha_h - \frac{1}{2} \sum_{h=1}^m \sum_{k=1}^m \alpha_h \alpha_k y^h y^k \langle \phi(\mathbf{x}^h), \phi(\mathbf{x}^k) \rangle_{\mathcal{H}} \\ \text{s.t.} \quad & \sum_{h=1}^m y^h \alpha_h = 0, \quad 0 \leq \alpha_h \leq C \quad \forall h = 1, \dots, m. \end{aligned} \quad (13)$$

Because data points appear solely through inner products, Mercer's Theorem permits evaluating them via a **kernel function** $K(\mathbf{x}^h, \mathbf{x}^k) = \langle \phi(\mathbf{x}^h), \phi(\mathbf{x}^k) \rangle_{\mathcal{H}}$. For example, the **Gaussian RBF Kernel**:

$$K(\mathbf{x}, \mathbf{z}) = \exp\left(-\gamma \|\mathbf{x} - \mathbf{z}\|_2^2\right) \quad (14)$$

corresponds to an infinite-dimensional Hilbert space $\mathcal{H}$, evaluated in $\mathcal{O}(n)$ time without ever forming $\phi(\mathbf{x})$ explicitly.

0.4.4 Example 4: Clustering and the K-Means Algorithm

Given an unlabelled dataset $X = \{\mathbf{x}^1, \dots, \mathbf{x}^m\} \subset \mathbb{R}^h$, clustering seeks to partition X into k disjoint subsets $\{X_p\}_{p=1}^k$ maximizing intra-cluster cohesion.

Biconvex Optimization Formulation

Introducing k cluster centroids $\mathbf{c}^1, \dots, \mathbf{c}^k \in \mathbb{R}^h$ and binary assignment variables $z_{ip} \in \{0, 1\}$:

$$\min_{\mathbf{c}, Z} \sum_{i=1}^m \sum_{p=1}^k z_{ip} \|\mathbf{c}^p - \mathbf{x}^i\|_2^2 \quad \text{s.t.} \quad \sum_{p=1}^k z_{ip} = 1 \quad \forall i, \quad z_{ip} \in \{0, 1\} \quad \forall i, p. \quad (15)$$

The problem is non-convex due to binary integrality $z_{ip} \in \{0, 1\}$ and bilinear terms $z_{ip} \mathbf{c}^p$. However, its biconvex structure enables **Block Gauss-Seidel alternating optimization**:

1. **Fixing $\mathbf{c}$, optimize Z :** Each point is assigned to its nearest Euclidean centroid ($z_{i\bar{p}} = 1 \iff \bar{p} = \arg \min_p \|\mathbf{c}^p - \mathbf{x}^i\|_2^2$).

2. **Fixing Z , optimize $\mathbf{c}$:** Vanishing the convex quadratic gradient with respect to $\mathbf{c}^p$:

$$2 \sum_{i \in I(p)} (\mathbf{c}^p - \mathbf{x}^i) = \mathbf{0} \implies (\mathbf{c}^p)^* = \frac{1}{|I(p)|} \sum_{i \in I(p)} \mathbf{x}^i. \quad (16)$$

The optimal centroid is the **vector mean (barycenter)** of all points belonging to cluster p .

Algorithm 1: K-Means Algorithm for Euclidean Clustering

Input: Dataset $X = \{\mathbf{x}^1, \dots, \mathbf{x}^m\} \subset \mathbb{R}^h$, number of clusters k , initial centroids $\mathbf{c}^1, \dots, \mathbf{c}^k$, tolerance $\epsilon > 0$

Output: Final centroids $\mathbf{c}^1, \dots, \mathbf{c}^k$ and partition $\{I(p)\}_{p=1}^k$

```

1  $v \leftarrow \infty$ 
2 repeat
    // Step 1: Point assignment to nearest centroid
3   for  $p \leftarrow 1$  to  $k$  do
4      $I(p) \leftarrow \emptyset$ 
5   for  $i \leftarrow 1$  to  $m$  do
6      $p^* \leftarrow \arg \min_{p \in \{1, \dots, k\}} \|\mathbf{c}^p - \mathbf{x}^i\|_2^2$ 
7      $I(p^*) \leftarrow I(p^*) \cup \{i\}$ 
    // Step 2: Centroid re-computation (mean)
8   for  $p \leftarrow 1$  to  $k$  do
9     if  $I(p) \neq \emptyset$  then
10       $\mathbf{c}^p \leftarrow \frac{1}{|I(p)|} \sum_{i \in I(p)} \mathbf{x}^i$ 
    // Step 3: Objective value evaluation
11    $v_{\text{new}} \leftarrow \sum_{p=1}^k \sum_{i \in I(p)} \|\mathbf{c}^p - \mathbf{x}^i\|_2^2$ 
12    $v \leftarrow v_{\text{new}}$ 
13 until  $(v - v_{\text{new}}) \leq \epsilon$ 
14 return  $\mathbf{c}^1, \dots, \mathbf{c}^k$ 

```

Why Scientific Computing Violates Agile Modularity

In standard software engineering (Agile), systems are decomposed into decoupled, independently swappable modules. In scientific computing, this separation fails:

- Replacing the Euclidean metric ℓ_2 with the Manhattan distance ℓ_1 completely changes the closed-form centroid: the ℓ_1 optimum is no longer the arithmetic mean, but the **component-wise median**.
- The tight interdependence between loss formulation and update routines demands an integrated, holistic mathematical design.

0.5 Methodological Principles of Computational Learning

0.5.1 The “Magic Academy” Metaphor

To illustrate the difference between superficial library usage and rigorous computational competence, consider the **Magic Academy**:

- **The Sorcerer’s Apprentice (*Black-Box Practitioner*)**: Invokes library routines (`import sklearn, fit()`) as pre-packaged incantations. In nominal conditions, output appears satisfactory; however, when confronted with numerical instability or ill-conditioning, they lack the insight required to diagnose the root cause.
- **The Master Magician (*Computational Scientist*)**: Understands spectral properties and convergence theory, detects ill-posed systems, tracks floating-point error amplification, and replaces fragile algorithms with robust, numerically sound alternatives.

0.5.2 The Three Guiding Principles of Machine Learning

1. **Learning is Numerical Computation**: Every AI algorithm fundamentally reduces to solving a linear algebraic system or minimizing an objective loss function.
2. **The Dimensional Scale Challenge**: Machine learning problems are rarely mathematically intractable; they are predominantly linear,

convex, or quadratic. **Their difficulty stems entirely from their colossal scale** (*«hard because large, not hard because hard»*), spanning millions to billions of parameters.

3. **Linear Algebra as the Engine of Optimization:** Optimizing at extreme scales demands ruthless hardware efficiency: numerical linear algebra provides the computational engine that makes optimization fast, scalable, and numerically stable.

0.6 Synoptic Formula Reference of Introductory Models

Table 2 summarizes the foundational mathematical formulations and their computational paradigms.

Model	Formulation	Paradigm	Computational Notes
Least Squares	$\min_{\mathbf{w}} \ \mathbf{y} - X\mathbf{w}\ _2^2$	Normal equations: $\mathbf{w}^* = (X^T X)^{-1} X^T \mathbf{y}$	Solve via QR factorization; never invert $X^T X$ explicitly.
Low-Rank	$\min_{A,B} \ M - AB\ _F^2$	Truncated SVD (Eckart-Young)	Compression $nm \rightarrow k(n+m)$; avoid Gramian $M^T M$ on sparse data.
SVM (Hard)	$\min \frac{1}{2} \ \mathbf{w}_+\ _2^2$ s.t. $y^h(\mathbf{w}_+^T \mathbf{x}^h - w_0) \geq 1$	Convex Quadratic Program (QP)	Maximizes geometric margin $2/\ \mathbf{w}_+\ _2$ under strict separation.
SVM (Soft)	$\min \frac{1}{2} \ \mathbf{w}_+\ _2^2 + C \sum \xi_h$ s.t. $y^h(\mathbf{w}_+^T \mathbf{x}^h - w_0) \geq 1 - \xi_h$	QP with linear constraints (equivalent to Hinge loss)	Trade-off parameter $C > 0$ balances bias-variance via slack variables.
SVM (Kernel)	$\max_{\boldsymbol{\alpha}} \mathbf{1}^T \boldsymbol{\alpha} - \frac{1}{2} \boldsymbol{\alpha}^T \tilde{K} \boldsymbol{\alpha}$ s.t. $\mathbf{y}^T \boldsymbol{\alpha} = 0, 0 \leq \alpha_h \leq C$	QP with box constraints	Inner products in infinite-dimensional Hilbert spaces (RBF) in $\mathcal{O}(n)$.
K-Means	$\min \sum_{i,p} z_{ip} \ \mathbf{c}^p - \mathbf{x}^i\ _2^2$ s.t. $\sum_p z_{ip} = 1, z_{ip} \in \{0, 1\}$	Alternating minimization (Block Gauss-Seidel)	Finite convergence to local minimum; cluster centroids given in closed form.

Table 2: Synoptic reference of introductory mathematical models and computational properties.

Part I

**Numerical Linear Algebra and
Data Analysis**

Chapter 1

Linear Algebra Review: Vectors, Norms, Bases, and Transformations

1.1 Introduction to Numerical Computation and Applied Motivations

This course addresses the computational and efficient resolution of linear algebra problems on digital computers [2, 7]. In applied mathematics, data science, and machine learning, the general formulation of a decision problem typically takes the form of an optimization task:

$$\min_{\mathbf{x} \in S} f(\mathbf{x}), \quad (1.1)$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ denotes the objective function and $S \subseteq \mathbb{R}^n$ defines the feasible set of constraints.

Within numerical linear algebra, this general problem is frequently reformulated or locally approximated by a linear least squares problem:

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2, \quad (1.2)$$

where $A \in \mathbb{R}^{m \times n}$ represents the model or design matrix, $\mathbf{b} \in \mathbb{R}^m$ the observed target vector, and $\mathbf{x} \in \mathbb{R}^n$ the unknown parameter vector to estimate.

Why Focus on Linear Problems?

Key motivations establishing linear problems as fundamental to scientific computing include:

1. **Theoretical tractability and analytical clarity:** Linear problems exhibit clean, well-understood algebraic structures, enabling rigorous guarantees on existence, uniqueness, and algorithm convergence.
2. **Efficiency and machine accuracy:** Specialized direct and iterative algorithms can compute solutions up to the limits of machine floating-point arithmetic (IEEE 754 standard, $\epsilon_{\text{mach}} \approx 10^{-16}$ in double precision).
3. **Exploitation of special matrix structures:** Matrices endowed with specific properties (symmetric, positive definite, orthogonal, banded, Toeplitz, or sparse) permit specialized algorithms with significantly reduced asymptotic and memory complexity.
4. **Building blocks for non-linear algorithms:** Nearly all sophisticated numerical methods for non-linear problems (Newton's method, Quasi-Newton schemes like BFGS, constrained optimization, and PDE solvers) solve one or more linear systems or least squares subproblems at every iteration.
5. **Stability and error control:** Numerical practitioners demand algorithms resilient against catastrophic round-off error accumulation and algorithmic instability.

1.2 Vectors, Vector Spaces, and Geometric Representation

1.2.1 Notation and Conventions

Unless stated otherwise, **all vectors are assumed to be column vectors**. A vector $\mathbf{x} \in \mathbb{R}^n$ is denoted:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, \quad \mathbf{x}^T = \begin{bmatrix} x_1 & x_2 & \dots & x_n \end{bmatrix}. \quad (1.3)$$

Consider, for example in $\mathbb{R}^2$, the two vectors:

$$\mathbf{x} = \begin{bmatrix} 2 \\ 4 \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}. \quad (1.4)$$

The fundamental operation defining vector space addition is component-wise summation:

$$\mathbf{x} + \mathbf{y} = \begin{bmatrix} 2 \\ 4 \end{bmatrix} + \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 2+2 \\ 4+1 \end{bmatrix} = \begin{bmatrix} 4 \\ 5 \end{bmatrix}. \quad (1.5)$$

1.2.2 Geometric Interpretation: Points, Arrows, and the Parallelogram Rule

A vector $\mathbf{x} = [x_1, x_2]^T \in \mathbb{R}^2$ admits two complementary geometric perspectives:

- as a **point** in the Cartesian plane with coordinates (x_1, x_2) relative to the origin;
- as a **directed segment (arrow)** originating at $(0,0)$ and terminating at (x_1, x_2) .

Vector addition $\mathbf{x} + \mathbf{y}$ corresponds to the classical **parallelogram law**: translating the tail of $\mathbf{y}$ to the tip of $\mathbf{x}$, the resulting vector connects the origin to the opposite vertex of the parallelogram formed by the two vectors, terminating at $(4,5)$. This notion generalizes directly to n -dimensional Euclidean spaces $\mathbb{R}^n$ ($n \geq 3$).




Figure 1.1: Geometric interpretation of vector addition in $\mathbb{R}^2$ via the parallelogram rule.

1.2.3 Measuring the “Size” of a Vector

Consider the following question:

Given the vectors $\mathbf{v}, \mathbf{w} \in \mathbb{R}^4$:

$$\mathbf{v} = \begin{bmatrix} 2 \\ 2 \\ 2 \\ 2 \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} 4 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad (1.6)$$

which of these two vectors is “larger”?

The question is fundamentally ill-posed until a precise formal criterion of measurement is chosen:

- Looking at the **maximum individual component**, $\mathbf{w}$ contains an entry of 4, whereas all entries of $\mathbf{v}$ are 2.
- Looking at the **sum of components**, $\mathbf{v}$ sums to $2 + 2 + 2 + 2 = 8$, while $\mathbf{w}$ only reaches 4.
- Looking at the **Euclidean geometric length**, both measure $\sqrt{16} = 4$.

To answer such questions unequivocally, the formal concept of a **vector norm** is required.

1.3 Vector Norms and Normed Spaces

Definition 1.3.1 (Vector Norm). Let V be a real (or complex) vector space. A function $\|\cdot\| : V \rightarrow \mathbb{R}$ is called a **vector norm** if and only if it satisfies the following three axioms for all $\mathbf{v}, \mathbf{w} \in V$ and every scalar $\lambda \in \mathbb{R}$:

1. **Non-negativity and positive definiteness:**

$$\|\mathbf{v}\| \geq 0 \quad \forall \mathbf{v} \in V, \quad \|\mathbf{v}\| = 0 \iff \mathbf{v} = \mathbf{0}_V. \quad (1.7)$$

2. **Absolute homogeneity:**

$$\|\lambda \mathbf{v}\| = |\lambda| \cdot \|\mathbf{v}\| \quad \forall \lambda \in \mathbb{R}, \forall \mathbf{v} \in V. \quad (1.8)$$

3. **Triangle inequality:**

$$\|\mathbf{v} + \mathbf{w}\| \leq \|\mathbf{v}\| + \|\mathbf{w}\| \quad \forall \mathbf{v}, \mathbf{w} \in V. \quad (1.9)$$

A vector space V paired with a norm $\|\cdot\|$ is termed a **normed space**.

1.3.1 The ℓ_p -Norm Family on $\mathbb{R}^n$

Let $\mathbf{v} = [v_1, v_2, \dots, v_n]^T \in \mathbb{R}^n$. The norms most commonly employed in numerical analysis and machine learning belong to the parameterized ℓ_p -norm family:

$$\|\mathbf{v}\|_p = \left(\sum_{i=1}^n |v_i|^p \right)^{1/p} \quad (p \geq 1). \quad (1.10)$$

Three canonical special cases are:

1. Euclidean Norm (ℓ_2). Represents standard geometric distance in Euclidean affine space:

$$\|\mathbf{v}\|_2 = \sqrt{\sum_{i=1}^n v_i^2} = \sqrt{\mathbf{v}^T \mathbf{v}}. \quad (1.11)$$

2. Manhattan Norm (ℓ_1 or Taxi-cab). Sums the absolute values of all coordinates:

$$\|\mathbf{v}\|_1 = \sum_{i=1}^n |v_i|. \quad (1.12)$$

3. Maximum Norm (ℓ_∞ or Chebyshev Norm). Captures the largest absolute component:

$$\|\mathbf{v}\|_\infty = \lim_{p \rightarrow \infty} \|\mathbf{v}\|_p = \max_{1 \leq i \leq n} |v_i|. \quad (1.13)$$

1.3.2 Rigorous Comparison of Vectors $\mathbf{v}$ and $\mathbf{w}$

Revisiting vectors $\mathbf{v} = [2, 2, 2, 2]^T$ and $\mathbf{w} = [4, 0, 0, 0]^T \in \mathbb{R}^4$:

Example 1.3.1 (Comparative Evaluation Across the Three Norms).

- **In Euclidean ℓ_2 norm:**

$$\|\mathbf{v}\|_2 = \sqrt{2^2 + 2^2 + 2^2 + 2^2} = \sqrt{16} = 4, \quad (1.14)$$

$$\|\mathbf{w}\|_2 = \sqrt{4^2 + 0^2 + 0^2 + 0^2} = \sqrt{16} = 4. \quad (1.15)$$

The two vectors have the exact same **Euclidean length**: $\|\mathbf{v}\|_2 = \|\mathbf{w}\|_2 = 4$.

- **In ℓ_1 norm:**

$$\|\mathbf{v}\|_1 = |2| + |2| + |2| + |2| = 8, \quad (1.16)$$

$$\|\mathbf{w}\|_1 = |4| + |0| + |0| + |0| = 4. \quad (1.17)$$

Under ℓ_1 , we have $\|\mathbf{v}\|_1 = 8 > 4 = \|\mathbf{w}\|_1$.

- **In ℓ_∞ norm:**

$$\|\mathbf{v}\|_\infty = \max\{|2|, |2|, |2|, |2|\} = 2, \quad (1.18)$$

$$\|\mathbf{w}\|_\infty = \max\{|4|, |0|, |0|, |0|\} = 4. \quad (1.19)$$

Under ℓ_∞ , the ranking reverses: $\|\mathbf{v}\|_\infty = 2 < 4 = \|\mathbf{w}\|_\infty$.

Norm Selection Criteria in Real Problems

The choice of norm fundamentally influences algorithmic formulation and model behavior:

- **Differentiability and smoothness:** The squared Euclidean norm $\|\cdot\|_2^2$ is smooth everywhere, leading to closed-form normal equations in least squares.
- **Sparsity promotion:** The non-smooth ℓ_1 norm induces compact support (sparse solutions), pivotal in Lasso regression and compressed sensing.
- **Peak error bounding:** The ℓ_∞ norm is used in minimax approximation and robust control, where bounding worst-case deviation is paramount.

1.4 Standard Inner Product

Definition 1.4.1 (Standard Inner Product in $\mathbb{R}^n$). For two column vectors $\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$, the **standard inner product** (or dot product) is the bilinear map:

$$\langle \mathbf{v}, \mathbf{w} \rangle = \mathbf{v}^T \mathbf{w} = \sum_{i=1}^n v_i w_i = v_1 w_1 + v_2 w_2 + \cdots + v_n w_n. \quad (1.20)$$

1.4.1 Matrix Notation and Output Dimensions

Since $\mathbf{v}, \mathbf{w} \in \mathbb{R}^{n \times 1}$, the transpose $\mathbf{v}^T$ is a $1 \times n$ row matrix:

$$\mathbf{v}^T = \begin{bmatrix} v_1 & v_2 & \dots & v_n \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix}. \quad (1.21)$$

Matrix multiplication $\mathbf{v}^T \mathbf{w}$ yields a 1×1 scalar:

$$\mathbf{v}^T \mathbf{w} = \sum_{i=1}^n v_i w_i \in \mathbb{R}. \quad (1.22)$$

Property 1.4.1 (Direct Link with the Euclidean Norm). *The inner product of a vector with itself recovers the square of its Euclidean norm:*

$$\langle \mathbf{v}, \mathbf{v} \rangle = \mathbf{v}^T \mathbf{v} = \sum_{i=1}^n v_i^2 = \|\mathbf{v}\|_2^2 \implies \|\mathbf{v}\|_2 = \sqrt{\langle \mathbf{v}, \mathbf{v} \rangle} = \sqrt{\mathbf{v}^T \mathbf{v}}. \quad (1.23)$$

Furthermore, two non-zero vectors are **orthogonal** if and only if $\langle \mathbf{v}, \mathbf{w} \rangle = 0$.

1.5 Bases of Vector Spaces and Linear Independence

Definition 1.5.1 (Linear Combination). Given vectors $\{\mathbf{v}_1, \dots, \mathbf{v}_k\} \subseteq V$ and scalars $\alpha_1, \dots, \alpha_k \in \mathbb{R}$, their **linear combination** is:

$$\mathbf{v} = \sum_{i=1}^k \alpha_i \mathbf{v}_i = \alpha_1 \mathbf{v}_1 + \alpha_2 \mathbf{v}_2 + \dots + \alpha_k \mathbf{v}_k. \quad (1.24)$$

Definition 1.5.2 (Basis of a Vector Space). An ordered subset

$$\mathcal{B} = \{\mathbf{v}_1, \dots, \mathbf{v}_k\} \subseteq V \quad (1.25)$$

is a **basis** of V if every element $\mathbf{v} \in V$ can be uniquely expressed as a linear combination of vectors in $\mathcal{B}$.

1.5.1 Prominent Examples of Bases

1. Standard (Canonical) Basis. In $\mathbb{R}^n$, the canonical basis vectors $\mathbf{e}_1, \dots, \mathbf{e}_n$ have a 1 at the corresponding index and 0 elsewhere:

$$\mathbf{e}_1 = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \quad \mathbf{e}_2 = \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \quad \dots, \quad \mathbf{e}_n = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}. \quad (1.26)$$

Any vector $\mathbf{x} = [x_1, \dots, x_n]^T$ is uniquely expressed as $\mathbf{x} = \sum_{i=1}^n x_i \mathbf{e}_i$.

2. Non-Standard Basis in $\mathbb{R}^2$. Consider the two vectors in $\mathbb{R}^2$:

$$\mathcal{B} = \left\{ \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 \\ -1 \end{bmatrix} \right\}. \quad (1.27)$$

To confirm $\mathcal{B}$ forms a basis, we show that for every $[x, y]^T \in \mathbb{R}^2$ there exist unique scalars $\alpha_1, \alpha_2 \in \mathbb{R}$ satisfying:

$$\begin{bmatrix} x \\ y \end{bmatrix} = \alpha_1 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \alpha_2 \begin{bmatrix} 1 \\ -1 \end{bmatrix} = \begin{bmatrix} \alpha_1 + \alpha_2 \\ -\alpha_2 \end{bmatrix}. \quad (1.28)$$

The resulting triangular linear system:

$$\begin{cases} \alpha_1 + \alpha_2 = x \\ -\alpha_2 = y \end{cases} \implies \begin{cases} \alpha_2 = -y \\ \alpha_1 = x + y \end{cases} \quad (1.29)$$

has a unique solution $(\alpha_1, \alpha_2) = (x + y, -y)$ for any arbitrary (x, y) , proving $\mathcal{B}$ is a valid basis.

1.5.2 Linear Independence and Cardinality Constraints

Definition 1.5.3 (Linear Independence). A set of vectors $\{\mathbf{v}_1, \dots, \mathbf{v}_k\}$ is **linearly independent** if the only linear combination yielding the zero vector is the trivial one:

$$\sum_{i=1}^k \alpha_i \mathbf{v}_i = \mathbf{0}_V \implies \alpha_1 = \alpha_2 = \dots = \alpha_k = 0. \quad (1.30)$$

If non-trivial scalars exist such that the sum vanishes, the vectors are **linearly dependent**.

Example 1.5.1 (Checking Linear Independence). We illustrate the verification of linear independence and the notion of basis through two concrete questions:

1. **Three vectors in $\mathbb{R}^2$.** Given the set of vectors:

$$\mathcal{S} = \{[1, 0]^T, [2, 1]^T, [1, 1]^T\} \subset \mathbb{R}^2,$$

determine whether it forms a basis.

Answer: NO.

- *Dimension constraint:* The dimension of $\mathbb{R}^2$ is 2. Any basis of $\mathbb{R}^2$ must contain exactly two vectors. Any three vectors in $\mathbb{R}^2$ are necessarily linearly dependent.
- *Explicit dependence:* Observe that $[1, 1]^T = -[1, 0]^T + [2, 1]^T$, yielding the non-trivial relation:

$$1 \begin{bmatrix} 1 \\ 0 \end{bmatrix} - 1 \begin{bmatrix} 2 \\ 1 \end{bmatrix} + 1 \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}. \quad (1.31)$$

2. **Two collinear vectors in $\mathbb{R}^2$.** Given the set of vectors:

$$\mathcal{S}' = \{[2, 1]^T, [-4, -2]^T\} \subset \mathbb{R}^2,$$

determine whether it forms a basis.

Answer: NO.

The vectors are collinear:

$$\begin{bmatrix} -4 \\ -2 \end{bmatrix} = -2 \begin{bmatrix} 2 \\ 1 \end{bmatrix} \implies 2 \begin{bmatrix} 2 \\ 1 \end{bmatrix} + \begin{bmatrix} -4 \\ -2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}. \quad (1.32)$$

Their span is the line $y = \frac{1}{2}x$, not the full space $\mathbb{R}^2$.

1.6 Classroom Exercises: Unit Spheres and Norm Geometry

1.6.1 Exercise 1: Comparative Norm Calculation in $\mathbb{R}^3$

Given $\mathbf{v} = [1, 2, -3]^T \in \mathbb{R}^3$, we compute the three core norms:

1. **ℓ_1 Norm:**

$$\|\mathbf{v}\|_1 = |1| + |2| + |-3| = 1 + 2 + 3 = 6. \quad (1.33)$$

2. ℓ_2 **Euclidean Norm:**

$$\|\mathbf{v}\|_2 = \sqrt{1^2 + 2^2 + (-3)^2} = \sqrt{1 + 4 + 9} = \sqrt{14} \approx 3.742. \quad (1.34)$$

3. ℓ_∞ **Maximum Norm:**

$$\|\mathbf{v}\|_\infty = \max\{|1|, |2|, |-3|\} = \max\{1, 2, 3\} = 3. \quad (1.35)$$

1.6.2 Exercise 2: Plotting Unit Spheres in $\mathbb{R}^2$

Consider the locus of unit-norm vectors for $\mathbf{v} = [x, y]^T \in \mathbb{R}^2$:

1. Euclidean Unit Sphere $S_2 = \{\mathbf{v} \in \mathbb{R}^2 : \|\mathbf{v}\|_2 = 1\}$. The equation is $\sqrt{x^2 + y^2} = 1 \iff x^2 + y^2 = 1$, which defines the **standard unit circle** centered at $(0, 0)$ with radius $R = 1$.

2. Manhattan Unit Sphere $S_1 = \{\mathbf{v} \in \mathbb{R}^2 : \|\mathbf{v}\|_1 = 1\}$. The equation $|x| + |y| = 1$ expands piecewise in the four Cartesian quadrants:

$$\begin{cases} x + y = 1 \implies y = 1 - x & (x \geq 0, y \geq 0), \\ -x + y = 1 \implies y = 1 + x & (x \leq 0, y \geq 0), \\ -x - y = 1 \implies y = -1 - x & (x \leq 0, y \leq 0), \\ x - y = 1 \implies y = x - 1 & (x \geq 0, y \leq 0). \end{cases} \quad (1.36)$$

Geometrically, this is a **square rotated by 45°** (diamond) with vertices at $(\pm 1, 0)$ and $(0, \pm 1)$.

3. Maximum Unit Sphere $S_\infty = \{\mathbf{v} \in \mathbb{R}^2 : \|\mathbf{v}\|_\infty = 1\}$. The condition $\max\{|x|, |y|\} = 1$ requires one coordinate to have absolute value 1 while the other lies in $[-1, 1]$. Geometrically, this forms an **axis-aligned square** of side length 2 with vertices at $(\pm 1, \pm 1)$.




Figure 1.2: Geometric comparison of unit spheres S_1 (green diamond), S_2 (blue circle), and S_∞ (dashed orange square) in $\mathbb{R}^2$. The four red points mark the intersection $S_1 \cap S_2 \cap S_\infty = \{\pm \mathbf{e}_1, \pm \mathbf{e}_2\}$.

Property 1.6.1 (Common Intersection of Unit Spheres). *The intersection of all three unit loci consists exactly of the canonical basis vectors and their opposites:*

$$S_1 \cap S_2 \cap S_\infty = \left\{ \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \end{bmatrix}, \begin{bmatrix} -1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ -1 \end{bmatrix} \right\} = \{\pm \mathbf{e}_1, \pm \mathbf{e}_2\}. \quad (1.37)$$

Only for these four points do all three norms simultaneously evaluate to unity: $\|\mathbf{v}\|_1 = \|\mathbf{v}\|_2 = \|\mathbf{v}\|_\infty = 1$.

1.6.3 Exercise 3: Basis Verification in $\mathbb{R}^3$

Consider the three vectors in $\mathbb{R}^3$:

$$\mathbf{v}_1 = \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}, \quad \mathbf{v}_2 = \begin{bmatrix} 2 \\ 2 \\ 0 \end{bmatrix}, \quad \mathbf{v}_3 = \begin{bmatrix} -3 \\ -4 \\ 2 \end{bmatrix}. \quad (1.38)$$

Testing whether $\mathbf{v}_3$ lies in $\text{span}(\mathbf{v}_1, \mathbf{v}_2)$:

$$\alpha \mathbf{v}_1 + \beta \mathbf{v}_2 = \mathbf{v}_3 \iff \begin{cases} \alpha + 2\beta = -3 & \text{(component 1),} \\ 2\beta = -4 & \text{(component 2),} \\ 2\alpha = 2 & \text{(component 3).} \end{cases} \quad (1.39)$$

From the third equation, $\alpha = 1$; from the second, $\beta = -2$. Substituting into the first equation:

$$\alpha + 2\beta = 1 + 2(-2) = 1 - 4 = -3. \quad (1.40)$$

The system is consistent, confirming the non-trivial relation:

$$\mathbf{v}_1 - 2\mathbf{v}_2 - \mathbf{v}_3 = \mathbf{0}_{\mathbb{R}^3}. \quad (1.41)$$

Thus the vectors are linearly dependent, and $\{\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3\}$ **does not form a basis of $\mathbb{R}^3$** .

1.7 Three Perspectives on Matrix-Vector Products

1.7.1 Definition and Dimensional Compatibility

A matrix $A \in \mathbb{R}^{n \times m}$ is an ordered rectangular array with n rows and m columns:

$$A = \begin{bmatrix} A_{11} & A_{12} & \cdots & A_{1m} \\ A_{21} & A_{22} & \cdots & A_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ A_{n1} & A_{n2} & \cdots & A_{nm} \end{bmatrix}. \quad (1.42)$$

For the matrix-vector product $\mathbf{c} = A\mathbf{b}$, dimension compatibility requires:

$$A \in \mathbb{R}^{n \times m}, \quad \mathbf{b} \in \mathbb{R}^{m \times 1} \implies \mathbf{c} = A\mathbf{b} \in \mathbb{R}^{n \times 1}. \quad (1.43)$$

The number of columns in A must equal the length of $\mathbf{b}$.

1.7.2 The Three Equivalent Perspectives

Consider the concrete numerical example:

$$A = \begin{bmatrix} 2 & 0 & 1 \\ 3 & -1 & 0 \end{bmatrix} \in \mathbb{R}^{2 \times 3}, \quad \mathbf{b} = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \in \mathbb{R}^3. \quad (1.44)$$

Computational linear algebra examines this product through three complementary lenses:

Perspective 1: Component-Wise Scalar Formula. Each entry c_i of the output accumulates along row and column:

$$c_i = \sum_{j=1}^m A_{ij} b_j \quad (i = 1, \dots, n). \quad (1.45)$$

In our example:

$$c_1 = 2 \cdot 1 + 0 \cdot 2 + 1 \cdot 3 = 2 + 0 + 3 = 5, \quad (1.46)$$

$$c_2 = 3 \cdot 1 + (-1) \cdot 2 + 0 \cdot 3 = 3 - 2 + 0 = 1 \implies \mathbf{c} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}. \quad (1.47)$$

Perspective 2: Row Perspective (Inner Products). Expressing A by its rows $\mathbf{u}_1^T, \dots, \mathbf{u}_n^T$ with $\mathbf{u}_i \in \mathbb{R}^m$:

$$A = \begin{bmatrix} - & \mathbf{u}_1^T & - \\ - & \mathbf{u}_2^T & - \\ & \vdots & \\ - & \mathbf{u}_n^T & - \end{bmatrix} \implies c_i = \langle \mathbf{u}_i, \mathbf{b} \rangle = \mathbf{u}_i^T \mathbf{b}. \quad (1.48)$$

Each output component is the Euclidean dot product between the corresponding row of A and vector $\mathbf{b}$.

Perspective 3: Column Perspective (Linear Combinations). Expressing A through its m columns $\mathbf{w}_1, \dots, \mathbf{w}_m \in \mathbb{R}^n$:

$$A = \begin{bmatrix} | & | & & | \\ \mathbf{w}_1 & \mathbf{w}_2 & \dots & \mathbf{w}_m \\ | & | & & | \end{bmatrix} \implies \mathbf{c} = A\mathbf{b} = \sum_{j=1}^m b_j \mathbf{w}_j. \quad (1.49)$$

Output vector $\mathbf{c}$ is a **linear combination of the columns of A** weighted by the scalars of $\mathbf{b}$:

$$\mathbf{c} = 1 \begin{bmatrix} 2 \\ 3 \end{bmatrix} + 2 \begin{bmatrix} 0 \\ -1 \end{bmatrix} + 3 \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 2 + 0 + 3 \\ 3 - 2 + 0 \end{bmatrix} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}. \quad (1.50)$$

Significance of the Column Perspective

The column perspective is foundational across this course: it reveals that the set of all possible outputs $A\mathbf{b}$ as $\mathbf{b}$ varies is precisely the **column space** $\text{range}(A) = \text{span}(\mathbf{w}_1, \dots, \mathbf{w}_m)$. Solving $A\mathbf{x} = \mathbf{b}$ means finding coefficients $\mathbf{x}$ to reconstruct $\mathbf{b}$ from A 's columns.

1.8 Matrices as Linear Transformations in the Plane

A square matrix $A \in \mathbb{R}^{2 \times 2}$ acts on plane vectors as a linear operator $T_A : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ defined by $T_A(\mathbf{v}) = A\mathbf{v}$. Its geometric behavior is entirely determined by the images of the standard basis vectors:

$$A\mathbf{e}_1 = \text{first column of } A, \quad A\mathbf{e}_2 = \text{second column of } A. \quad (1.51)$$

1.8.1 Core Geometric Examples

1. 90° Clockwise Rotation. Consider:

$$A_{\text{rot}} = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}. \quad (1.52)$$

Acting on the canonical basis:

$$A_{\text{rot}}\mathbf{e}_1 = \begin{bmatrix} 0 \\ -1 \end{bmatrix} = -\mathbf{e}_2, \quad A_{\text{rot}}\mathbf{e}_2 = \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \mathbf{e}_1. \quad (1.53)$$

Horizontal vector $\mathbf{e}_1$ rotates downward to $-\mathbf{e}_2$, while vertical vector $\mathbf{e}_2$ rotates rightward to $\mathbf{e}_1$. For general $\mathbf{v} = [v_1, v_2]^T$, it outputs $[v_2, -v_1]^T$, rotating the entire plane by 90° clockwise.

2. Horizontal Shear Mapping. Consider the upper-triangular matrix:

$$A_{\text{shear}} = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}. \quad (1.54)$$

Its action on the basis vectors gives:

$$A_{\text{shear}}\mathbf{e}_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \mathbf{e}_1, \quad A_{\text{shear}}\mathbf{e}_2 = \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \mathbf{e}_1 + \mathbf{e}_2. \quad (1.55)$$

The horizontal axis remains fixed, while the horizontal coordinate shifts proportionally to vertical position ($x \mapsto x + y$, $y \mapsto y$). Squares deform into parallelograms of equal area ($\det A_{\text{shear}} = 1$).

3. Uniform Scaling and Reflection.

- $A = \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix} = 2I$: uniformly dilates every vector by a factor of 2 ($A\mathbf{v} = 2\mathbf{v}$), preserving directions and angles.
- $A = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$: swaps coordinates ($x \mapsto y$, $y \mapsto x$), performing reflection across the line $y = x$.

1.9 Introduction to Computational Complexity and Flops

In numerical linear algebra, establishing theoretical solvability is insufficient: one must quantify the required elementary floating-point operations (**flops** — additions, subtractions, multiplications, divisions).

1.9.1 Operation Count for the Inner Product

For $\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$, computing $\langle \mathbf{v}, \mathbf{w} \rangle = \sum_{i=1}^n v_i w_i$ requires:

1. n multiplications ($v_1 w_1, \dots, v_n w_n$);
2. $n - 1$ additions to sum the partial products.

The exact count is:

$$\text{Flops}(n) = n + (n - 1) = 2n - 1 \implies \mathcal{O}(n). \quad (1.56)$$

1.9.2 Linear Time Scaling Model

Assuming execution time is proportional to operation count: if computing an inner product of length $n = 10^7$ takes 3 seconds, tripling the size to $n' = 3 \cdot 10^7$ yields:

$$\text{Estimated time} = 3 \times (\text{base time}) = 3 \times 3 \text{ s} = 9 \text{ s}. \quad (1.57)$$

Linear complexity $\mathcal{O}(n)$ guarantees runtime scales proportionally with dimension.

Remark 1.9.1 (Open Question: The Cost of Matrix-Vector Multiplication). For a square matrix $A \in \mathbb{R}^{n \times n}$ and vector $\mathbf{b} \in \mathbb{R}^n$:

- Evaluating $A\mathbf{b}$ requires computing n independent dot products of length n , each costing $2n - 1$:

$$\text{Flops}_{\text{mat-vec}}(n) = n(2n - 1) = 2n^2 - n \implies \mathcal{O}(n^2). \quad (1.58)$$

- If computation for $n = 1\,000$ takes 3 seconds, doubling to $n' = 2\,000$ quadruples the expected time: $(2)^2 \times 3 \text{ s} = 12 \text{ s}$.

1.10 Summary of Core Mathematical Concepts

Object / Concept	Mathematical Definition	Key Property / Complexity
ℓ_1 Norm	$\ \mathbf{v}\ _1 = \sum_{i=1}^n v_i $	Unit sphere is a 45° diamond; induces sparsity in models.
ℓ_2 Norm	$\ \mathbf{v}\ _2 = \sqrt{\sum_{i=1}^n v_i^2} = \sqrt{\mathbf{v}^T \mathbf{v}}$	Standard rotation-invariant Euclidean length; smooth and differentiable.
ℓ_∞ Norm	$\ \mathbf{v}\ _\infty = \max_{1 \leq i \leq n} v_i $	Unit sphere is an axis-aligned square of side length 2.
Inner Product	$\langle \mathbf{v}, \mathbf{w} \rangle = \mathbf{v}^T \mathbf{w} = \sum_{i=1}^n v_i w_i$	Arithmetic operations: $2n - 1 \implies \mathcal{O}(n)$.
Basis	Linearly independent set with span = V	Unique coordinate representation for every vector.
Product $A\mathbf{b}$ (Row view)	$c_i = \mathbf{u}_i^T \mathbf{b}$	n dot products between rows of A and vector $\mathbf{b}$.
Product $A\mathbf{b}$ (Column view)	$\mathbf{c} = \sum_{j=1}^m b_j \mathbf{w}_j$	Linear combination of columns weighted by $\mathbf{b}$; spans $\text{range}(A)$.
90° Clockwise Rotation	$A = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$	Maps $\mathbf{e}_1 \mapsto -\mathbf{e}_2$ and $\mathbf{e}_2 \mapsto \mathbf{e}_1$.
Reflection across $y = x$	$A = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$	Maps $\mathbf{e}_1 \mapsto \mathbf{e}_2$ and $\mathbf{e}_2 \mapsto \mathbf{e}_1$, swapping coordinates.

Table 1.1: Summary of algebraic, geometric, and computational concepts covered in Chapter 1.

Chapter 2

Matrix Multiplication, Computational Complexity, and Orthogonal Matrices

This chapter investigates higher-level algebraic operations between matrices, focusing on matrix-matrix multiplication, its computational complexity and scaling laws, and the structural properties of the fundamental subspaces associated with a matrix. We then formalize vector orthogonality and orthogonal matrices, which constitute cornerstone building blocks for designing stable numerical algorithms and canonical matrix factorizations in numerical linear algebra [2].

2.1 Matrix-Matrix Multiplication

2.1.1 Dimensional Compatibility Condition

Let A and B be real matrices. The matrix product $C = AB$ is well-defined **if and only if** the number of columns in the first factor A equals the number of rows in the second factor B .

Definition 2.1.1 (Dimensional Compatibility of Matrix Products). Given $n, m, p \in \mathbb{N}$, let $A \in \mathbb{R}^{n \times m}$ and $B \in \mathbb{R}^{m \times p}$. Matrix multiplication produces a matrix $C = AB \in \mathbb{R}^{n \times p}$ whose dimensions match the rows of the first factor and the columns of the second:

$$A \in \mathbb{R}^{n \times m}, \quad B \in \mathbb{R}^{m \times p} \implies C = AB \in \mathbb{R}^{n \times p}. \quad (2.1)$$

If the inner dimensions differ, the product is undefined. In general, even

for square matrices ($n = m = p$), matrix multiplication is non-commutative:

$$AB \neq BA \quad (\text{in general}). \quad (2.2)$$




Figure 2.1: Dimensional compatibility in matrix multiplication: entry C_{ij} is formed by the inner product of row $\mathbf{u}_i^T$ of A and column $\mathbf{w}_j$ of B .

2.1.2 Component-Wise Definition

Each entry C_{ij} located at row i ($1 \leq i \leq n$) and column j ($1 \leq j \leq p$) of the resulting matrix is computed as the standard Euclidean inner product between the i -th row of A and the j -th column of B :

Definition 2.1.2 (Component-Wise Matrix Multiplication Formula). Given $A \in \mathbb{R}^{n \times m}$ and $B \in \mathbb{R}^{m \times p}$, entry (i, j) of $C = AB$ is:

$$C_{ij} = \sum_{k=1}^m A_{ik} B_{kj} = A_{i1} B_{1j} + A_{i2} B_{2j} + \cdots + A_{im} B_{mj}. \quad (2.3)$$

In vector notation, denoting the i -th row of A by $\mathbf{u}_i^T$ and the j -th column of B by $\mathbf{w}_j$:

$$C_{ij} = \langle \mathbf{u}_i, \mathbf{w}_j \rangle = \mathbf{u}_i^T \mathbf{w}_j. \quad (2.4)$$

2.1.3 Step-by-Step Numerical Example

Example 2.1.1 (Detailed Rectangular Matrix Multiplication). Consider the matrices:

$$A = \begin{bmatrix} 1 & 2 \\ 0 & 3 \end{bmatrix} \in \mathbb{R}^{2 \times 2}, \quad B = \begin{bmatrix} 0 & 1 & -1 \\ 1 & 0 & 2 \end{bmatrix} \in \mathbb{R}^{2 \times 3}. \quad (2.5)$$

Since A has $m = 2$ columns and B has $m = 2$ rows, product $C = AB$ is well-defined and yields a 2×3 matrix. We compute each component analytically:

First row of C (row 1 of A against columns of B):

$$C_{11} = A_{11}B_{11} + A_{12}B_{21} = 1 \cdot 0 + 2 \cdot 1 = 0 + 2 = 2, \quad (2.6)$$

$$C_{12} = A_{11}B_{12} + A_{12}B_{22} = 1 \cdot 1 + 2 \cdot 0 = 1 + 0 = 1, \quad (2.7)$$

$$C_{13} = A_{11}B_{13} + A_{12}B_{23} = 1 \cdot (-1) + 2 \cdot 2 = -1 + 4 = 3. \quad (2.8)$$

Second row of C (row 2 of A against columns of B):

$$C_{21} = A_{21}B_{11} + A_{22}B_{21} = 0 \cdot 0 + 3 \cdot 1 = 0 + 3 = 3, \quad (2.9)$$

$$C_{22} = A_{21}B_{12} + A_{22}B_{22} = 0 \cdot 1 + 3 \cdot 0 = 0 + 0 = 0, \quad (2.10)$$

$$C_{23} = A_{21}B_{13} + A_{22}B_{23} = 0 \cdot (-1) + 3 \cdot 2 = 0 + 6 = 6. \quad (2.11)$$

Assembling the calculated components yields:

$$C = AB = \begin{bmatrix} 2 & 1 & 3 \\ 3 & 0 & 6 \end{bmatrix} \in \mathbb{R}^{2 \times 3}. \quad (2.12)$$

2.2 Computational Complexity and Scaling Law

2.2.1 Rigorous FLOP Counting

Assessing numerical algorithm execution time requires quantifying floating-point operations (FLOPs).

Property 2.2.1 (Computational Cost of Matrix Multiplication).

1. **Cost of a single entry:** Evaluating $C_{ij} = \sum_{k=1}^m A_{ik}B_{kj}$ requires m scalar multiplications and $m - 1$ additions:

$$\text{FLOPs per entry} = m + (m - 1) = 2m - 1 \implies \mathcal{O}(m). \quad (2.13)$$

2. **General rectangular case** ($A \in \mathbb{R}^{n \times m}, B \in \mathbb{R}^{m \times p}$): Computing all $n \cdot p$ independent entries yields:

$$\text{Total FLOPs} = n \cdot p \cdot (2m - 1) = 2nmp - np \implies \mathcal{O}(nmp). \quad (2.14)$$

3. **Square matrix case** ($n = m = p$): When both matrices belong to $\mathbb{R}^{n \times n}$, matrix C has n^2 entries:

$$\text{Total FLOPs}(n) = n^2 \cdot (2n - 1) = 2n^3 - n^2 \implies \mathcal{O}(n^3). \quad (2.15)$$

The leading term $2n^3$ dictates cubic growth with respect to dimension n .

2.2.2 Cubic Time Scaling Law

Let $T(n) \approx \gamma n^3$ model execution time on fixed hardware, where $\gamma > 0$ absorbs cache architecture, CPU clock speed, and compiler optimization.

Example 2.2.1 (Predicting Execution Time under Doubled Dimension). Suppose multiplying two dense square matrices of size $n = 1000$ takes 3 seconds of CPU time:

$$T(1000) \approx \gamma \cdot (1000)^3 = 3 \text{ s.} \quad (2.16)$$

Doubling the dimension to $n' = 2000$, the cubic scaling law implies:

$$T(2000) \approx \gamma \cdot (2000)^3 = \gamma \cdot (2 \cdot 1000)^3 = 8 \cdot (\gamma \cdot 1000^3) = 8 \cdot T(1000). \quad (2.17)$$

The execution time scales by a factor of 8:

$$T(2000) \approx 8 \cdot 3 \text{ s} = 24 \text{ seconds.} \quad (2.18)$$

2.2.3 Comparative Overview of Core Linear Algebra Operations

Table 2.1 illustrates how the three foundational linear algebra operations scale when the linear data dimension n increases by a factor of 10 (from $n = 10^3$ to $n = 10^4$).

Operation	Operands	FLOP Formula	Cost	Scaling Factor (10 $\times$)
Inner Product	$\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$	$2n - 1$	$\mathcal{O}(n)$	$10^1 = 10$
Matrix-Vector	$A \in \mathbb{R}^{n \times n}, \mathbf{b} \in \mathbb{R}^n$	$2n^2 - n$	$\mathcal{O}(n^2)$	$10^2 = 100$
Matrix-Matrix	$A, B \in \mathbb{R}^{n \times n}$	$2n^3 - n^2$	$\mathcal{O}(n^3)$	$10^3 = 1000$

Table 2.1: Complexity comparison and scaling factors for a tenfold increase in dimension.

If an $n = 1000$ matrix-matrix multiplication runs in 1 second, the same task for $n = 10\,000$ requires 1000 seconds (≈ 16.6 minutes).

2.3 Fundamental Properties: Rank, Subspaces, and Invertibility

2.3.1 Subspaces Associated with a Linear Map

Let $A \in \mathbb{R}^{n \times m}$ define a linear operator $T_A : \mathbb{R}^m \rightarrow \mathbb{R}^n$ via $T_A(\mathbf{x}) = A\mathbf{x}$.

Definition 2.3.1 (Rank of a Matrix). The **rank** of A , denoted by $\text{rank}(A)$, is the dimension of the vector space spanned by its columns. By the fundamental theorem of linear algebra, column rank and row rank always coincide:

$$\text{rank}(A) = \dim(\text{col}(A)) = \dim(\text{row}(A)) \leq \min(n, m). \quad (2.19)$$

When $\text{rank}(A) = \min(n, m)$, the matrix is said to have **full rank**.

Definition 2.3.2 (Kernel or Nullspace). The **kernel** (or *nullspace*) of A is the subspace of $\mathbb{R}^m$ containing all vectors mapped to the zero vector in $\mathbb{R}^n$:

$$\ker(A) = \{\mathbf{x} \in \mathbb{R}^m : A\mathbf{x} = \mathbf{0}_{\mathbb{R}^n}\} \subseteq \mathbb{R}^m. \quad (2.20)$$

Definition 2.3.3 (Image or Range). The **image** (or *range*) of A is the subspace of $\mathbb{R}^n$ spanned by all linear combinations of the columns of A :

$$\text{Im}(A) = \text{span}(\text{columns of } A) = \{A\mathbf{x} : \mathbf{x} \in \mathbb{R}^m\} \subseteq \mathbb{R}^n. \quad (2.21)$$

We have $\dim(\text{Im}(A)) = \text{rank}(A)$.

Theorem 2.3.1 (Rank-Nullity Theorem). *For any matrix $A \in \mathbb{R}^{n \times m}$ with m columns, the dimension of the image plus the dimension of the kernel equals the dimension of the domain:*

$$\dim(\text{Im}(A)) + \dim(\ker(A)) = m. \quad (2.22)$$




Figure 2.2: Geometric representation of fundamental subspaces and the Rank-Nullity Theorem: $\dim(\ker(A)) + \dim(\text{Im}(A)) = m$. The restriction of A to the orthogonal complement $(\ker(A))^\perp$ is an isomorphism onto $\text{Im}(A)$.

2.3.2 Invertibility of Square Matrices

For square matrices $A \in \mathbb{R}^{n \times n}$, rank, nullity, and invertibility merge into an overarching equivalence theorem.

Definition 2.3.4 (Invertible or Non-Singular Matrix). A matrix $A \in \mathbb{R}^{n \times n}$ is called **invertible** if there exists a unique matrix $A^{-1} \in \mathbb{R}^{n \times n}$ satisfying:

$$AA^{-1} = I_n \quad \text{and} \quad A^{-1}A = I_n. \quad (2.23)$$

Theorem 2.3.2 (Equivalent Characterizations of Invertibility). *Let $A \in \mathbb{R}^{n \times n}$. The following statements are equivalent:*

1. A is invertible.
2. $\text{rank}(A) = n$ (A has full rank).
3. $\det(A) \neq 0$.
4. $\ker(A) = \{\mathbf{0}_{\mathbb{R}^n}\}$ (trivial nullspace).
5. The columns of A form a basis of $\mathbb{R}^n$.
6. Linear system $A\mathbf{x} = \mathbf{b}$ has a unique solution $\mathbf{x} = A^{-1}\mathbf{b}$ for every $\mathbf{b} \in \mathbb{R}^n$.

2.3.3 Guided Example: Rank, Kernel, and Image Calculation

Example 2.3.1 (Full Subspace Analysis of a Singular 3×3 Matrix). Consider the matrix:

$$A = \begin{bmatrix} 1 & 2 & 0 \\ 2 & 4 & 0 \\ 0 & 0 & 1 \end{bmatrix} \in \mathbb{R}^{3 \times 3}. \quad (2.24)$$

1. Rank via Gaussian Elimination: Apply elementary row operation $R_2 \leftarrow R_2 - 2R_1$:

$$\begin{bmatrix} 1 & 2 & 0 \\ 2 & 4 & 0 \\ 0 & 0 & 1 \end{bmatrix} \xrightarrow{R_2 \rightarrow R_2 - 2R_1} \begin{bmatrix} 1 & 2 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \xrightarrow{R_2 \leftrightarrow R_3} \begin{bmatrix} 1 & 2 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}. \quad (2.25)$$

The row echelon form contains 2 non-zero pivots. Therefore:

$$\text{rank}(A) = 2. \quad (2.26)$$

Since $\text{rank}(A) = 2 < 3$, matrix A is **not invertible** ($\det(A) = 0$).

2. Image Space $\text{Im}(A)$: Examining the columns of A :

$$\mathbf{a}_1 = \begin{bmatrix} 1 \\ 2 \\ 0 \end{bmatrix}, \quad \mathbf{a}_2 = \begin{bmatrix} 2 \\ 4 \\ 0 \end{bmatrix} = 2\mathbf{a}_1, \quad \mathbf{a}_3 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}. \quad (2.27)$$

The second column is collinear to the first and hence redundant. Columns $\mathbf{a}_1$ and $\mathbf{a}_3$ are linearly independent and form a basis for the image:

$$\text{Im}(A) = \text{span} \left\{ \begin{bmatrix} 1 \\ 2 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \right\}, \quad \dim(\text{Im}(A)) = 2. \quad (2.28)$$

3. Kernel Space $\ker(A)$: Solve homogeneous linear system $A\mathbf{x} = \mathbf{0}_{\mathbb{R}^3}$ with $\mathbf{x} = [x_1, x_2, x_3]^T$:

$$\begin{bmatrix} 1 & 2 & 0 \\ 2 & 4 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} \implies \begin{cases} x_1 + 2x_2 = 0 \\ x_3 = 0 \end{cases} \implies \begin{cases} x_1 = -2x_2 \\ x_3 = 0 \end{cases}. \quad (2.29)$$

Setting free variable $x_2 = t \in \mathbb{R}$:

$$\mathbf{x} = \begin{bmatrix} -2t \\ t \\ 0 \end{bmatrix} = t \begin{bmatrix} -2 \\ 1 \\ 0 \end{bmatrix} \implies \ker(A) = \text{span} \left\{ \begin{bmatrix} -2 \\ 1 \\ 0 \end{bmatrix} \right\}. \quad (2.30)$$

The nullity is $\dim(\ker(A)) = 1$. Verifying the Rank-Nullity Theorem:

$$\dim(\text{Im}(A)) + \dim(\ker(A)) = 2 + 1 = 3 = m. \quad (2.31)$$

2.4 Vector Orthogonality and Orthonormality

2.4.1 Core Definitions

Definition 2.4.1 (Orthogonal Vectors). Two vectors $\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$ are **orthogonal** ($\mathbf{v} \perp \mathbf{w}$) if and only if their standard dot product vanishes:

$$\langle \mathbf{v}, \mathbf{w} \rangle = \mathbf{v}^T \mathbf{w} = 0. \quad (2.32)$$

Definition 2.4.2 (Orthogonal and Orthonormal Sets). A set of vectors $\{\mathbf{v}_1, \dots, \mathbf{v}_k\} \subset \mathbb{R}^n$ is called:

1. **Orthogonal** if vectors are mutually perpendicular:

$$\langle \mathbf{v}_i, \mathbf{v}_j \rangle = 0 \quad \forall i \neq j. \quad (2.33)$$

2. **Orthonormal** if it is orthogonal and every vector has unit Euclidean norm:

$$\langle \mathbf{v}_i, \mathbf{v}_j \rangle = \delta_{ij} = \begin{cases} 1 & \text{if } i = j, \\ 0 & \text{if } i \neq j, \end{cases} \quad (2.34)$$

where δ_{ij} is the Kronecker delta.

2.4.2 Constructive Normalization Exercises in $\mathbb{R}^2$

Example 2.4.1 (Constructing and Normalizing an Orthogonal Pair).

1. **Finding an orthogonal vector:** Given vector $\mathbf{v}_1 = [2, 1]^T \in \mathbb{R}^2$, find a non-zero vector $\mathbf{v}_2 = [x, y]^T$ such that $\langle \mathbf{v}_1, \mathbf{v}_2 \rangle = 0$:

$$2x + y = 0 \implies y = -2x. \quad (2.35)$$

Choosing $x = 1$ yields $y = -2$, producing the orthogonal pair:

$$\mathbf{v}_1 = \begin{bmatrix} 2 \\ 1 \end{bmatrix}, \quad \mathbf{v}_2 = \begin{bmatrix} 1 \\ -2 \end{bmatrix}, \quad \langle \mathbf{v}_1, \mathbf{v}_2 \rangle = 2(1) + 1(-2) = 0. \quad (2.36)$$

2. **Euclidean Normalization (Orthonormalization):** Dividing each vector by its Euclidean ℓ_2 norm:

$$\|\mathbf{v}_1\|_2 = \sqrt{2^2 + 1^2} = \sqrt{5} \implies \mathbf{u}_1 = \frac{\mathbf{v}_1}{\|\mathbf{v}_1\|_2} = \begin{bmatrix} 2/\sqrt{5} \\ 1/\sqrt{5} \end{bmatrix}, \quad (2.37)$$

$$\|\mathbf{v}_2\|_2 = \sqrt{1^2 + (-2)^2} = \sqrt{5} \implies \mathbf{u}_2 = \frac{\mathbf{v}_2}{\|\mathbf{v}_2\|_2} = \begin{bmatrix} 1/\sqrt{5} \\ -2/\sqrt{5} \end{bmatrix}. \quad (2.38)$$

Verifying orthonormality:

$$\|\mathbf{u}_1\|_2 = 1, \quad \|\mathbf{u}_2\|_2 = 1, \quad (2.39)$$

$$\langle \mathbf{u}_1, \mathbf{u}_2 \rangle = \frac{2}{\sqrt{5}} \frac{1}{\sqrt{5}} + \frac{1}{\sqrt{5}} \left(-\frac{2}{\sqrt{5}} \right) = \frac{2}{5} - \frac{2}{5} = 0. \quad (2.40)$$

The set $\{\mathbf{u}_1, \mathbf{u}_2\}$ forms an orthonormal basis of $\mathbb{R}^2$.

2.5 Orthogonal Matrices

2.5.1 Definition and Transpose Properties

Recall that for general matrix $A \in \mathbb{R}^{n \times m}$, the transpose $A^T \in \mathbb{R}^{m \times n}$ has entries $(A^T)_{ij} = A_{ji}$. For any compatible matrices, the transpose of a product reverses order:

$$(AB)^T = B^T A^T. \quad (2.41)$$

Definition 2.5.1 (Orthogonal Matrix). A square matrix $U \in \mathbb{R}^{n \times n}$ is **orthogonal** if and only if its transpose equals its inverse:

$$U^T U = U U^T = I_n \iff U^{-1} = U^T. \quad (2.42)$$

2.5.2 Column and Row Characterization

Theorem 2.5.1 (Column and Row Orthonormality). *A square matrix $U \in \mathbb{R}^{n \times n}$ is orthogonal if and only if its columns form an orthonormal basis of $\mathbb{R}^n$. Equivalently, U is orthogonal if and only if its rows form an orthonormal basis of $\mathbb{R}^n$.*

Proof. Partition U into column vectors $\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_n \in \mathbb{R}^n$:

$$U = [\mathbf{u}_1 \quad \mathbf{u}_2 \quad \dots \quad \mathbf{u}_n]. \quad (2.43)$$

Then U^T is partitioned by rows:

$$U^T = \begin{bmatrix} \mathbf{u}_1^T \\ \mathbf{u}_2^T \\ \vdots \\ \mathbf{u}_n^T \end{bmatrix}. \quad (2.44)$$

Computing block matrix product $U^T U$:

$$\begin{aligned} U^T U &= \begin{bmatrix} \mathbf{u}_1^T \\ \mathbf{u}_2^T \\ \vdots \\ \mathbf{u}_n^T \end{bmatrix} [\mathbf{u}_1 \quad \mathbf{u}_2 \quad \dots \quad \mathbf{u}_n] \\ &= \begin{bmatrix} \mathbf{u}_1^T \mathbf{u}_1 & \mathbf{u}_1^T \mathbf{u}_2 & \dots & \mathbf{u}_1^T \mathbf{u}_n \\ \mathbf{u}_2^T \mathbf{u}_1 & \mathbf{u}_2^T \mathbf{u}_2 & \dots & \mathbf{u}_2^T \mathbf{u}_n \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{u}_n^T \mathbf{u}_1 & \mathbf{u}_n^T \mathbf{u}_2 & \dots & \mathbf{u}_n^T \mathbf{u}_n \end{bmatrix} = [\langle \mathbf{u}_i, \mathbf{u}_j \rangle]_{i,j=1}^n. \end{aligned} \quad (2.45)$$

For this product to equal identity matrix I_n , we must have:

$$\langle \mathbf{u}_i, \mathbf{u}_j \rangle = \delta_{ij} \quad \forall i, j \in \{1, \dots, n\}. \quad (2.46)$$

Thus the columns are mutually orthogonal and have unit Euclidean norm. The identical reasoning applied to $U U^T = I_n$ proves that rows are orthonormal. $\square$

2.5.3 Prototypical Example: Planar Rotation Matrix

Example 2.5.1 (Planar Rotation Matrix). Let $\theta \in \mathbb{R}$. Consider clockwise rotation matrix:

$$U = \begin{bmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{bmatrix}. \quad (2.47)$$

Verifying $U^T U = I_2$:

$$\begin{aligned} U^T U &= \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{bmatrix} \\ &= \begin{bmatrix} \cos^2 \theta + \sin^2 \theta & \cos \theta \sin \theta - \sin \theta \cos \theta \\ \sin \theta \cos \theta - \cos \theta \sin \theta & \sin^2 \theta + \cos^2 \theta \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = I_2. \end{aligned} \quad (2.48)$$

Matrix U is orthogonal for all angles θ .

2.5.4 Closure under Matrix Products

Proposition 2.5.1 (Orthogonal Group and Product Closure). *If $U, V \in \mathbb{R}^{n \times n}$ are orthogonal matrices, their product $Q = UV$ is also orthogonal.*

Proof. Compute QQ^T using the reversal rule $(UV)^T = V^T U^T$:

$$QQ^T = (UV)(UV)^T = (UV)(V^T U^T). \quad (2.49)$$

By associativity of matrix products:

$$QQ^T = U(VV^T)U^T. \quad (2.50)$$

Since V is orthogonal, $VV^T = I_n$:

$$QQ^T = UI_n U^T = UU^T = I_n. \quad (2.51)$$

Similarly, $Q^T Q = V^T (U^T U) V = V^T I_n V = I_n$. Hence $Q = UV$ is orthogonal. $\square$

2.6 The Four Noteworthy Questions on Orthogonal Matrices

2.6.1 Question 1: Upper Bound on Individual Matrix Entries

Example 2.6.1 (Can an Orthogonal Matrix Entry Exceed 1?). **Question:** Can an orthogonal matrix $U \in \mathbb{R}^{n \times n}$ have an entry $U_{ij} > 1$?

Answer: NO.

Proof: Let $\mathbf{u}_j$ denote the j -th column of U . Since U is orthogonal, $\mathbf{u}_j$ has unit Euclidean norm:

$$\|\mathbf{u}_j\|_2^2 = \sum_{k=1}^n U_{kj}^2 = 1. \quad (2.52)$$

Since each squared entry $U_{kj}^2 \geq 0$, every single term is bounded by the total sum:

$$U_{ij}^2 \leq \sum_{k=1}^n U_{kj}^2 = 1 \implies |U_{ij}| \leq 1 \implies -1 \leq U_{ij} \leq 1. \quad (2.53)$$

No entry of an orthogonal matrix can exceed 1 in absolute value.

2.6.2 Question 2: Orthogonal Diagonal Matrices

Example 2.6.2 (Characterizing Orthogonal Diagonal Matrices). **Question:** When is a diagonal matrix $D \in \mathbb{R}^{n \times n}$ orthogonal?

Solution: Let $D = \text{diag}(d_1, d_2, \dots, d_n)$. Being diagonal, $D^T = D$. Computing $D^T D$:

$$D^T D = D^2 = \text{diag}(d_1^2, d_2^2, \dots, d_n^2). \quad (2.54)$$

For $D^T D = I_n$, each diagonal entry must satisfy:

$$d_i^2 = 1 \iff d_i = \pm 1 \quad \forall i = 1, \dots, n. \quad (2.55)$$

Conclusion: A diagonal matrix is orthogonal if and only if all diagonal entries are $+1$ or -1 . There exist exactly 2^n distinct orthogonal diagonal matrices in $\mathbb{R}^{n \times n}$.

2.6.3 Question 3: Orthogonal Upper Triangular Matrices

Example 2.6.3 (Characterizing Orthogonal Upper Triangular Matrices).

Question: When is an upper triangular matrix $R \in \mathbb{R}^{n \times n}$ orthogonal?

Proof by column induction: Let R be upper triangular ($R_{ij} = 0$ for all $i > j$):

$$R = \begin{bmatrix} r_{11} & r_{12} & r_{13} & \cdots & r_{1n} \\ 0 & r_{22} & r_{23} & \cdots & r_{2n} \\ 0 & 0 & r_{33} & \cdots & r_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & r_{nn} \end{bmatrix}. \quad (2.56)$$

Since R is orthogonal, its columns $\mathbf{r}_1, \dots, \mathbf{r}_n$ must form an orthonormal basis:

1. **First column** $\mathbf{r}_1 = [r_{11}, 0, \dots, 0]^T$: Unit norm requires $r_{11}^2 = 1 \implies r_{11} = \pm 1$.
2. **Second column** $\mathbf{r}_2 = [r_{12}, r_{22}, 0, \dots, 0]^T$: Orthogonality to $\mathbf{r}_1$ requires:
$$\langle \mathbf{r}_1, \mathbf{r}_2 \rangle = r_{11}r_{12} = 0 \implies r_{12} = 0 \quad (\text{since } r_{11} = \pm 1 \neq 0). \quad (2.57)$$
Unit norm gives $\|\mathbf{r}_2\|_2^2 = r_{12}^2 + r_{22}^2 = 0 + r_{22}^2 = 1 \implies r_{22} = \pm 1$.
3. **Third column** $\mathbf{r}_3 = [r_{13}, r_{23}, r_{33}, 0, \dots, 0]^T$: Orthogonality to $\mathbf{r}_1$ and $\mathbf{r}_2$ forces $r_{11}r_{13} = 0 \implies r_{13} = 0$ and $r_{22}r_{23} = 0 \implies r_{23} = 0$. Hence $r_{33} = \pm 1$.
4. **Inductive step**: For every column j , all off-diagonal entries vanish ($r_{ij} = 0$ for $i < j$), and diagonal entries satisfy $r_{jj} = \pm 1$.

Conclusion: An upper triangular matrix is orthogonal if and only if it is diagonal with diagonal entries ± 1 .

2.6.4 Question 4: Euclidean Isometry and Norm Preservation

Theorem 2.6.1 (Isometry Property of Orthogonal Matrices). *Orthogonal matrices strictly preserve Euclidean vector lengths and standard inner products:*

$$\forall U \in \mathbb{R}^{n \times n} \text{ orthogonal}, \quad \forall \mathbf{x} \in \mathbb{R}^n \implies \|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2. \quad (2.58)$$

Furthermore, inner products between arbitrary pairs of vectors are preserved:

$$\langle U\mathbf{x}, U\mathbf{y} \rangle = \langle \mathbf{x}, \mathbf{y} \rangle \quad \forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^n. \quad (2.59)$$

Proof. Recall that for any vector $\mathbf{v}$, $\|\mathbf{v}\|_2^2 = \mathbf{v}^T \mathbf{v}$. Setting $\mathbf{v} = U\mathbf{x}$:

$$\|U\mathbf{x}\|_2^2 = (U\mathbf{x})^T (U\mathbf{x}). \quad (2.60)$$

Applying the transpose product rule $(U\mathbf{x})^T = \mathbf{x}^T U^T$:

$$\|U\mathbf{x}\|_2^2 = \mathbf{x}^T (U^T U) \mathbf{x}. \quad (2.61)$$

Since U is orthogonal, $U^T U = I_n$:

$$\|U\mathbf{x}\|_2^2 = \mathbf{x}^T I_n \mathbf{x} = \mathbf{x}^T \mathbf{x} = \|\mathbf{x}\|_2^2. \quad (2.62)$$

Taking the positive square root yields $\|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2$. Similarly:

$$\langle U\mathbf{x}, U\mathbf{y} \rangle = (U\mathbf{x})^T (U\mathbf{y}) = \mathbf{x}^T (U^T U) \mathbf{y} = \mathbf{x}^T I_n \mathbf{y} = \mathbf{x}^T \mathbf{y} = \langle \mathbf{x}, \mathbf{y} \rangle. \quad (2.63)$$

□




Figure 2.3: Euclidean isometry of orthogonal transformation U : length conservation and rigid rotation.

Fundamental Numerical Significance of Isometry

The identity $\|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2$ is the theoretical backbone of backward error stability in numerical linear algebra. Since orthogonal matrices cannot amplify vector norms, multiplying an error or perturbation $\mathbf{e}$ by U :

$$\|U(\mathbf{x} + \mathbf{e}) - U\mathbf{x}\|_2 = \|U\mathbf{e}\|_2 = \|\mathbf{e}\|_2 \quad (2.64)$$

preserves the exact magnitude of the perturbation without magnification. Algorithms built on orthogonal transformations (Householder reflectors, Givens rotations, and QR factorizations) provide exceptional backward stability guarantees.

2.7 Chapter Summary Reference

Table 2.2 summarizes the key analytical formulas, complexities, and properties developed in this chapter.

Property / Object	Mathematical Formulation	Key Characteristic
Matrix Multiplication	$C_{ij} = \sum_{k=1}^m A_{ik} B_{kj}$	Requires $A \in \mathbb{R}^{n \times m}, B \in \mathbb{R}^{m \times p}$
Square Product Cost	$2n^3 - n^2$ FLOPs $\implies \mathcal{O}(n^3)$	Cubic law: $T(2n) \approx 8T(n)$
Rank-Nullity Theorem	$\dim(\text{Im}(A)) + \dim(\ker(A)) = m$	m is number of columns of A
Orthonormal Vectors	$\langle \mathbf{u}_i, \mathbf{u}_j \rangle = \delta_{ij}$	Mutually orthogonal, $\ \mathbf{u}_i\ _2 = 1$
Orthogonal Matrix	$U^T U = U U^T = I_n$	Equivalent to $U^{-1} = U^T$
Rows and Columns of U	Form orthonormal bases of $\mathbb{R}^n$	Direct consequence of $U^T U = I_n$
Entries of Orthogonal U	$-1 \leq U_{ij} \leq 1$	Derived from $\sum_{i=1}^n U_{ij}^2 = 1$
Orthogonal Diagonal Matrix	$D = \text{diag}(\pm 1, \dots, \pm 1)$	2^n distinct matrices in $\mathbb{R}^{n \times n}$
Orthogonal Triangular Matrix	$R = \text{diag}(\pm 1, \dots, \pm 1)$	Column induction proof
Euclidean Isometry	$\ U\mathbf{x}\ _2 = \ \mathbf{x}\ _2$	Rigid map: length-preserving
Inner Product Invariance	$\langle U\mathbf{x}, U\mathbf{y} \rangle = \langle \mathbf{x}, \mathbf{y} \rangle$	Preserves angles between vectors
Transpose of Product	$(AB)^T = B^T A^T$	Reverses factor order

Table 2.2: Comprehensive reference table of definitions, computational complexities, and orthogonal matrix properties.

Chapter 3

Matrices with Orthonormal Columns, Eigenvalues, Spectral Decomposition, and Quadratic Forms

This chapter investigates the algebraic and geometric properties of rectangular matrices with orthonormal columns, systematically comparing the products $V^T V$ and $V V^T$. We then introduce spectral theory, the notions of eigenvalue and eigenvector, diagonalization, and outer product (dyadic) expansion, culminating in the Spectral Theorem for symmetric matrices. Finally, we analyze quadratic forms, the Rayleigh quotient inequality, the classification of symmetric positive definite and semidefinite matrices (SPD and SPSD), and the natural emergence of positive semidefinite structures via the Gram matrices $B^T B$ and $B B^T$ [2].

3.1 Review of Square Orthogonal Matrices and Isometry

In the preceding chapter, we introduced the concept of an orthogonal matrix in Euclidean space $\mathbb{R}^n$. We briefly summarize the equivalent formulations that govern its behavior.

Definition 3.1.1 (Equivalent Characterizations of a Square Orthogonal Matrix). A square matrix $U \in \mathbb{R}^{n \times n}$ is said to be **orthogonal** if it satisfies any of the following mutually equivalent conditions:

1. $U^T U = I_n$, meaning its inverse coincides with its transpose ($U^{-1} = U^T$).

2. $UU^T = I_n$.

3. The set of columns of $U = [\mathbf{u}_1, \dots, \mathbf{u}_n]$ forms an orthonormal basis of $\mathbb{R}^n$:

$$\langle \mathbf{u}_i, \mathbf{u}_j \rangle = \mathbf{u}_i^T \mathbf{u}_j = \delta_{ij} = \begin{cases} 1 & \text{if } i = j, \\ 0 & \text{if } i \neq j. \end{cases} \quad (3.1)$$

4. The set of rows of U forms an orthonormal basis of $\mathbb{R}^n$.

Historical Terminological Clarification

In mathematical literature, the term *orthogonal matrix* is standard, although the columns (and rows) of such a matrix are not merely pairwise orthogonal, but also **normalized to unit length**, forming an explicitly **orthonormal** system.

The key geometric property of an orthogonal transformation is the preservation of Euclidean distances and norms, formalized as isometry.

Property 3.1.1 (Euclidean Isometry). *Let $U \in \mathbb{R}^{n \times n}$ be an orthogonal matrix. For every vector $\mathbf{x} \in \mathbb{R}^n$, the linear map preserves the Euclidean norm exactly:*

$$\|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2. \quad (3.2)$$

Proof. Evaluating the squared Euclidean norm of $U\mathbf{x}$ and applying the identity $(U\mathbf{x})^T = \mathbf{x}^T U^T$:

$$\|U\mathbf{x}\|_2^2 = (U\mathbf{x})^T (U\mathbf{x}) = \mathbf{x}^T (U^T U) \mathbf{x} = \mathbf{x}^T I_n \mathbf{x} = \mathbf{x}^T \mathbf{x} = \|\mathbf{x}\|_2^2. \quad (3.3)$$

Taking the principal square root yields $\|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2$. □

3.2 Rectangular Matrices with Orthonormal Columns: $V^T V$ vs VV^T

In computational settings (such as dimensionality reduction, principal component analysis, and thin singular value decomposition), one routinely works with rectangular matrices whose columns form an orthonormal system.

3.2.1 Definition and Dimensional Constraint

Definition 3.2.1 (Rectangular Matrix with Orthonormal Columns). A matrix $V \in \mathbb{R}^{n \times k}$ with $k < n$ (a tall-skinny matrix) has orthonormal columns if:

$$V = \begin{bmatrix} | & | & \dots & | \\ \mathbf{v}_1 & \mathbf{v}_2 & \dots & \mathbf{v}_k \\ | & | & \dots & | \end{bmatrix}, \quad \mathbf{v}_i \in \mathbb{R}^n, \quad (3.4)$$

with column vectors satisfying the orthonormality condition:

$$\langle \mathbf{v}_i, \mathbf{v}_j \rangle = \mathbf{v}_i^T \mathbf{v}_j = \delta_{ij} \quad \forall i, j \in \{1, \dots, k\}. \quad (3.5)$$

Dimensional Constraint: Why Must $k \leq n$?

In an n -dimensional vector space such as $\mathbb{R}^n$, any subspace has dimension at most n . Because any finite set of orthonormal vectors is inherently linearly independent, there cannot exist more than n orthonormal vectors in $\mathbb{R}^n$. Consequently, a matrix with orthonormal columns can possess at most as many columns as it has rows ($k \leq n$).

3.2.2 The Inner Product $V^T V$ (Column Gramian)

Consider the product of the transpose $V^T \in \mathbb{R}^{k \times n}$ and $V \in \mathbb{R}^{n \times k}$. The result is a square matrix of order k :

$$V^T V \in \mathbb{R}^{k \times k}. \quad (3.6)$$

Expanding the block multiplication, representing V^T by rows and V by columns:

$$V^T V = \begin{bmatrix} - & \mathbf{v}_1^T & - \\ - & \mathbf{v}_2^T & - \\ & \vdots & \\ - & \mathbf{v}_k^T & - \end{bmatrix} \begin{bmatrix} | & | & \dots & | \\ \mathbf{v}_1 & \mathbf{v}_2 & \dots & \mathbf{v}_k \\ | & | & \dots & | \end{bmatrix} = \begin{bmatrix} \mathbf{v}_1^T \mathbf{v}_1 & \mathbf{v}_1^T \mathbf{v}_2 & \dots & \mathbf{v}_1^T \mathbf{v}_k \\ \mathbf{v}_2^T \mathbf{v}_1 & \mathbf{v}_2^T \mathbf{v}_2 & \dots & \mathbf{v}_2^T \mathbf{v}_k \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{v}_k^T \mathbf{v}_1 & \mathbf{v}_k^T \mathbf{v}_2 & \dots & \mathbf{v}_k^T \mathbf{v}_k \end{bmatrix}. \quad (3.7)$$

Applying orthonormality $\mathbf{v}_i^T \mathbf{v}_j = \delta_{ij}$, this matrix evaluates precisely to the $k \times k$ identity matrix:

$$V^T V = I_k. \quad (3.8)$$

3.2.3 The Reversed Product VV^T (Orthogonal Projector)

Reversing the order of multiplication yields a square matrix of substantially larger size, $n \times n$:

$$VV^T \in \mathbb{R}^{n \times n}. \quad (3.9)$$

Theorem 3.2.1 (Non-Identity of VV^T). *Let $V \in \mathbb{R}^{n \times k}$ have orthonormal columns with $k < n$. Then:*

$$VV^T \neq I_n. \quad (3.10)$$

Proof. The proof relies on the rank inequality for matrix products:

1. The rank of a matrix product is bounded by the rank of its factors:

$$\text{rank}(VV^T) \leq \min(\text{rank}(V), \text{rank}(V^T)). \quad (3.11)$$

2. Since the k columns of V are orthonormal and hence linearly independent, $\text{rank}(V) = \text{rank}(V^T) = k$.
3. Therefore, $\text{rank}(VV^T) \leq k$. In fact, because $V^TV = I_k$, we have exactly $\text{rank}(VV^T) = k$.
4. On the other hand, the identity matrix $I_n \in \mathbb{R}^{n \times n}$ has full rank n .
5. Because $k < n$ by assumption:

$$\text{rank}(VV^T) = k < n = \text{rank}(I_n) \implies VV^T \neq I_n. \quad (3.12)$$

The matrix VV^T is singular and possesses a non-trivial null space of dimension $n - k > 0$. $\square$

Outer Product Expansion and Geometric Interpretation

By block partition, the product VV^T can be expressed as a sum of outer products (dyads):

$$VV^T = \sum_{i=1}^k \mathbf{v}_i \mathbf{v}_i^T = \mathbf{v}_1 \mathbf{v}_1^T + \mathbf{v}_2 \mathbf{v}_2^T + \cdots + \mathbf{v}_k \mathbf{v}_k^T. \quad (3.13)$$

Each term $\mathbf{v}_i \mathbf{v}_i^T$ is a symmetric rank-1 matrix of size $n \times n$. The operator $P = VV^T$:

- Is idempotent: $P^2 = (VV^T)(VV^T) = V(V^TV)V^T = VI_kV^T = VV^T = P$.

- Is symmetric: $P^T = (VV^T)^T = (V^T)^T V^T = VV^T = P$.

Hence, $P = VV^T$ is the **orthogonal projector** from $\mathbb{R}^n$ onto the k -dimensional subspace:

$$\mathcal{S} = \text{span}\{\mathbf{v}_1, \dots, \mathbf{v}_k\}. \quad (3.14)$$




Figure 3.1: Dimension and rank comparison: the product VV^T yields an $n \times n$ orthogonal projector with strictly deficient rank $k < n$.

3.3 Partial Isometry for Rectangular Orthonormal Matrices

In the rectangular setting ($V \in \mathbb{R}^{n \times k}$ with $k < n$), the isometric property applies asymmetrically depending on the domain of the input vector.

3.3.1 Question (a): Norm Preservation for Vectors in $\mathbb{R}^k$

Proposition 3.3.1 (Right Isometry). *Let $V \in \mathbb{R}^{n \times k}$ have orthonormal columns with $k < n$. For any vector $\mathbf{x} \in \mathbb{R}^k$:*

$$\|V\mathbf{x}\|_2 = \|\mathbf{x}\|_2. \quad (3.15)$$

Proof. Compute the squared Euclidean norm of $V\mathbf{x} \in \mathbb{R}^n$:

$$\|V\mathbf{x}\|_2^2 = (V\mathbf{x})^T(V\mathbf{x}) = \mathbf{x}^T(V^T V)\mathbf{x}. \quad (3.16)$$

Since the columns of V are orthonormal, $V^T V = I_k$. Substituting:

$$\|V\mathbf{x}\|_2^2 = \mathbf{x}^T I_k \mathbf{x} = \mathbf{x}^T \mathbf{x} = \|\mathbf{x}\|_2^2. \quad (3.17)$$

Taking the positive square root yields $\|V\mathbf{x}\|_2 = \|\mathbf{x}\|_2$ for all $\mathbf{x} \in \mathbb{R}^k$. $\square$

Geometric Interpretation

The mapping $V : \mathbb{R}^k \rightarrow \mathbb{R}^n$ is an **isometric embedding**: it embeds the lower-dimensional space $\mathbb{R}^k$ into a k -dimensional subspace of $\mathbb{R}^n$ while strictly preserving all vector lengths and angles.

3.3.2 Question (b): Norm Preservation for Vectors in $\mathbb{R}^n$ under V^T

Proposition 3.3.2 (Failure of Left Isometry). *Let $V \in \mathbb{R}^{n \times k}$ have orthonormal columns with $k < n$. The transposed operator $V^T : \mathbb{R}^n \rightarrow \mathbb{R}^k$ does **not** preserve the Euclidean norm across all of $\mathbb{R}^n$:*

$$\exists \mathbf{y} \in \mathbb{R}^n \quad \text{such that} \quad \|V^T \mathbf{y}\|_2 \neq \|\mathbf{y}\|_2. \quad (3.18)$$

Proof. Evaluate the squared Euclidean norm of $V^T \mathbf{y}$:

$$\|V^T \mathbf{y}\|_2^2 = (V^T \mathbf{y})^T (V^T \mathbf{y}) = \mathbf{y}^T (V V^T) \mathbf{y}. \quad (3.19)$$

Because $V V^T \neq I_n$, this expression does not collapse identically to $\|\mathbf{y}\|_2^2$.

Null-Space Counterexample: Consider the matrix $V^T \in \mathbb{R}^{k \times n}$. By the Rank-Nullity Theorem:

$$\dim(\ker(V^T)) = n - \text{rank}(V^T) = n - k. \quad (3.20)$$

Since $k < n$, the null-space dimension is strictly positive ($n - k > 0$). There exists a non-zero vector $\tilde{\mathbf{y}} \in \mathbb{R}^n \setminus \{\mathbf{0}\}$ such that:

$$\tilde{\mathbf{y}} \in \ker(V^T) \implies V^T \tilde{\mathbf{y}} = \mathbf{0} \in \mathbb{R}^k. \quad (3.21)$$

Computing norms:

$$\|V^T \tilde{\mathbf{y}}\|_2 = \|\mathbf{0}\|_2 = 0 \neq \|\tilde{\mathbf{y}}\|_2. \quad (3.22)$$

Thus, norm shrinkage is maximal for any vector orthogonal to the subspace spanned by the columns of V . $\square$

3.4 Eigenvalues and Eigenvectors: Foundations and Step-by-Step Example

3.4.1 Fundamental Definitions

Let $A \in \mathbb{R}^{n \times n}$ be a **square** matrix.

Definition 3.4.1 (Eigenvalue, Eigenvector, and Eigenpair). A scalar $\lambda \in \mathbb{R}$ (or in $\mathbb{C}$) is called an **eigenvalue** of A if there exists a **non-zero** vector $\mathbf{v} \in \mathbb{R}^n \setminus \{\mathbf{0}\}$ (or $\mathbb{C}^n \setminus \{\mathbf{0}\}$) such that:

$$A\mathbf{v} = \lambda\mathbf{v}. \quad (3.23)$$

The vector $\mathbf{v}$ is termed an **eigenvector** associated with λ . The pair $(\lambda, \mathbf{v})$ is called an **eigenpair**.

The eigenvalue equation $A\mathbf{v} = \lambda\mathbf{v}$ can be recast as a homogeneous linear system:

$$(A - \lambda I_n)\mathbf{v} = \mathbf{0}. \quad (3.24)$$

For non-trivial solutions ($\mathbf{v} \neq \mathbf{0}$) to exist, the coefficient matrix $(A - \lambda I_n)$ must be singular, meaning its determinant vanishes:

$$p(\lambda) = \det(A - \lambda I_n) = 0. \quad (3.25)$$

The polynomial $p(\lambda)$ is the **characteristic polynomial** of A , of degree n . The eigenvalues of A are precisely the roots of this characteristic polynomial.

3.4.2 Step-by-Step Worked Example

Example 3.4.1 (Computing Eigenvalues and Eigenvectors for a 2×2 Matrix). Consider the real symmetric matrix:

$$A = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \in \mathbb{R}^{2 \times 2}. \quad (3.26)$$

Step 1: Characteristic Polynomial and Eigenvalues.

$$p(\lambda) = \det(A - \lambda I_2) = \det \begin{bmatrix} 2 - \lambda & 1 \\ 1 & 2 - \lambda \end{bmatrix} = (2 - \lambda)^2 - 1. \quad (3.27)$$

Using the difference of squares factorization $a^2 - b^2 = (a - b)(a + b)$:

$$p(\lambda) = ((2 - \lambda) - 1)((2 - \lambda) + 1) = (1 - \lambda)(3 - \lambda) = 0. \quad (3.28)$$

Setting $p(\lambda) = 0$ yields the eigenvalues:

$$\lambda_1 = 3, \quad \lambda_2 = 1. \quad (3.29)$$

Step 2: Eigenvector for $\lambda_1 = 3$. We solve $(A - 3I_2)\mathbf{v}_1 = \mathbf{0}$:

$$\begin{bmatrix} 2 - 3 & 1 \\ 1 & 2 - 3 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -1 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix} \implies -x + y = 0 \iff y = x. \quad (3.30)$$

Choosing $x = 1$, we obtain the eigenvector:

$$\mathbf{v}_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}. \quad (3.31)$$

Step 3: Eigenvector for $\lambda_2 = 1$. We solve $(A - I_2)\mathbf{v}_2 = \mathbf{0}$:

$$\begin{bmatrix} 2-1 & 1 \\ 1 & 2-1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix} \implies x+y=0 \iff y=-x. \quad (3.32)$$

Choosing $x = 1$, we find the eigenvector:

$$\mathbf{v}_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix}. \quad (3.33)$$

Notice that $\mathbf{v}_1^T \mathbf{v}_2 = 1(1) + 1(-1) = 0$: the two eigenvectors are mutually orthogonal.

3.5 Diagonalization and Dyadic Expansion

3.5.1 Diagonalizable Matrices

Definition 3.5.1 (Diagonalizability). A matrix $A \in \mathbb{R}^{n \times n}$ is called **diagonalizable** if it possesses n linearly independent eigenvectors, that is, if there exists a basis of $\mathbb{R}^n$ consisting of eigenvectors of A .

In this case, assembling the n eigenvectors into the modal matrix $V = [\mathbf{v}_1, \dots, \mathbf{v}_n] \in \mathbb{R}^{n \times n}$ and the eigenvalues along the diagonal matrix $D = \text{diag}(\lambda_1, \dots, \lambda_n)$:

$$AV = A[\mathbf{v}_1, \dots, \mathbf{v}_n] = [\lambda_1 \mathbf{v}_1, \dots, \lambda_n \mathbf{v}_n] = [\mathbf{v}_1, \dots, \mathbf{v}_n] \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} = VD. \quad (3.34)$$

Because the eigenvectors are linearly independent, V is invertible. Multiplying on the right by V^{-1} :

$$A = VDV^{-1}. \quad (3.35)$$

3.5.2 Dyadic (Outer Product) Expansion

Let $W = V^{-1} \in \mathbb{R}^{n \times n}$, partitioned by rows:

$$W = V^{-1} = \begin{bmatrix} - & \mathbf{w}_1^T & - \\ - & \mathbf{w}_2^T & - \\ & \vdots & \\ - & \mathbf{w}_n^T & - \end{bmatrix}. \quad (3.36)$$

Because $WV = I_n$, the rows of W and columns of V satisfy biorthonormality: $\mathbf{w}_i^T \mathbf{v}_j = \delta_{ij}$. Expanding $A = VDW$:

$$\begin{aligned} A &= \begin{bmatrix} | & & | \\ \mathbf{v}_1 & \dots & \mathbf{v}_n \\ | & & | \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} \begin{bmatrix} - & \mathbf{w}_1^T & - \\ & \vdots & \\ - & \mathbf{w}_n^T & - \end{bmatrix} \\ &= \begin{bmatrix} | & & | \\ \lambda_1 \mathbf{v}_1 & \dots & \lambda_n \mathbf{v}_n \\ | & & | \end{bmatrix} \begin{bmatrix} - & \mathbf{w}_1^T & - \\ & \vdots & \\ - & \mathbf{w}_n^T & - \end{bmatrix}. \end{aligned} \quad (3.37)$$

Performing block column-by-row multiplication yields the dyadic expansion:

$$A = \sum_{i=1}^n \lambda_i \mathbf{v}_i \mathbf{w}_i^T = \lambda_1 \mathbf{v}_1 \mathbf{w}_1^T + \lambda_2 \mathbf{v}_2 \mathbf{w}_2^T + \dots + \lambda_n \mathbf{v}_n \mathbf{w}_n^T. \quad (3.38)$$

Significance of the Dyadic Expansion

The matrix A is represented as a weighted sum of n elementary rank-1 matrices, each formed by the outer product $\mathbf{v}_i \mathbf{w}_i^T$. The eigenvalues λ_i serve as modulation weights. This structure serves as the direct precursor to the Singular Value Decomposition (SVD) and the Eckart-Young-Mirsky low-rank approximation theorem.

3.5.3 Pathologies: When Diagonalization Fails

Not all matrices admit a decomposition $A = VDV^{-1}$ over the real numbers. Two distinct theoretical obstacles exist.

Case 1: Absence of Real Eigenvalues

Consider the 90° planar rotation matrix:

$$A = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}. \quad (3.39)$$

Its characteristic polynomial is:

$$p(\lambda) = \det \begin{bmatrix} -\lambda & -1 \\ 1 & -\lambda \end{bmatrix} = \lambda^2 + 1 = 0 \implies \lambda_{1,2} = \pm i \notin \mathbb{R}. \quad (3.40)$$

Although A has purely real entries, it **admits no real eigenvalues or real eigenvectors**. Geometrically, a non-trivial planar rotation does not preserve the direction of any real vector in $\mathbb{R}^2$.

Case 2: Defective Matrices (Geometric Multiplicity Strictly Less than Algebraic)

Consider the shear matrix:

$$A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}. \quad (3.41)$$

1. The characteristic polynomial is $p(\lambda) = (1 - \lambda)^2 = 0$. There is a single eigenvalue $\lambda = 1$ with **algebraic multiplicity** $m_a(1) = 2$.
2. We determine its **geometric multiplicity** $m_g(1) = \dim(\ker(A - I_2))$:

$$A - I_2 = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \implies \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix} \implies y = 0, \quad x \in \mathbb{R} \text{ free.} \quad (3.42)$$

The eigenspace is one-dimensional: $\ker(A - I_2) = \text{span} \left\{ \begin{bmatrix} 1 \\ 0 \end{bmatrix} \right\}$, so $m_g(1) = 1$.

Since $m_g(1) = 1 < 2 = m_a(1)$, it is impossible to find 2 linearly independent eigenvectors to form V . The matrix A is termed **defective** and cannot be diagonalized (it admits only a Jordan canonical form).

3.6 Guided Exercises on Eigenvalues (Q1 – Q4)

3.6.1 Exercise Q1: Eigenpairs of a Diagonal Matrix

Example 3.6.1 (Spectrum of a Diagonal Matrix). Given a generic diagonal matrix $D = \text{diag}(d_1, d_2, \dots, d_n) \in \mathbb{R}^{n \times n}$, we determine its eigenpairs.

Solution: Examine the action of D on the i -th standard basis vector $\mathbf{e}_i = [0, \dots, 1, \dots, 0]^T$:

$$D\mathbf{e}_i = \begin{bmatrix} d_1 & & & \\ & d_2 & & \\ & & \ddots & \\ & & & d_n \end{bmatrix} \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ \vdots \\ d_i \\ \vdots \\ 0 \end{bmatrix} = d_i \mathbf{e}_i. \quad (3.43)$$

Therefore:

- The eigenvalues are precisely the entries along the main diagonal: $\lambda_i = d_i$ for $i = 1, \dots, n$.

- The corresponding eigenvectors are the standard basis vectors: $\mathbf{v}_i = \mathbf{e}_i$.
- The canonical eigenpairs are $(d_1, \mathbf{e}_1), (d_2, \mathbf{e}_2), \dots, (d_n, \mathbf{e}_n)$.

Property 3.6.1 (Coinciding Eigenvalues and Eigenspace Structure). *What happens if two or more diagonal entries coincide, for example $d_1 = d_2 = \lambda$? In this case, **any linear combination** $\alpha\mathbf{e}_1 + \beta\mathbf{e}_2$ (with α, β not both zero) remains an eigenvector corresponding to λ .*

*General proof: Let $\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$ be two eigenvectors of an arbitrary matrix A associated with the **same** eigenvalue λ ($A\mathbf{v} = \lambda\mathbf{v}$ and $A\mathbf{w} = \lambda\mathbf{w}$). For any scalars $\alpha, \beta \in \mathbb{R}$:*

$$A(\alpha\mathbf{v} + \beta\mathbf{w}) = \alpha(A\mathbf{v}) + \beta(A\mathbf{w}) = \alpha(\lambda\mathbf{v}) + \beta(\lambda\mathbf{w}) = \lambda(\alpha\mathbf{v} + \beta\mathbf{w}). \quad (3.44)$$

*Hence, the set of eigenvectors associated with λ , together with the zero vector, forms a closed vector subspace known as the **eigenspace**:*

$$\mathcal{E}(\lambda) = \ker(A - \lambda I_n). \quad (3.45)$$

3.6.2 Exercise Q2: Eigenpairs of the Identity Matrix

Example 3.6.2 (Spectrum of the Identity Matrix). Determine the eigenvalues and eigenvectors of the identity matrix $I_n \in \mathbb{R}^{n \times n}$.

Solution: By definition of the identity operator, for every non-zero vector $\mathbf{x} \in \mathbb{R}^n \setminus \{\mathbf{0}\}$:

$$I_n \mathbf{x} = 1 \cdot \mathbf{x}. \quad (3.46)$$

- **Eigenvalue:** $\lambda = 1$ with algebraic and geometric multiplicity n ($m_a = m_g = n$).
- **Eigenvectors:** every non-zero vector in $\mathbb{R}^n$.
- The entire space $\mathbb{R}^n$ constitutes the eigenspace: $\mathcal{E}(1) = \mathbb{R}^n$.

3.6.3 Exercise Q3: Sum and Product of Matrices Sharing an Eigenvalue

Example 3.6.3 (Eigenvalues of Matrix Sum and Product). Let $\lambda \in \mathbb{R}$ be a common eigenvalue of $A \in \mathbb{R}^{n \times n}$ and $B \in \mathbb{R}^{n \times n}$.

1. Is 2λ necessarily an eigenvalue of $A + B$?
2. Is λ^2 necessarily an eigenvalue of AB ?

Analytical Answer: NO to both questions. For eigenvalue addition or multiplication to hold, A and B must share the **same eigenvector** associated with λ . If their eigenvectors are misaligned, the property fails.

Counterexample for Summation: Consider:

$$A = \begin{bmatrix} 5 & 0 \\ 0 & 1 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & 0 \\ 0 & 5 \end{bmatrix}. \quad (3.47)$$

Both matrices possess the eigenvalue set $\{1, 5\}$ (in particular, sharing $\lambda = 1$ and $\lambda = 5$). Compute the sum:

$$A + B = \begin{bmatrix} 6 & 0 \\ 0 & 6 \end{bmatrix} = 6I_2. \quad (3.48)$$

The eigenvalues of $A + B$ are both equal to 6. However, doubling the original eigenvalues gives:

$$2 \cdot 1 = 2 \neq 6, \quad 2 \cdot 5 = 10 \neq 6. \quad (3.49)$$

Thus, 2λ does not belong to the spectrum of the sum matrix.

Counterexample for Multiplication: Consider:

$$A = \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}, \quad B = \begin{bmatrix} 10 & 0 \\ 0 & 2 \end{bmatrix}. \quad (3.50)$$

Both matrices share the eigenvalue $\lambda = 2$. Their product evaluates to:

$$AB = \begin{bmatrix} 20 & 0 \\ 0 & 6 \end{bmatrix}. \quad (3.51)$$

The eigenvalues of AB are $\{20, 6\}$, neither of which matches $\lambda^2 = 2^2 = 4$.

3.6.4 Exercise Q4: Eigenpairs of a Matrix Power

Theorem 3.6.1 (Spectrum of Matrix Power A^k). *Let $A \in \mathbb{R}^{n \times n}$ and let $(\lambda, \mathbf{v})$ be an eigenpair of A ($A\mathbf{v} = \lambda\mathbf{v}$). Then for every positive integer exponent $k \in \mathbb{N}$:*

$$(\lambda^k, \mathbf{v}) \text{ is an eigenpair of } A^k, \quad \text{that is, } A^k \mathbf{v} = \lambda^k \mathbf{v}. \quad (3.52)$$

Proof. The proof proceeds by induction on k :

1. **Base Case** ($k = 1$): $A^1 \mathbf{v} = A\mathbf{v} = \lambda\mathbf{v} = \lambda^1 \mathbf{v}$.

2. **Inductive Step:** Assume $A^k \mathbf{v} = \lambda^k \mathbf{v}$. Multiplying on the left by A :

$$A^{k+1} \mathbf{v} = A(A^k \mathbf{v}) = A(\lambda^k \mathbf{v}) = \lambda^k (A \mathbf{v}) = \lambda^k (\lambda \mathbf{v}) = \lambda^{k+1} \mathbf{v}. \quad (3.53)$$

The statement holds for all $k \geq 1$. $\square$

Corollary 3.6.1 (Power Factorization for Diagonalizable Matrices). *If $A = VDV^{-1}$ is diagonalizable, evaluating powers telescopically yields:*

$$\begin{aligned} A^k &= (VDV^{-1})(VDV^{-1}) \cdots (VDV^{-1}) \\ &= VD(V^{-1}V)D \cdots (V^{-1}V)DV^{-1} = VD^kV^{-1}. \end{aligned} \quad (3.54)$$

Since $D = \text{diag}(\lambda_1, \dots, \lambda_n)$ is diagonal, its k -th power is simply:

$$D^k = \text{diag}(\lambda_1^k, \lambda_2^k, \dots, \lambda_n^k). \quad (3.55)$$

Warning on Eigenvector Invariance

Take careful note that when raising a matrix to a power, **only the eigenvalues** are raised to the power k , while the **eigenvectors remain strictly unchanged**. The notation $\mathbf{v}^k$ has no meaning in vector algebra.

3.7 Symmetric Matrices and the Spectral Theorem

Symmetric matrices represent the most important class of matrices in numerical linear algebra, optimization, and machine learning.

Definition 3.7.1 (Symmetric Matrix). A square matrix $A \in \mathbb{R}^{n \times n}$ is said to be **symmetric** if it coincides with its transpose:

$$A^T = A \iff A_{ij} = A_{ji} \quad \forall i, j \in \{1, \dots, n\}. \quad (3.56)$$

Theorem 3.7.1 (Spectral Theorem for Real Symmetric Matrices). *Let $A \in \mathbb{R}^{n \times n}$ be a real symmetric matrix ($A^T = A$). Then:*

1. All n eigenvalues of A are **real numbers**: $\lambda_1, \lambda_2, \dots, \lambda_n \in \mathbb{R}$.
2. Eigenvectors corresponding to distinct eigenvalues are mutually **orthogonal**.

3. The matrix A is **orthogonally diagonalizable**: there exist an orthogonal matrix $U \in \mathbb{R}^{n \times n}$ ($U^T U = U U^T = I_n$) and a real diagonal matrix $D \in \mathbb{R}^{n \times n}$ such that:

$$A = U D U^T. \quad (3.57)$$

Equivalently, there exists an **orthonormal basis of $\mathbb{R}^n$ consisting entirely of eigenvectors of A** .

Distinct Eigenvalues. Let $(\lambda_1, \mathbf{u}_1)$ and $(\lambda_2, \mathbf{u}_2)$ be two eigenpairs of A with $\lambda_1 \neq \lambda_2$. Consider the scalar quantity $\mathbf{u}_1^T A \mathbf{u}_2$:

1. Acting with A to the right:

$$\mathbf{u}_1^T (A \mathbf{u}_2) = \mathbf{u}_1^T (\lambda_2 \mathbf{u}_2) = \lambda_2 (\mathbf{u}_1^T \mathbf{u}_2). \quad (3.58)$$

2. Using symmetry $A = A^T$ and acting to the left:

$$\mathbf{u}_1^T A \mathbf{u}_2 = (A^T \mathbf{u}_1)^T \mathbf{u}_2 = (A \mathbf{u}_1)^T \mathbf{u}_2 = (\lambda_1 \mathbf{u}_1)^T \mathbf{u}_2 = \lambda_1 (\mathbf{u}_1^T \mathbf{u}_2). \quad (3.59)$$

Equating the two expressions:

$$\lambda_1 (\mathbf{u}_1^T \mathbf{u}_2) = \lambda_2 (\mathbf{u}_1^T \mathbf{u}_2) \iff (\lambda_1 - \lambda_2) (\mathbf{u}_1^T \mathbf{u}_2) = 0. \quad (3.60)$$

Since $\lambda_1 \neq \lambda_2$ by assumption, it necessarily follows that:

$$\mathbf{u}_1^T \mathbf{u}_2 = 0, \quad (3.61)$$

proving that the eigenvectors are orthogonal. $\square$

3.8 Quadratic Forms and the Rayleigh Quotient Inequality

3.8.1 Definition of a Quadratic Form

Definition 3.8.1 (Quadratic Form). Let $A \in \mathbb{R}^{n \times n}$ be a real symmetric matrix. The function $q : \mathbb{R}^n \rightarrow \mathbb{R}$ defined by:

$$q(\mathbf{x}) = \mathbf{x}^T A \mathbf{x} = \sum_{i=1}^n \sum_{j=1}^n A_{ij} x_i x_j \quad (3.62)$$

is called the **quadratic form** associated with matrix A .

3.8.2 Spectral Bounds Theorem (Rayleigh Inequality)

Theorem 3.8.1 (Spectral Bounds on Quadratic Forms). *Let $A \in \mathbb{R}^{n \times n}$ be real symmetric, with eigenvalues arranged in descending order:*

$$\lambda_{\max} = \lambda_1 \geq \lambda_2 \geq \cdots \geq \lambda_n = \lambda_{\min}. \quad (3.63)$$

Then for every vector $\mathbf{x} \in \mathbb{R}^n$, the double inequality holds:

$$\lambda_{\min} \|\mathbf{x}\|_2^2 \leq \mathbf{x}^T A \mathbf{x} \leq \lambda_{\max} \|\mathbf{x}\|_2^2. \quad (3.64)$$

Proof. The proof proceeds systematically:

1. **Spectral Factorization:** By the Spectral Theorem 3.7.1, we factorize $A = U D U^T$, where $U \in \mathbb{R}^{n \times n}$ is orthogonal and $D = \text{diag}(\lambda_1, \dots, \lambda_n)$.

2. **Substitution into the Quadratic Form:**

$$\mathbf{x}^T A \mathbf{x} = \mathbf{x}^T (U D U^T) \mathbf{x} = (U^T \mathbf{x})^T D (U^T \mathbf{x}). \quad (3.65)$$

3. **Orthogonal Change of Coordinates:** Define $\mathbf{y} \in \mathbb{R}^n$ via:

$$\mathbf{y} = U^T \mathbf{x} \iff \mathbf{x} = U \mathbf{y}. \quad (3.66)$$

Because U is orthogonal, the map preserves Euclidean norm (Property 3.1.1):

$$\|\mathbf{y}\|_2^2 = \|U^T \mathbf{x}\|_2^2 = \|\mathbf{x}\|_2^2 \implies \sum_{i=1}^n y_i^2 = \|\mathbf{x}\|_2^2. \quad (3.67)$$

4. **Component Expansion:**

$$\mathbf{x}^T A \mathbf{x} = \mathbf{y}^T D \mathbf{y} = \sum_{i=1}^n \lambda_i y_i^2 = \lambda_1 y_1^2 + \lambda_2 y_2^2 + \cdots + \lambda_n y_n^2. \quad (3.68)$$

5. **Upper Bound:** Since $\lambda_i \leq \lambda_{\max}$ for all $i \in \{1, \dots, n\}$ and $y_i^2 \geq 0$:

$$\mathbf{x}^T A \mathbf{x} = \sum_{i=1}^n \lambda_i y_i^2 \leq \sum_{i=1}^n \lambda_{\max} y_i^2 = \lambda_{\max} \sum_{i=1}^n y_i^2 = \lambda_{\max} \|\mathbf{y}\|_2^2 = \lambda_{\max} \|\mathbf{x}\|_2^2. \quad (3.69)$$

6. **Lower Bound:** Analogously, because $\lambda_i \geq \lambda_{\min}$ for all i :

$$\mathbf{x}^T A \mathbf{x} = \sum_{i=1}^n \lambda_i y_i^2 \geq \sum_{i=1}^n \lambda_{\min} y_i^2 = \lambda_{\min} \sum_{i=1}^n y_i^2 = \lambda_{\min} \|\mathbf{y}\|_2^2 = \lambda_{\min} \|\mathbf{x}\|_2^2. \quad (3.70)$$

Combining both bounds completes the proof for all $\mathbf{x} \in \mathbb{R}^n$. $\square$

3.8.3 Equality Conditions and the Rayleigh Quotient

Property 3.8.1 (Equality Conditions and Extremal Vectors).

1. The upper equality $\mathbf{x}^T A \mathbf{x} = \lambda_{\max} \|\mathbf{x}\|_2^2$ holds **if and only if** $\mathbf{x}$ is an eigenvector of A corresponding to $\lambda_{\max}$.
2. The lower equality $\mathbf{x}^T A \mathbf{x} = \lambda_{\min} \|\mathbf{x}\|_2^2$ holds **if and only if** $\mathbf{x}$ is an eigenvector of A corresponding to $\lambda_{\min}$.

Proof. Examining the upper bound difference:

$$\lambda_{\max} \sum_{i=1}^n y_i^2 - \sum_{i=1}^n \lambda_i y_i^2 = \sum_{i=1}^n (\lambda_{\max} - \lambda_i) y_i^2 = 0. \quad (3.71)$$

Since each term $(\lambda_{\max} - \lambda_i) y_i^2 \geq 0$, the sum vanishes if and only if $y_i = 0$ whenever $\lambda_i < \lambda_{\max}$. Thus $\mathbf{y}$ is supported only on coordinates corresponding to $\lambda_{\max}$. In original coordinates $\mathbf{x} = U \mathbf{y}$, this precisely states that $\mathbf{x} \in \mathcal{E}(\lambda_{\max})$. $\square$

Definition 3.8.2 (Rayleigh Quotient). For any non-zero vector $\mathbf{x} \in \mathbb{R}^n \setminus \{\mathbf{0}\}$, the **Rayleigh quotient** of a symmetric matrix A is defined as:

$$R_A(\mathbf{x}) = \frac{\mathbf{x}^T A \mathbf{x}}{\mathbf{x}^T \mathbf{x}} = \frac{\mathbf{x}^T A \mathbf{x}}{\|\mathbf{x}\|_2^2}. \quad (3.72)$$

Theorem 3.8.1 yields the variational characterization of extremal eigenvalues:

$$\lambda_{\min} = \min_{\mathbf{x} \neq \mathbf{0}} R_A(\mathbf{x}), \quad \lambda_{\max} = \max_{\mathbf{x} \neq \mathbf{0}} R_A(\mathbf{x}). \quad (3.73)$$

3.9 Symmetric Positive Definite and Semidefinite Matrices

3.9.1 Axiomatic Definitions

Let $A \in \mathbb{R}^{n \times n}$ be a real symmetric matrix ($A^T = A$).

Definition 3.9.1 (Symmetric Positive Definite Matrix – SPD). The matrix A is called **symmetric positive definite** (SPD) if its quadratic form is strictly positive for every non-zero vector:

$$\mathbf{v}^T A \mathbf{v} > 0 \quad \forall \mathbf{v} \in \mathbb{R}^n \setminus \{\mathbf{0}\}. \quad (3.74)$$

Definition 3.9.2 (Symmetric Positive Semi-Definite Matrix – SPSD). The matrix A is called **symmetric positive semi-definite** (SPSD) if its quadratic form is non-negative for every vector:

$$\mathbf{v}^T A \mathbf{v} \geq 0 \quad \forall \mathbf{v} \in \mathbb{R}^n. \quad (3.75)$$

3.9.2 Equivalent Spectral Characterization

Theorem 3.9.1 (Spectral Criterion for SPD and SPSD Matrices). *Let $A \in \mathbb{R}^{n \times n}$ be real symmetric with eigenvalues $\lambda_1, \dots, \lambda_n$. Then:*

1. A is **SPD** $\iff \lambda_i > 0$ for all $i \in \{1, \dots, n\}$ (all eigenvalues are strictly positive).
2. A is **SPSD** $\iff \lambda_i \geq 0$ for all $i \in \{1, \dots, n\}$ (all eigenvalues are non-negative).

Proof. Directly from the Rayleigh quotient inequality:

- If all $\lambda_i > 0$, then $\lambda_{\min} > 0$. By $\mathbf{v}^T A \mathbf{v} \geq \lambda_{\min} \|\mathbf{v}\|_2^2$, for any $\mathbf{v} \neq \mathbf{0}$ we have $\mathbf{v}^T A \mathbf{v} > 0$, so A is SPD.
- Conversely, if A is SPD, setting $\mathbf{v} = \mathbf{u}_i$ (unit eigenvector for λ_i):

$$\mathbf{u}_i^T A \mathbf{u}_i = \lambda_i \|\mathbf{u}_i\|_2^2 = \lambda_i > 0. \quad (3.76)$$

- Replacing strict inequalities with weak ones establishes the SPSD case. □

3.9.3 Matrix Class Hierarchy and Nested Inclusions




Figure 3.2: Structural hierarchy and progressive refinement of matrix classes in linear algebra.

The inclusion is strict:

$$\text{SPD} \subset \text{SPSD} \subset \text{Symmetric}. \quad (3.77)$$

An SPSD matrix belongs to the SPD class if and only if it is non-singular, meaning it has no zero eigenvalue ($\lambda_{\min} > 0 \iff \det(A) = \prod_{i=1}^n \lambda_i > 0$).

3.9.4 Illustrative Examples in $\mathbb{R}^{2 \times 2}$

Example 3.9.1 (Verification of SPD and SPSD Matrices).

1. **SPD Examples:**

- $A_1 = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$: symmetric diagonal with eigenvalues $\lambda_1 = 1 > 0, \lambda_2 = 2 > 0 \implies$ **SPD**.
- $A_2 = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$: symmetric with eigenvalues calculated in Example 3.4.1 as $\lambda_1 = 3 > 0, \lambda_2 = 1 > 0 \implies$ **SPD**.

2. **SPSD but Non-SPD Examples:**

- $A_3 = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$: symmetric diagonal with eigenvalues $\lambda_1 = 1 \geq 0, \lambda_2 = 0 \geq 0 \implies$ **SPSD**. It is not SPD because $\det(A_3) = 0$.
- $A_4 = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$: rows are identical, so $\text{rank}(A_4) = 1 < 2$, implying $\lambda_1 = 0$. Because the trace equals the sum of eigenvalues, $\text{Tr}(A_4) = 1 + 1 = 2 \implies \lambda_1 + \lambda_2 = 2 \implies \lambda_2 = 2 \geq 0$. The eigenvalues are $\{2, 0\} \implies$ **SPSD**.

3. **Counterexample of an Indefinite Matrix:** Consider $A_5 = \begin{bmatrix} 1 & 2 \\ 2 & 0 \end{bmatrix}$.

Although its diagonal entries are non-negative, the determinant evaluates to:

$$\det(A_5) = 1(0) - 2(2) = -4 < 0. \quad (3.78)$$

Because $\det(A_5) = \lambda_1 \lambda_2 = -4 < 0$, the eigenvalues have opposite signs ($\lambda_1 > 0$ and $\lambda_2 < 0$). The matrix is **indefinite** (neither SPD nor SPSD).

3.10 Natural Generation of SPSD Matrices: $B^T B$ and BB^T

In scientific computing and statistics, positive semidefinite matrices are rarely assembled manually; rather, they arise naturally from multiplying a rectangular matrix by its transpose.

Proposition 3.10.1 (Positive Semidefiniteness of Gram Matrices). *Let $B \in \mathbb{R}^{n \times m}$ be an arbitrary matrix (square or rectangular). Then:*

1. The square matrix $B^T B \in \mathbb{R}^{m \times m}$ is symmetric positive semi-definite (**SPSD**).
2. The square matrix $BB^T \in \mathbb{R}^{n \times n}$ is symmetric positive semi-definite (**SPSD**).

Proof. We establish symmetry and positive semidefiniteness:

1. Verification of Symmetry: Using the transpose of a matrix product:

$$(B^T B)^T = B^T (B^T)^T = B^T B, \quad (BB^T)^T = (B^T)^T B^T = BB^T. \quad (3.79)$$

Both matrices are formally symmetric.

2. Positive Semidefiniteness of $B^T B$: Let $\mathbf{v} \in \mathbb{R}^m$ be an arbitrary vector. Consider the quadratic form:

$$\mathbf{v}^T (B^T B) \mathbf{v} = (\mathbf{v}^T B^T) (B \mathbf{v}) = (B \mathbf{v})^T (B \mathbf{v}). \quad (3.80)$$

Recognizing the inner product $(B \mathbf{v})^T (B \mathbf{v})$ as the squared Euclidean norm of $\mathbf{w} = B \mathbf{v} \in \mathbb{R}^n$:

$$\mathbf{v}^T (B^T B) \mathbf{v} = \|B \mathbf{v}\|_2^2. \quad (3.81)$$

By non-negativity of vector norms, $\|B \mathbf{v}\|_2^2 \geq 0$ for all $\mathbf{v} \in \mathbb{R}^m$. Therefore:

$$\mathbf{v}^T (B^T B) \mathbf{v} \geq 0 \quad \forall \mathbf{v} \in \mathbb{R}^m \implies B^T B \text{ is SPSPD}. \quad (3.82)$$

3. Positive Semidefiniteness of BB^T : Let $\mathbf{u} \in \mathbb{R}^n$ be an arbitrary vector. In an identical manner:

$$\mathbf{u}^T (BB^T) \mathbf{u} = (\mathbf{u}^T B) (B^T \mathbf{u}) = (B^T \mathbf{u})^T (B^T \mathbf{u}) = \|B^T \mathbf{u}\|_2^2 \geq 0 \quad \forall \mathbf{u} \in \mathbb{R}^n. \quad (3.83)$$

Hence BB^T is also SPSPD. $\square$

Corollary 3.10.1 (Condition for Strict Positive Definiteness of $B^T B$). *The matrix $B^T B \in \mathbb{R}^{m \times m}$ is **strictly positive definite (SPD)** if and only if B has full column rank:*

$$B^T B \in \text{SPD} \iff \text{rank}(B) = m \iff \ker(B) = \{\mathbf{0}\}. \quad (3.84)$$

Proof. From the previous proof, $\mathbf{v}^T (B^T B) \mathbf{v} = \|B \mathbf{v}\|_2^2$. This quantity vanishes if and only if $\|B \mathbf{v}\|_2 = 0 \iff B \mathbf{v} = \mathbf{0} \iff \mathbf{v} \in \ker(B)$. When $\ker(B) = \{\mathbf{0}\}$, no non-zero vector can zero out the quadratic form, making $B^T B$ strictly positive definite. $\square$

Relevance to Normal Equations and SVD

This fundamental property underpins two cornerstones of computational mathematics:

1. **Linear Least Squares:** Solving $\min_{\mathbf{x}} \|\mathbf{Ax} - \mathbf{b}\|_2$ yields the normal equations:

$$A^T A \mathbf{x} = A^T \mathbf{b}. \quad (3.85)$$

When A has full column rank, $A^T A$ is SPD, guaranteeing unique solvability via Cholesky factorization.

2. **Singular Value Decomposition (SVD):** The singular values σ_i of an arbitrary matrix B are defined as the non-negative square roots of the eigenvalues of the SPSP matrix $B^T B$:

$$\sigma_i = \sqrt{\lambda_i(B^T B)} \geq 0. \quad (3.86)$$

3.11 Chapter Summary Formulary

Table 3.1 provides an organized summary of definitions, matrix structures, and algebraic properties covered in this chapter.

Concept / Matrix	Mathematical Formulation	Key Property / Spectrum
Orthonormal Columns ($k < n$)	$V \in \mathbb{R}^{n \times k}, \mathbf{v}_i^T \mathbf{v}_j = \delta_{ij}$	$V^T V = I_k$, while $V V^T \neq I_n$
Orthogonal Projector	$P = V V^T = \sum_{i=1}^k \mathbf{v}_i \mathbf{v}_i^T$	Symmetric ($P^T = P$), idempotent ($P^2 = P$), $\text{rank}(P) = k$
Partial Isometry	$\ V \mathbf{x}\ _2 = \ \mathbf{x}\ _2 \quad (\forall \mathbf{x} \in \mathbb{R}^k)$	Does not preserve norm on $\mathbb{R}^n$: $\ V^T \tilde{\mathbf{y}}\ _2 = 0$ for $\tilde{\mathbf{y}} \in \ker(V^T)$
Eigenpair	$A \mathbf{v} = \lambda \mathbf{v}, \quad \mathbf{v} \neq \mathbf{0}$	λ root of $\det(A - \lambda I_n) = 0$
Diagonalization	$A = V D V^{-1}$	Requires n linearly independent eigenvectors
Dyadic Expansion	$A = \sum_{i=1}^n \lambda_i \mathbf{v}_i \mathbf{v}_i^T$	Linear combination of n rank-1 matrices
Matrix Power	$A^k = V D^k V^{-1}$	Eigenpairs $(\lambda_i^k, \mathbf{v}_i)$: eigenvectors unchanged
Spectral Theorem	$A^T = A \implies A = U D U^T$	Real eigenvalues, orthonormal eigenvectors ($U^T U = I_n$)
Rayleigh Inequality	$\lambda_{\min} \ \mathbf{x}\ _2^2 \leq \mathbf{x}^T A \mathbf{x} \leq \lambda_{\max} \ \mathbf{x}\ _2^2$	Bounds attained on corresponding eigenspaces
SPD Matrix	$A^T = A, \mathbf{x}^T A \mathbf{x} > 0 \quad (\forall \mathbf{x} \neq \mathbf{0})$	All $\lambda_i > 0$; invertible with $\det(A) > 0$
SPSD Matrix	$A^T = A, \mathbf{x}^T A \mathbf{x} \geq 0 \quad (\forall \mathbf{x})$	All $\lambda_i \geq 0$; may have zero eigenvalues
Gramians $B^T B$ and $B B^T$	$\mathbf{v}^T (B^T B) \mathbf{v} = \ B \mathbf{v}\ _2^2 \geq 0$	Always SPSPD for any $B \in \mathbb{R}^{n \times m}$

Table 3.1: Comprehensive overview of spectral properties, orthogonal projectors, and quadratic forms.

Part II

Continuous Optimization and Machine Learning

Chapter 4

Elementary Univariate Optimization Problems

4.1 Introduction, Motivation, and Learning Objectives

Mathematical optimization constitutes, alongside numerical linear algebra, the second computational pillar of machine learning and data science [4, 5]. In practical machine learning pipelines, virtually every model—from linear regression to deep neural architectures parameterizing hundreds of billions of weights—is conceptually formalized as the minimization of an empirical loss augmented by regularizing penalties over a designated admissible domain:

$$\min_{w \in \mathcal{W}} \mathcal{L}(w; \mathcal{D}) + \lambda \Omega(w) \quad (4.1)$$

where $\mathcal{L}$ represents the empirical loss over dataset $\mathcal{D}$, Ω imposes geometric regularization, and $\mathcal{W}$ delimits the feasible parameter set.

However, generic non-linear optimization is an intrinsically intractable endeavor: devoid of structural assumptions, locating the global minimum of an arbitrary function is formally unsolvable within finite time. To build a robust computational foundation, the pedagogical methodology adopted in this second part of the course follows an incremental, modular paradigm:

1. **Elementary univariate problem analysis:** examining linear and quadratic functions over $\mathbb{R}$, endowed with analytical structures so transparent that they admit closed-form solutions with $\mathcal{O}(1)$ computational complexity.
2. **Developing geometric and algebraic intuition:** gaining deep insight into gradients, curvatures, level sets, and the active influence of

boundary constraints.

3. **Identifying boundaries of computational tractability:** understanding precisely how relaxing a structural assumption or expanding into multidimensional spaces transforms elementary decisions into NP-hard challenges.
4. **Simple problems as algorithmic building blocks:** in virtually all state-of-the-art algorithms for non-linear and large-scale optimization (such as Newton, Quasi-Newton, Sequential Quadratic Programming, and line search techniques), the inner computational core repeatedly solves an elementary local approximation (linear or quadratic) of the complex global objective.

Methodological Philosophy

Mastering univariate linear and quadratic optimization is not merely an introductory exercise: it represents the foundational vocabulary in which advanced numerical algorithms for modern machine learning are written.

4.2 Fundamental Definitions: Spaces, Functions, Graphs, and Level Sets

4.2.1 Input and Output Spaces

Consider the simplest scalar unconstrained setting:

- **Input space:** $x \in \mathbb{R}$
- **Output space:** $y \in \mathbb{R}$

A scalar mapping $f : \mathbb{R} \rightarrow \mathbb{R}$ assigns to each input value x a unique real output $y = f(x)$.

4.2.2 Graph of a Function

Definition 4.2.1 (Graph). Given a scalar function $f : \mathbb{R} \rightarrow \mathbb{R}$, the **graph** of f , denoted by $\text{gr}(f)$, is the subset of the two-dimensional Cartesian product space $\mathbb{R}^2$ defined as:

$$\text{gr}(f) = \{(x, f(x)) \in \mathbb{R}^2 : x \in \mathbb{R}\} \subset \mathbb{R}^2 \quad (4.2)$$

Geometrically, for univariate functions, $\text{gr}(f)$ traces a planar curve in the (x, y) coordinate plane.

4.2.3 Image (Effective Codomain)

Definition 4.2.2 (Image). The **image** of a function f , denoted by $\text{im}(f)$, is the set of all real output values achieved by the mapping:

$$\text{im}(f) = \{y \in \mathbb{R} : \exists x \in \mathbb{R} \text{ such that } y = f(x)\} = f(\mathbb{R}) \subseteq \mathbb{R} \quad (4.3)$$

Geometrically, $\text{im}(f)$ corresponds to the orthogonal projection of the graph $\text{gr}(f)$ onto the vertical ordinate axis (y -axis).

4.2.4 Level Sets and Sublevel Sets

Definition 4.2.3 (Level Set). Given a target value $v \in \mathbb{R}$, the **level set** at value v , denoted by $L(f, v)$, is the pre-image of v within the domain:

$$L(f, v) = \{x \in \mathbb{R} : f(x) = v\} \subset \mathbb{R} \quad (4.4)$$

The geometric interpretation is direct: $L(f, v)$ represents the intersection between the curve $\text{gr}(f)$ and the horizontal line $y = v$, projected onto the abscissa axis (x -axis).

Remark 4.2.1 (Roots or Zeros of a Function). The special case where $v = 0$ identifies the set of roots of the function:

$$\text{roots}(f) = L(f, 0) = \{x \in \mathbb{R} : f(x) = 0\} \quad (4.5)$$

Definition 4.2.4 (Sublevel Set). Given a threshold $v \in \mathbb{R}$, the **sublevel set** at level v , denoted by $S(f, v)$, is defined as:

$$S(f, v) = \{x \in \mathbb{R} : f(x) \leq v\} \subset \mathbb{R} \quad (4.6)$$

Sublevel sets play a central role in convexity characterizations and convergence proofs for descent algorithms.

4.3 The Unconstrained Univariate Optimization Problem

4.3.1 Formal Problem Statement

Given an objective function $f : \mathbb{R} \rightarrow \mathbb{R}$, the unconstrained univariate minimization problem is formally stated as:

$$(P) \quad f^* = \min\{f(x) : x \in \mathbb{R}\} \quad (4.7)$$

4.3.2 Optimal Value vs Optimal Solution

A rigorous conceptual distinction must be maintained between the optimal value and an optimal minimizer:

Definition 4.3.1 (Optimal Value). The **optimal value** (denoted by $f^* \in \mathbb{R}$ or $\nu(P)$) is the scalar representing the infimum achieved by the image of the objective function:

$$f^* = \min(\text{im}(f)) = \min\{v \in \mathbb{R} : L(f, v) \neq \emptyset\} \quad (4.8)$$

If the minimum exists, the optimal value f^* is **strictly unique**.

Definition 4.3.2 (Optimal Solution or Minimizer). An **optimal solution** (denoted by x^*) is a point in the domain $x^* \in \mathbb{R}$ at which the function attains its optimal value:

$$f(x^*) = f^* \iff f(x^*) \leq f(x), \quad \forall x \in \mathbb{R} \quad (4.9)$$

The set of all minimizers is denoted using the set-valued arg min operator:

$$X^* = \arg \min\{f(x) : x \in \mathbb{R}\} = L(f, f^*) \quad (4.10)$$

Property 4.3.1 (Non-Uniqueness of Minimizers). *While the optimal value f^* is always a unique scalar, the set of minimizers X^* may:*

1. *Be empty ($X^* = \emptyset$, when the minimum is not attained).*
2. *Contain a single isolated point ($|X^*| = 1$, unique global minimizer).*
3. *Contain a discrete set of multiple points (e.g., $f(x) = \sin(x)$ with $f^* = -1$ and $X^* = \{-\pi/2 + 2k\pi : k \in \mathbb{Z}\}$).*
4. *Contain a continuum of points (e.g., a constant function over an interval).*

4.3.3 Indifference Between Equivalent Minimizers

From an optimization perspective:

$$\forall x_1^*, x_2^* \in X^* \implies f(x_1^*) = f(x_2^*) = f^* \quad (4.11)$$

All points within X^* are perfectly indistinguishable to the objective function. Barring secondary criteria or explicit regularizations (such as minimal ℓ_2 -norm or ℓ_1 -sparsity preferences), any numerical solver is mathematically authorized to return an arbitrary element of X^* .

4.4 Simple Reformulations and Equivalence

Two optimization problems are termed **equivalent** if solving one directly yields the solution of the other without requiring non-trivial iterative computations.

4.4.1 Min–Max Duality

Without loss of generality (w.l.o.g.), optimization theory can be formulated strictly in terms of minimization problems. Any maximization problem can be transformed by negating the objective function:

$$\min\{f(x) : x \in \mathbb{R}\} = -\max\{-f(x) : x \in \mathbb{R}\} \quad (4.12)$$

$$\arg \min\{f(x) : x \in \mathbb{R}\} = \arg \max\{-f(x) : x \in \mathbb{R}\} \quad (4.13)$$

Conceptual Caution

Do not confuse this equivalence with the relationship between $\min f(x)$ and $\max f(x)$: minimizing $f(x)$ and maximizing $f(x)$ are fundamentally distinct problems with opposite convexity and curvature properties. The equivalence holds exclusively between minimizing $f(x)$ and maximizing $-f(x)$.

4.4.2 Invariance Under Vertical Shifts (Additive Constant)

Adding an arbitrary real constant $c \in \mathbb{R}$ shifts the graph vertically along the ordinate axis:

$$\min\{f(x) + c : x \in \mathbb{R}\} = \min\{f(x) : x \in \mathbb{R}\} + c = f^* + c \quad (4.14)$$

$$\arg \min\{f(x) + c : x \in \mathbb{R}\} = \arg \min\{f(x) : x \in \mathbb{R}\} = X^* \quad (4.15)$$

The optimal value shifts by c , while the set of optimal minimizers X^* remains invariant.

4.4.3 Invariance Under Positive Scaling (Multiplicative Constant)

Multiplying the objective by a strictly positive scalar $c > 0$ vertically dilates or compresses the graph:

$$\min\{c \cdot f(x) : x \in \mathbb{R}\} = c \cdot \min\{f(x) : x \in \mathbb{R}\} = cf^* \quad (c > 0) \quad (4.16)$$

$$\arg \min\{c \cdot f(x) : x \in \mathbb{R}\} = \arg \min\{f(x) : x \in \mathbb{R}\} = X^* \quad (4.17)$$

If $c < 0$, the negative scalar inverts the orientation, turning minimization into maximization:

$$\min\{c \cdot f(x) : x \in \mathbb{R}\} = c \cdot \max\{f(x) : x \in \mathbb{R}\} = -|c| \max\{f(x) : x \in \mathbb{R}\} \quad (4.18)$$

4.5 Constrained Optimization and the Feasible Region

4.5.1 General Formulation of Constrained Optimization

In engineering applications and machine learning, decision parameters are bounded by physical limitations, resource budgets, or algorithmic stability bounds. We define the **feasible region** $X \subseteq \mathbb{R}$:

$$(P_X) \quad f^* = \min\{f(x) : x \in X\} \quad (4.19)$$

- A point $x \in X$ is a **feasible solution**.
- A point $x \in \mathbb{R} \setminus X$ is an **infeasible solution**.
- The constrained optimal value is $f^* = \min(\text{im}(X, f)) = \min\{f(x) : x \in X\}$.
- The constrained optimal solution set is given by:

$$X^* = L(f, f^*) \cap X \quad (4.20)$$

4.5.2 Monotone Non-Improvement Under Constraint Restriction

Theorem 4.5.1 (Fundamental Constraint Inequality). *Let $X_1 \subseteq X_2 \subseteq \mathbb{R}$ be nested feasible regions. Then:*

$$\min_{x \in X_1} f(x) \geq \min_{x \in X_2} f(x) \quad (4.21)$$

In particular, for any constrained set $X \subseteq \mathbb{R}$:

$$\min_{x \in X} f(x) \geq \min_{x \in \mathbb{R}} f(x) \quad (4.22)$$

Proof. Let $x_1^* \in \arg \min_{x \in X_1} f(x)$. Since $X_1 \subseteq X_2$, it follows that $x_1^* \in X_2$. By definition of global minimization over X_2 , $f(x) \geq \min_{z \in X_2} f(z)$ for any $x \in X_2$. Evaluating at $x = x_1^*$ completes the proof. $\square$

Operational Insight

Imposing additional constraints can never improve the optimal objective value in a minimization problem; it can only leave it unchanged or worsen it (i.e., increase it).

4.5.3 Taxonomy of Redundant Constraints

Given a constrained problem $\min_{x \in X} f(x)$ with unconstrained minimizer set $X_{\text{free}}^* = \arg \min_{x \in \mathbb{R}} f(x)$:

- **Completely redundant (inactive) constraint:** occurs when $X_{\text{free}}^* \subseteq X$. The constrained optimal value equals the unconstrained value ($f_X^* = f_{\mathbb{R}}^*$) and the set of optimal solutions is unchanged.
- **Partially redundant constraint:** occurs when $X_{\text{free}}^* \cap X \neq \emptyset$, but $X_{\text{free}}^* \not\subseteq X$. The constraint eliminates some alternative optimal solutions without worsening the optimal value ($f_X^* = f_{\mathbb{R}}^*$ remains intact, while $X_X^* \subset X_{\text{free}}^*$ shrinks).
- **Active and essential constraint:** occurs when $X_{\text{free}}^* \cap X = \emptyset$. The constraint forces the solution onto the boundary of X , strictly increasing the optimal value ($f_X^* > f_{\mathbb{R}}^*$).

4.5.4 Functional Representation of Feasible Sets

In computational practice, the set X is not defined as an abstract collection, but via systems of algebraic equations and inequalities based on constraint functions $g_i : \mathbb{R} \rightarrow \mathbb{R}$:

1. Equality Constraints:

$$g(x) = v \iff X = L(g, v) \quad (4.23)$$

In canonical homogeneous form:

$$g(x) = 0 \quad (4.24)$$

2. Inequality Constraints:

$$g(x) \leq v \iff X = S(g, v) = \{x \in \mathbb{R} : g(x) \leq v\} \quad (4.25)$$

In canonical normalized form:

$$g(x) \leq 0 \quad (4.26)$$

When the constraint is originally given as $g(x) \geq 0$, multiplying by -1 recovers the canonical format $-g(x) \leq 0$.

3. Conjunction of Multiple Constraints:

Multiple constraints correspond to the intersection of their respective sublevel sets:

$$\begin{cases} g_1(x) \leq 0 \\ g_2(x) \leq 0 \end{cases} \iff X = S(g_1, 0) \cap S(g_2, 0) \quad (4.27)$$

4. Box Constraints (Closed Intervals):

Representing the fundamental class in univariate optimization:

$$x_- \leq x \leq x_+ \iff X = [x_-, x_+] \quad (4.28)$$

which is canonicalized as the two-inequality system:

$$\begin{cases} x - x_+ \leq 0 \\ x_- - x \leq 0 \end{cases} \quad (4.29)$$

4.6 Why Optimization is Difficult: Theory vs Computational Practice

4.6.1 Modes of Non-Existence of the Minimum

An optimization problem can fail to admit an optimal minimizer due to two distinct topological reasons:

Case 1: Unbounded Below Problem ($f^* = -\infty$)

The objective decreases indefinitely across the domain:

$$\forall M \in \mathbb{R}, \quad \exists x \in X \quad \text{such that} \quad f(x) \leq M \quad (4.30)$$

A canonical example is $f(x) = x$ as $x \rightarrow -\infty$. In this setting, we write $f^* = -\infty$ and $\arg \min f(x) = \emptyset$.

Case 2: Finite Infimum Not Attained (f^* finite, but $x^* \notin X$)

The image $\text{im}(f)$ is bounded below in $\mathbb{R}$, but the infimum does not belong to the image:

$$\inf\{f(x) : x \in X\} > -\infty \quad \text{yet} \quad \nexists x^* \in X \text{ such that } f(x^*) = \inf_{x \in X} f(x) \quad (4.31)$$

Two typical manifestations include:

- **Smooth asymptotic settings:** $f(x) = e^x$ on $X = \mathbb{R}$. The infimum is $\inf_{x \in \mathbb{R}} e^x = 0$, but $e^x = 0$ has no finite solution in $\mathbb{R}$. The minimum is approached only asymptotically as $x \rightarrow -\infty$.
- **Discontinuous pathological settings:**

$$f(x) = \begin{cases} 1 & \text{if } x = 0 \\ |x| & \text{if } x \neq 0 \end{cases} \quad (4.32)$$

The infimum is 0, but at the natural candidate $x = 0$ the function jumps to 1. No point attains $f(x) = 0$.

4.6.2 Formalization via Extended Real Numbers $\overline{\mathbb{R}}$

To mathematically handle these asymptotic limits, real space is extended by incorporating infinity:

$$\overline{\mathbb{R}} = \mathbb{R} \cup \{-\infty, +\infty\} \quad (4.33)$$

In $\overline{\mathbb{R}}$, every subset $S \subseteq \mathbb{R}$ **always** possesses both an infimum and a supremum:

$$\inf S = \max\{v \in \overline{\mathbb{R}} : v \leq s, \forall s \in S\} \quad (4.34)$$

$$\sup S = \min\{v \in \overline{\mathbb{R}} : v \geq s, \forall s \in S\} \quad (4.35)$$

With empty set conventions:

$$\inf \emptyset = +\infty, \quad \sup \emptyset = -\infty \quad (4.36)$$

4.6.3 Optimization in Finite-Precision Floating-Point Arithmetic

In computational hardware (modern CPUs and GPUs), the continuous real line $\mathbb{R}$ does not exist: algorithms operate strictly in finite-precision floating-point arithmetic (IEEE 754 standard ‘double’ 64-bit, ‘single’ 32-bit, or reduced formats like FP8/INT4 in deep learning accelerators).

This imposes key algorithmic realities:

1. **Finite discrete grid:** between any two consecutive representable numbers lies an indivisible gap determined by machine precision ($\varepsilon_{\text{mach}} \approx 10^{-16}$ in ‘double’).
2. **Impossibility of exact zero-error convergence:** seeking an exact minimizer x^* is computationally meaningless. Stopping criteria must accept an ε -approximate solution (**ε -optimum**):

$$f^* \leq f(x_\varepsilon) \leq f^* + \varepsilon \quad (4.37)$$

4.6.4 Error Metrics: Absolute Gap vs Relative Gap

To assess solution accuracy:

Definition 4.6.1 (Absolute Gap).

$$A(x) = f(x) - f^* \geq 0 \quad (4.38)$$

Limitation: absolute error is scale-dependent. If $f^* = 10^7$, an error $A(x) = 1$ is exceptional; if $f^* = 10^{-7}$, $A(x) = 1$ renders the solution completely useless.

Definition 4.6.2 (Normalized Relative Gap).

$$R(x) = \frac{f(x) - f^*}{\max(|f^*|, 1)} \geq 0 \quad (4.39)$$

Advantages: dimension-free and invariant to positive scaling of the objective. The term $\max(|f^*|, 1)$ in the denominator guards against numerical division-by-zero when $f^* \rightarrow 0$.

4.6.5 The Impossibility of Generic Black-Box Optimization

Suppose an objective $f : [x_-, x_+] \rightarrow \mathbb{R}$ is accessible solely through a black-box oracle: given input x , the oracle returns only the scalar value $f(x)$, with zero access to gradients or structural properties.

Theorem 4.6.1 (Informal Impossibility Theorem for Adversarial Oracles). *It is impossible to identify the global minimum of an arbitrary continuous function on a bounded interval $[x_-, x_+]$ in a finite number of oracle queries.*

Adversarial Oracle Argument (Needle in a Haystack). Consider an interval $[x_-, x_+]$ and a function identically zero $f(x) = 0$ everywhere except at an unknown single point x_0 , where it exhibits a sharp negative spike:

$$f(x_0) = -1, \quad f(x) = 0 \quad \forall x \neq x_0 \quad (4.40)$$

Any algorithm querying a finite sample $\{x_1, \dots, x_K\}$ receives return value 0. An adversarial oracle can always place x_0 in $[x_-, x_+] \setminus \{x_1, \dots, x_K\}$.

Even if we enforce continuity by connecting $(x_0, -1)$ via a very narrow cusp of base width $\delta > 0$, any discrete sampling step $\Delta x > \delta$ fails to detect the spike. $\square$

Necessary Condition for Computational Tractability

Optimization becomes computationally viable only when the objective function exhibits **exploitable mathematical structure**: Lipschitz continuity, differentiability, convexity, or linearity.

4.7 Univariate Linear Optimization

4.7.1 Geometric and Analytical Properties

The simplest non-constant univariate function is the homogeneous linear function:

$$f(x) = bx, \quad b \in \mathbb{R} \quad (4.41)$$

The parameter b defines the **slope**:

- If $b > 0$: f is *strictly increasing*:

$$x_1 > x_2 \iff x_1 - x_2 > 0 \implies b(x_1 - x_2) > 0 \iff f(x_1) > f(x_2) \quad (4.42)$$

- If $b < 0$: f is *strictly decreasing*.
- If $b = 0$: f degenerates to the constant zero function $f(x) = 0$.

4.7.2 Unconstrained Optimization over $\mathbb{R}$

The unconstrained problem $\min_{x \in \mathbb{R}} bx$ is trivial:

- If $b > 0$: $\lim_{x \rightarrow -\infty} bx = -\infty \implies f^* = -\infty$ (unbounded below).
- If $b < 0$: $\lim_{x \rightarrow +\infty} bx = -\infty \implies f^* = -\infty$ (unbounded below).
- If $b = 0$: $f^* = 0$, $X^* = \mathbb{R}$ (every real point is an optimal minimizer).

4.7.3 Optimization Over a Bounded Box $[x_-, x_+]$

Consider the constrained problem over a closed bounded interval $[x_-, x_+]$:

$$\min\{bx : x \in [x_-, x_+]\} \quad (4.43)$$

Owing to global strict monotonicity, the minimum and maximum necessarily lie on the boundary points. The decision rule operates in $\mathcal{O}(1)$ time:

Condition	Minimum x^*	Min Value f^*	Maximum $x_{\max}^*$	Max Value $f_{\max}^*$
$b > 0$	x_-	bx_-	x_+	bx_+
$b < 0$	x_+	bx_+	x_-	bx_-
$b = 0$	Indifferent	0	Indifferent	0

Table 4.1: Decision rules at $\mathcal{O}(1)$ computational cost for box-constrained linear optimization.

4.8 The Illusion of Open Sets in Numerical Computation

4.8.1 The Analytical Failure of Open Intervals

Consider an open feasible interval:

$$X = (x_-, x_+) = \{x \in \mathbb{R} : x_- < x < x_+\} \quad (4.44)$$

Assuming $b > 0$, the function $f(x) = bx$ has infimum $\inf_{x \in X} f(x) = bx_-$. However:

$$\nexists x \in (x_-, x_+) \text{ such that } f(x) = bx_- \quad (4.45)$$

In pure analysis, the problem **has no solution**.

4.8.2 The Practical Numerical Resolution

In applied numerical computation:

1. **Physical parameters do not possess infinite precision:** no engineering constraint requires $x > x_-$ without an explicit safety tolerance.
2. **Floating-point arithmetic contains no open sets:** between x_- and the adjacent machine number $x_- + \varepsilon_{\text{mach}}$, no values exist.

3. **Formulation with explicit buffer tolerances:** boundaries are backed off via a positive numerical threshold:

$$X = [x_- + \varepsilon_-, x_+ - \varepsilon_+] \quad (4.46)$$

Golden Methodological Principle

Throughout this course, feasible regions and box constraints will always be modeled as **closed and bounded**.

4.9 Univariate Optimization: The Homogeneous Quadratic Case

4.9.1 Definition and Curvature

A homogeneous quadratic function is given by:

$$f(x) = ax^2, \quad a \in \mathbb{R} \quad (4.47)$$

The parameter a represents the **curvature**:

- $a > 0$ (**Convex Function**): the parabola curves upwards. The function is strictly decreasing for $x \leq 0$ and strictly increasing for $x \geq 0$. The vertex $(0, 0)$ is the global minimum. As a grows, the parabola steepens; as $a \rightarrow 0^+$, it flattens.
- $a < 0$ (**Concave Function**): the parabola curves downwards with vertex $(0, 0)$ as the global maximum.
- $a = 0$: the function collapses to the constant zero line.

4.9.2 Unconstrained Optimization over $\mathbb{R}$

- If $a > 0$: $x^* = 0$, $f^* = 0$, with $\max_{x \in \mathbb{R}} ax^2 = +\infty$.
- If $a < 0$: $x_{\max}^* = 0$, $f_{\max}^* = 0$, with $\min_{x \in \mathbb{R}} ax^2 = -\infty$.
- If $a = 0$: $f^* = 0$, any point $x \in \mathbb{R}$ is optimal.

4.9.3 Box-Constrained Optimization $[x_-, x_+]$ (for $a > 0$)

For a closed bounded interval $[x_-, x_+]$, the position of the vertex ($x = 0$) relative to the interval defines three mutually exclusive configurations:




Figure 4.1: The three topological configurations for box-constrained homogeneous quadratic optimization over $[x_-, x_+]$.

1. **Strictly Negative Configuration** ($x_+ < 0$): the entire box lies to the left of the vertex. The convex parabola is strictly decreasing across the box:

$$x^* = x_+, \quad f^* = a(x_+)^2 \quad (4.48)$$

$$x_{\max}^* = x_-, \quad f_{\max}^* = a(x_-)^2 \quad (4.49)$$

2. **Strictly Positive Configuration** ($x_- > 0$): the entire box lies to the right of the vertex. The objective is strictly increasing:

$$x^* = x_-, \quad f^* = a(x_-)^2 \quad (4.50)$$

$$x_{\max}^* = x_+, \quad f_{\max}^* = a(x_+)^2 \quad (4.51)$$

3. **Interior Vertex Configuration** ($x_- \leq 0 \leq x_+$): the vertex is feasible, so the constrained minimum equals the unconstrained global minimum:

$$x^* = 0, \quad f^* = 0 \quad (4.52)$$

The maximum is achieved at whichever boundary point lies furthest from the origin:

$$x_{\max}^* = \arg \max \{a(x_-)^2, a(x_+)^2\} = \begin{cases} x_- & \text{if } |x_-| > |x_+| \\ x_+ & \text{if } |x_+| > |x_-| \end{cases} \quad (4.53)$$

4.10 Univariate Optimization: The Non-Homogeneous Quadratic Case

4.10.1 Definition and Non-Separability of the Minimum Operator

Consider the complete quadratic scalar function with a linear term:

$$f(x) = ax^2 + bx, \quad a, b \in \mathbb{R}, \quad a \neq 0, b \neq 0 \quad (4.54)$$

Fundamental Non-Additivity of the Minimum Operator

In general, minimization does not distribute over sums:

$$\min\{ax^2 + bx\} \neq \min\{ax^2\} + \min\{bx\} \quad (4.55)$$

The minimum of the sum of two functions does not equal the sum of their individual minima, unless they happen to share the exact same minimizer (which is violated here because the linear term bx has no finite minimum on $\mathbb{R}$).

4.10.2 The Coordinate Translation Principle (Origin Shift)

To solve this problem without relying on differential calculus, we apply one of the most powerful algebraic techniques in optimization: **coordinate translation**. We define a shifted variable z around a center $\bar{x} \in \mathbb{R}$ to be determined:

$$z = x - \bar{x} \quad \Longleftrightarrow \quad x = z + \bar{x} \quad (4.56)$$

such that, when expressed in terms of z , the quadratic function becomes **purely homogeneous** (canceling the linear term in z).

4.10.3 Step-by-Step Algebraic Derivation

Substituting $x = z + \bar{x}$ into $f(x)$:

$$\begin{aligned} f(z + \bar{x}) &= a(z + \bar{x})^2 + b(z + \bar{x}) \\ &= a(z^2 + 2z\bar{x} + \bar{x}^2) + bz + b\bar{x} \\ &= az^2 + 2a\bar{x}z + a\bar{x}^2 + bz + b\bar{x} \end{aligned} \quad (4.57)$$

Grouping terms by descending powers of z :

$$f(z + \bar{x}) = az^2 + (2a\bar{x} + b)z + (a\bar{x}^2 + b\bar{x}) \quad (4.58)$$

Notice that the constant term independent of z is exactly the value of the original function evaluated at $\bar{x}$:

$$a\bar{x}^2 + b\bar{x} = f(\bar{x}) \quad (4.59)$$

Thus the transformed function is written as:

$$f(z + \bar{x}) = az^2 + (2a\bar{x} + b)z + f(\bar{x}) \quad (4.60)$$

To eliminate the linear term and render the function homogeneous in z , we set the coefficient of z to zero:

$$2a\bar{x} + b = 0 \implies \bar{x} = -\frac{b}{2a} \quad (4.61)$$

4.10.4 Solution in Shifted Coordinates and Inverse Mapping

Setting $\bar{x} = -\frac{b}{2a}$, the function in z becomes decoupled:

$$g(z) = f(z + \bar{x}) = az^2 + f(\bar{x}) \quad (4.62)$$

Since $f(\bar{x})$ is an additive constant independent of the optimization variable z :

$$\arg \min_{z \in \mathbb{R}} g(z) = \arg \min_{z \in \mathbb{R}} (az^2 + f(\bar{x})) = \arg \min_{z \in \mathbb{R}} (az^2) \quad (4.63)$$

For $a > 0$ (convex parabola), the minimizer in shifted space is immediate:

$$z^* = 0 \quad (4.64)$$

Mapping back to the original coordinate space:

$$x^* = z^* + \bar{x} = 0 + \bar{x} = -\frac{b}{2a} \quad (4.65)$$

The optimal value is found by evaluating f at $x^* = \bar{x}$:

$$\begin{aligned} f^* = f(\bar{x}) &= a \left(-\frac{b}{2a} \right)^2 + b \left(-\frac{b}{2a} \right) \\ &= a \frac{b^2}{4a^2} - \frac{b^2}{2a} = \frac{b^2}{4a} - \frac{2b^2}{4a} = -\frac{b^2}{4a} \end{aligned} \quad (4.66)$$




Figure 4.2: Geometric visualization of coordinate translation for the non-homogeneous parabola: the vertex x^* lies precisely at the midpoint between the two roots x_1 and x_2 .

4.10.5 Geometric Interpretation: The Midpoint of the Roots

The quadratic function $f(x) = ax^2 + bx = x(ax + b)$ always possesses two real roots (when $b \neq 0$):

$$x_1 = 0, \quad x_2 = -\frac{b}{a} \quad (4.67)$$

The optimal minimizer derived via coordinate shift lies exactly at the **midpoint** of the roots:

$$\bar{x} = \frac{x_1 + x_2}{2} = \frac{0 + (-b/a)}{2} = -\frac{b}{2a} \quad (4.68)$$

This universal symmetry property reflects the fact that the derivative of a degree-two polynomial is an affine line intersecting the horizontal axis at the exact arithmetic mean of its real zeros.

4.10.6 Box-Constrained Optimization for General Quadratic Functions

Under a box constraint $x \in [x_-, x_+]$, coordinate translation directly shifts the boundary endpoints for variable z :

$$x \in [x_-, x_+] \iff z \in [x_- - \bar{x}, x_+ - \bar{x}] \quad (4.69)$$

Setting $z_- = x_- - \bar{x}$ and $z_+ = x_+ - \bar{x}$, the problem reduces to the homogeneous quadratic problem over $[z_-, z_+]$ solved in Section 4.9:

$$z^* = \begin{cases} z_+ & \text{if } z_+ < 0 \\ z_- & \text{if } z_- > 0 \\ 0 & \text{if } z_- \leq 0 \leq z_+ \end{cases} \quad (4.70)$$

recovering the original variable solution via $x^* = z^* + \bar{x}$.

4.11 Synoptic Formula Reference of Univariate Optimization

Class	Form	Condition	Minimum x^*	Value f^*
Linear	bx	$b > 0$	Does not exist ($x \rightarrow -\infty$)	$-\infty$
		$b < 0$	Does not exist ($x \rightarrow +\infty$)	$-\infty$
		$b = 0$ (degenerate)	Indifferent ($\forall x \in \mathbb{R}$)	0
Homogeneous quadratic	ax^2	$a > 0$ (convex)	$x^* = 0$ (unique)	0
		$a < 0$ (concave)	Does not exist ($x \rightarrow \pm\infty$)	$-\infty$
		$a = 0$ (degenerate)	Indifferent ($\forall x \in \mathbb{R}$)	0
General quadratic	$ax^2 + bx$	$a > 0$ (convex)	$x^* = -b/(2a)$ (vertex)	$-b^2/(4a)$
		$a < 0$ (concave)	Does not exist ($x \rightarrow \pm\infty$)	$-\infty$
		$a = 0, b \neq 0$	Does not exist (linear reduction)	$-\infty$

Table 4.2: Comprehensive synoptic summary of univariate unconstrained optimization.

Chapter 5

Multivariate Optimization: Geometry, Linear Functions, and Quadratic Forms

5.1 From Univariate to Multivariate: Dimensionality of $\mathbb{R}^n$

5.1.1 Multivariate Input Space

In modern data science, machine learning, and computational engineering, models almost universally depend on a collection of mutually interacting variables. The primary mathematical object of study is the real-valued multivariate scalar function:

$$f : \mathbb{R}^n \rightarrow \mathbb{R}, \quad f(x) = f(x_1, x_2, \dots, x_n) \quad (5.1)$$

where the input argument x is a structured column vector of n real numbers:

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = [x_i]_{i=1}^n \in \mathbb{R}^n \quad (5.2)$$

Vector dimension n spans radically disparate orders of magnitude across application domains:

- **Small scale** ($n \in \{2, 3, 10^2\}$): Euclidean geometric models, inverse kinematics, elementary resource allocation.

- **Medium and large scale** ($n \in \{10^4, 10^5\}$): image and signal processing, finite element structural mechanics, network flow logistics.
- **Monstrous scale** (*heinously large*, $n \in \{10^9, 10^{11}\}$): training deep neural networks and Large Language Models (LLMs), where n enumerates synaptic weights and attention matrices.

5.1.2 The Curse of Dimensionality

The continuous vector space $\mathbb{R}^n$ is formally defined as the Cartesian product of $\mathbb{R}$ taken n times:

$$\mathbb{R}^n = \underbrace{\mathbb{R} \times \mathbb{R} \times \cdots \times \mathbb{R}}_{n \text{ times}} \quad (5.3)$$

Exponential Volumetric Expansion

The volume of available space within $\mathbb{R}^n$ grows exponentially with dimension n . Adding variables to an optimization model is never computationally neutral: it fundamentally changes the geometric density of the searchable space.

To quantitatively demonstrate this combinatorial explosion, consider a drastic simplification:

1. Restrict the continuous domain to a compact unit hypercube (*box*):

$$X = \{x \in \mathbb{R}^n : 0 \leq x_i \leq 1, i = 1, \dots, n\} = [0, 1]^n \quad (5.4)$$

2. Discard interior continuity entirely, sampling only the extreme discrete corner values for each coordinate, $x_i \in \{0, 1\}$.

The resulting set coincides with the vertices of the **Boolean hypercube** $\{0, 1\}^n$, whose cardinality is:

$$|V| = 2^n \text{ vertices} \quad (5.5)$$

Even for modest dimensions (such as $n = 100$), $2^{100} \approx 1.26 \times 10^{30}$. This number dwarfs the processing capacity of all existing computing clusters, rendering brute-force grid enumeration information-theoretically impossible.

5.2 Vector and Multi-Objective Optimization

5.2.1 The Limitation of a Single Scalar Objective

Standard optimization assumes that the overall merit of any decision $x \in X$ can be collapsed into a single real number $f(x) \in \mathbb{R}$.

Because the real line $\mathbb{R}$ is endowed with a **total order**:

$$\forall a, b \in \mathbb{R} \implies a \leq b \quad \vee \quad b \leq a \quad (5.6)$$

any two candidate solutions x' and x'' can always be unambiguously ranked by comparing $f(x')$ with $f(x'')$.

In practice, complex decision problems require balancing multiple competing, incommensurable criteria (comparing “apples to oranges”):

- **Vehicle procurement:** capital cost vs fuel economy vs safety ratings vs styling appeal.
- **Machine learning:** minimizing empirical data loss $\mathcal{L}(w)$ vs minimizing model complexity $R(w)$ to prevent overfitting.

A formal **multi-objective** (or vector) optimization problem is posed as:

$$\min_{x \in X} F(x) = \begin{bmatrix} f_1(x) \\ f_2(x) \\ \vdots \\ f_k(x) \end{bmatrix}, \quad F : X \rightarrow \mathbb{R}^k, \quad k \geq 2 \quad (5.7)$$

5.2.2 Illustrative Example: Markowitz Portfolio Selection

Consider an investor distributing capital across a set of financial assets:

- **Objective 1 (Expected Return):** $f_1(x) \in \mathbb{R}$, to be maximized (or $-f_1(x)$ minimized).
- **Objective 2 (Financial Risk / Volatility):** $f_2(x) \in \mathbb{R}$ (return variance), to be minimized.

Safe assets (e.g., short-term government bonds) deliver low returns, whereas high-yield equities carry significant volatility. No single portfolio simultaneously minimizes variance and maximizes return.

5.2.3 Partial Ordering and the Pareto Frontier

For $k \geq 2$, vector space $\mathbb{R}^k$ lacks a natural total ordering. Instead, it is governed by a **partial ordering** induced by the non-negative orthant cone:

$$u \leq v \iff u_i \leq v_i \quad \forall i = 1, \dots, k \quad (5.8)$$

Definition 5.2.1 (Pareto Dominance). Assuming by convention that all k objectives are to be minimized, a feasible point $x' \in X$ **Pareto-dominates** $x'' \in X$ (denoted $x' \succ_P x''$) if and only if:

$$f_i(x') \leq f_i(x'') \quad \forall i \in \{1, \dots, k\} \quad \wedge \quad \exists j \in \{1, \dots, k\} : f_j(x') < f_j(x'') \quad (5.9)$$

Definition 5.2.2 (Non-Dominated or Pareto-Optimal Solution). A point $x^* \in X$ is **Pareto-optimal** (or non-dominated) if no other point $x \in X$ dominates x^* . The collection of all objective values $F(x^*)$ forms the **Pareto Frontier**.




Figure 5.1: Geometric representation of the Pareto Frontier in Markowitz portfolio selection.

An optimization algorithm cannot objectively choose a single point on the Pareto frontier: the final selection rests with human decision-makers and their risk preferences.

5.2.4 Operational Reduction Techniques: Scalarization and Budgeting

To apply continuous optimization solvers to multi-objective problems, two strategies are standard:

1. **Scalarization:** Formulate a weighted linear or convex combination of objectives:

$$\min_{x \in X} \{f_1(x) + \alpha f_2(x)\}, \quad \alpha > 0 \quad (5.10)$$

Here α serves as an *exchange rate*: how much degradation in f_1 is tolerated to gain one unit of improvement in f_2 . Sweeping $\alpha \in (0, +\infty)$ traces the Pareto frontier. In machine learning, regularized empirical risk minimization:

$$\min_w \{\mathcal{L}(w) + \lambda R(w)\} \quad (5.11)$$

is a classic scalarization, where trade-off parameter λ is selected via cross-validation.

2. **Constraint Method (*Budgeting*):** Designate a primary objective to optimize, downgrading all other criteria to upper/lower bound constraints (budgets):

$$\min_{x \in X} \{f_1(x) : f_2(x) \leq \beta_2, \dots, f_k(x) \leq \beta_k\} \quad (5.12)$$

Systematically varying the budget vector β traces the non-dominated set.

Standard Convention

Unless stated otherwise, modeling is assumed to have already consolidated the problem into a single scalar objective $f : \mathbb{R}^n \rightarrow \mathbb{R}$.

5.3 Euclidean Geometry in $\mathbb{R}^n$ and Function Tomography

5.3.1 Standard Inner Product and Induced Norm

For any two vectors $x, z \in \mathbb{R}^n$, the **standard Euclidean inner product** is:

$$\langle x, z \rangle = x^T z = \sum_{i=1}^n x_i z_i = x_1 z_1 + x_2 z_2 + \dots + x_n z_n \quad (5.13)$$

The inner product satisfies symmetry ($\langle x, z \rangle = \langle z, x \rangle$), bilinearity, and positive definiteness ($\langle x, x \rangle > 0$ for all $x \neq 0$).

The **Euclidean norm** (ℓ_2 -norm) measures geometric length:

$$\|x\| = \sqrt{\langle x, x \rangle} = \sqrt{\sum_{i=1}^n x_i^2} \quad (5.14)$$

5.3.2 Angular Representation and the Cauchy-Schwarz Inequality

Geometrically, the inner product relates lengths to the angle $\theta \in [0, \pi]$ between x and z :

$$\langle x, z \rangle = \|x\| \|z\| \cos(\theta) \quad (5.15)$$

yielding three geometric orientations:

- $\langle x, z \rangle > 0 \iff \theta \in [0, \pi/2)$: acute angle (pointing in the same general direction).
- $\langle x, z \rangle = 0 \iff \theta = \pi/2$: **orthogonal** vectors ($x \perp z$).
- $\langle x, z \rangle < 0 \iff \theta \in (\pi/2, \pi]$: obtuse angle (pointing in opposing directions).

Since $|\cos(\theta)| \leq 1$, the fundamental **Cauchy-Schwarz inequality** follows directly:

$$|\langle x, z \rangle| \leq \|x\| \|z\|, \quad \forall x, z \in \mathbb{R}^n \quad (5.16)$$

with equality holding if and only if x and z are collinear ($x = \gamma z$ for some $\gamma \in \mathbb{R}$).

5.3.3 The Concept of Function Tomography

In multivariate space ($n \geq 3$), the graph $\text{gr}(f) = \{(x, f(x)) : x \in \mathbb{R}^n\} \subset \mathbb{R}^{n+1}$ cannot be visualized. The central analytical tool used to investigate its local and global behavior is **tomography**.

Definition 5.3.1 (Directional Tomography). Given a base point $x \in \mathbb{R}^n$ and a non-zero directional vector $d \in \mathbb{R}^n \setminus \{0\}$, the **tomography** of f from x along direction d is the univariate function $\phi_{x,d} : \mathbb{R} \rightarrow \mathbb{R}$:

$$\phi_{x,d}(\alpha) = f(x + \alpha d), \quad \alpha \in \mathbb{R} \quad (5.17)$$

Scalar α represents the signed displacement step along the line through x oriented by d .

Directional norm $\|d\|$ simply scales the parameter α :

$$\phi_{x,\beta d}(\alpha) = f(x + \alpha(\beta d)) = \phi_{x,d}(\alpha\beta) \quad (5.18)$$

To eliminate scaling ambiguity, it is standard practice to normalize the search direction:

$$\|d\| = 1 \quad (5.19)$$

5.4 Multivariate Linear Functions

5.4.1 Definition and Separability

A multivariate function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is **linear** if it can be represented as an inner product with a fixed coefficient vector $b \in \mathbb{R}^n$:

$$f(x) = \langle b, x \rangle = b^T x = \sum_{i=1}^n b_i x_i \quad (5.20)$$

Every multivariate linear function is inherently **separable**: it is the sum of n independent univariate linear functions:

$$f(x) = \sum_{i=1}^n f_i(x_i), \quad \text{with } f_i(x_i) = b_i x_i \quad (5.21)$$

There are no bilinear cross-terms between different variables.

5.4.2 Geometry of Graph and Level Sets

- **Graph:** $\text{gr}(f) = \{(x, b^T x) : x \in \mathbb{R}^n\}$ is an n -dimensional hyperplane in $\mathbb{R}^{n+1}$ passing through the origin.
- **Level Sets:** For any target value $v \in \mathbb{R}$, the level set is:

$$L(f, v) = \{x \in \mathbb{R}^n : b^T x = v\} \quad (5.22)$$

This set forms an $(n - 1)$ -dimensional affine hyperplane in $\mathbb{R}^n$. As v varies, level sets form an infinite family of parallel hyperplanes.

Theorem 5.4.1 (Orthogonality of Vector b to Level Sets). *The coefficient vector b is strictly orthogonal to any tangent direction within the level sets $L(f, v)$, for all $v \in \mathbb{R}$.*

Proof. Let $x, z \in L(f, v)$ be two arbitrary points on the same level set. By definition:

$$b^T x = v \quad \wedge \quad b^T z = v \quad (5.23)$$

Subtracting these expressions yields:

$$b^T z - b^T x = 0 \iff b^T(z - x) = 0 \iff \langle b, z - x \rangle = 0 \quad (5.24)$$

Vector $z - x$ represents an arbitrary displacement tangent to the level hyperplane. Because b is orthogonal to every displacement within the level set, $b \perp L(f, v)$. $\square$

5.5 Linear Tomography and Steepest Descent Direction

5.5.1 Tomographic Cross-Section Computation

Fixing the origin $x = 0$ and a unit direction $\|d\| = 1$, the tomography of linear function $f(x) = b^T x$ is:

$$\phi_{0,d}(\alpha) = f(0 + \alpha d) = b^T(\alpha d) = \alpha \langle b, d \rangle = \alpha \|b\| \|d\| \cos(\theta) = \alpha \|b\| \cos(\theta) \quad (5.25)$$

where θ is the angle between coefficient vector b and direction d .

The tomography is a **univariate line** through the origin whose slope is governed by $\langle b, d \rangle$.

5.5.2 Directional Analysis: The Radar Screen Analogy

Rotating direction d through all angles reveals five fundamental behaviors:

1. **d perfectly aligned with b ($\theta = 0 \implies \cos(\theta) = 1$):**

$$\phi_{0,d}(\alpha) = \alpha \|b\| \quad (5.26)$$

The line exhibits maximum positive slope: the function increases at the highest possible rate (**steepest ascent**).

2. **d acute to b ($\theta \in (0, \pi/2) \implies \cos(\theta) > 0$):** Positive slope, strictly bounded above by $\|b\|$.

3. **d orthogonal to b ($\theta = \pi/2 \implies \cos(\theta) = 0$):**

$$\phi_{0,d}(\alpha) = 0, \quad \forall \alpha \in \mathbb{R} \quad (5.27)$$

The tomography is identically flat. Moving orthogonally to b corresponds to walking along the level set $L(f, 0)$, leaving the objective value unchanged.

4. d **obtuse** to b ($\theta \in (\pi/2, \pi) \implies \cos(\theta) < 0$): Negative slope: advancing along positive α decreases the function.
5. d **opposite** to b ($\theta = \pi \implies \cos(\theta) = -1$):

$$\phi_{0,d}(\alpha) = -\alpha\|b\| \quad (5.28)$$

The line achieves the steepest negative slope (**steepest descent**).




Figure 5.2: Tomographic inspection of a linear function: rotating search directions and identifying slope bounds.

5.5.3 Formal Derivation of the Steepest Descent Direction

Suppose an algorithm stands at the origin and takes a single step of fixed length $\|\Delta x\| = \varepsilon > 0$. The resulting change in objective value is:

$$\Delta f = f(\varepsilon d) - f(0) = \varepsilon \langle b, d \rangle = \varepsilon \|b\| \cos(\theta) \quad (5.29)$$

To achieve maximal reduction ($\Delta f < 0$ minimized), with $\varepsilon > 0$ and $\|d\| = 1$, one solves:

$$\min_{\|d\|=1} \langle b, d \rangle = \min_{\|d\|=1} \|b\| \cos(\theta) \quad (5.30)$$

Cosine reaches its absolute minimum -1 uniquely at $\theta = \pi$, producing the canonical result:

$$d^* = -\frac{b}{\|b\|} \quad (5.31)$$

Antigradient Principle

The direction yielding the greatest instantaneous rate of decrease for a linear objective is strictly the **antigradient** $-b$.

5.6 From Linear to Quadratic Functions: The Separable Case

5.6.1 Separable Quadratic Functions

The most direct extension of univariate quadratic functions to n dimensions sums n independent scalar parabolas:

$$f(x) = \sum_{i=1}^n (a_i x_i^2 + b_i x_i), \quad a, b \in \mathbb{R}^n \quad (5.32)$$

When $a_i = 1$ and $b_i = 0$ for all i , the function is the squared Euclidean norm:

$$f(x) = \sum_{i=1}^n x_i^2 = \|x\|^2 \quad (5.33)$$

Its level sets $L(f, v) = \{x \in \mathbb{R}^n : \|x\|^2 = v\}$ are concentric hyperspheres of radius $r = \sqrt{v}$.

5.6.2 Anisotropic Curvature and Eccentricity

In dimension $n = 2$, setting $b = 0$, $a_2 = 1$ and varying parameter $a_1 = a > 0$:

$$f(x_1, x_2) = ax_1^2 + x_2^2 \quad (5.34)$$

Level sets for a positive value $v > 0$ satisfy:

$$ax_1^2 + x_2^2 = v \iff \frac{x_1^2}{v/a} + \frac{x_2^2}{v} = 1 \quad (5.35)$$

The geometry depends directly on the curvature ratio:

- **If $a = 1$:** equal semi-axes ($\sqrt{v/a} = \sqrt{v}$): the level set is a circle.
- **If $a > 1$:** curvature is higher along x_1 . Level curves reach v closer to the origin along x_1 ($|x_1| = \sqrt{v/a} < \sqrt{v}$): the ellipse is **vertically elongated**.

- **If $0 < a < 1$:** curvature along x_1 is gentler: the ellipse is **horizontally elongated**.
- **As $a \rightarrow 0^+$:** horizontal curvature vanishes; semi-axis $\sqrt{v/a} \rightarrow +\infty$: the ellipse degenerates into two parallel lines $x_2 = \pm\sqrt{v}$.

5.7 General Homogeneous Quadratic Forms and Symmetrization

5.7.1 Definition and Cross-Terms

General homogeneous quadratic functions contain bilinear cross-terms $x_i x_j$ with $i \neq j$:

$$f(x) = \frac{1}{2} x^T Q x = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n Q_{ij} x_i x_j = \frac{1}{2} \sum_{i=1}^n Q_{ii} x_i^2 + \frac{1}{2} \sum_{i=1}^n \sum_{j \neq i}^n Q_{ij} x_i x_j \quad (5.36)$$

where $Q \in \mathbb{R}^{n \times n}$ is a given square matrix. The $1/2$ prefactor ensures the Hessian matrix equals Q .

Diagonal elements Q_{ii} dictate curvature along coordinate axes, while off-diagonal elements Q_{ij} couple coordinates and rotate principal axes.

Example 5.7.1 (Evaluation in 2D Space). Let Q be non-symmetric and let x be an arbitrary vector:

$$Q = \begin{bmatrix} 1 & 3 \\ 2 & 2 \end{bmatrix}, \quad x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad (5.37)$$

Step-by-step evaluation yields:

1. Intermediate product Qx :

$$Qx = \begin{bmatrix} 1 & 3 \\ 2 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} x_1 + 3x_2 \\ 2x_1 + 2x_2 \end{bmatrix} \quad (5.38)$$

2. Left multiplication by x^T :

$$\begin{aligned} x^T Q x &= \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} x_1 + 3x_2 \\ 2x_1 + 2x_2 \end{bmatrix} = x_1(x_1 + 3x_2) + x_2(2x_1 + 2x_2) \\ &= x_1^2 + 3x_1x_2 + 2x_1x_2 + 2x_2^2 = x_1^2 + 5x_1x_2 + 2x_2^2 \end{aligned} \quad (5.39)$$

Off-diagonal entries ($Q_{12} = 3$ and $Q_{21} = 2$) aggregate into a single term with coefficient $3 + 2 = 5$.

5.7.2 Symmetrization Theorem

Theorem 5.7.1 (Invariance Under Symmetrization). *Let $Q \in \mathbb{R}^{n \times n}$ be an arbitrary non-symmetric square matrix ($Q \neq Q^T$). The quadratic form $f(x) = \frac{1}{2}x^T Qx$ is identical across all of $\mathbb{R}^n$ to the form defined by its symmetric part:*

$$Q_{\text{sym}} = \frac{Q + Q^T}{2} \quad (5.40)$$

that is:

$$\frac{1}{2}x^T Qx = \frac{1}{2}x^T Q_{\text{sym}}x, \quad \forall x \in \mathbb{R}^n \quad (5.41)$$

Proof. Quantity $x^T Qx$ is a scalar (1×1 matrix). Since a scalar equals its own transpose:

$$x^T Qx = (x^T Qx)^T \quad (5.42)$$

Applying the transpose rule $(ABC)^T = C^T B^T A^T$:

$$(x^T Qx)^T = x^T Q^T (x^T)^T = x^T Q^T x \quad (5.43)$$

Expressing $x^T Qx$ as the average of itself and its transpose:

$$x^T Qx = \frac{x^T Qx + x^T Q^T x}{2} = x^T \left(\frac{Q + Q^T}{2} \right) x = x^T Q_{\text{sym}}x \quad (5.44)$$

Multiplying by 1/2 establishes functional identity. $\square$

Memory Efficiency and Course Convention

A symmetric $n \times n$ matrix stores only $n(n+1)/2$ entries rather than n^2 , halving memory requirements at large scale.

Standard Convention: Throughout continuous optimization, every quadratic matrix Q is assumed to be **symmetric**:

$$Q = Q^T \quad (5.45)$$

5.8 Spectral Decomposition and Matrix Taxonomy

5.8.1 Central Symmetry and Tomography

Because $(-x)^T Q(-x) = (-1)^2 x^T Qx = x^T Qx$, homogeneous quadratic forms are centrally symmetric about the origin:

$$f(-x) = f(x), \quad \forall x \in \mathbb{R}^n \quad (5.46)$$

The origin $x = 0$ is always a stationary point. For unit direction $\|d\| = 1$, the tomography from the origin is:

$$\phi_{0,d}(\alpha) = \frac{1}{2}(\alpha d)^T Q(\alpha d) = \frac{1}{2}\alpha^2(d^T Q d) \quad (5.47)$$

This is a **homogeneous scalar parabola** whose curvature is determined by:

$$a_d = d^T Q d \quad (5.48)$$

5.8.2 The Spectral Theorem for Real Symmetric Matrices

Directional curvature a_d is fully governed by the spectral decomposition of Q . Since $Q = Q^T \in \mathbb{R}^{n \times n}$:

1. **Real Spectrum:** All n eigenvalues $\lambda_1, \dots, \lambda_n$ are **strictly real**.
2. **Orthonormal Eigenvector Basis:** $\mathbb{R}^n$ admits an orthonormal basis of eigenvectors:

$$H_i^T H_j = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases} \quad (5.49)$$

3. **Spectral Factorization:** With orthogonal matrix $H = [H_1, \dots, H_n] \in \mathbb{R}^{n \times n}$ ($H^T H = H H^T = I$) and diagonal matrix $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$:

$$Q = H \Lambda H^T = \sum_{i=1}^n \lambda_i H_i H_i^T \quad (5.50)$$

Eigenvalue Ordering

Eigenvalues are arranged in descending order:

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n \quad (5.51)$$

where $\lambda_{\max} = \lambda_1$ and $\lambda_{\min} = \lambda_n$.

5.8.3 Variational Characterization: The Rayleigh Quotient

For any non-zero vector $d \in \mathbb{R}^n \setminus \{0\}$, the **Rayleigh quotient** is bounded by extreme eigenvalues:

$$\lambda_n \leq \frac{d^T Q d}{\|d\|^2} \leq \lambda_1 \quad (5.52)$$

For unit vectors $\|d\| = 1$:

$$\max_{\|d\|=1} d^T Q d = \lambda_1 = H_1^T Q H_1 \quad (5.53)$$

$$\min_{\|d\|=1} d^T Q d = \lambda_n = H_n^T Q H_n \quad (5.54)$$

5.8.4 Taxonomy of Matrix Definiteness

Eigenvalue signs systematically characterize quadratic behavior:

Class	Notation	Spectrum	Curvature ($\forall d \neq 0$)
Positive Definite	$Q \succ 0$	$\lambda_n > 0$	$d^T Q d > 0$ (strictly convex)
Positive Semidefinite	$Q \succeq 0$	$\lambda_n = 0 \wedge \lambda_1 > 0$	$d^T Q d \geq 0$ (degenerate convex)
Negative Definite	$Q \prec 0$	$\lambda_1 < 0$	$d^T Q d < 0$ (strictly concave)
Negative Semidefinite	$Q \preceq 0$	$\lambda_1 = 0 \wedge \lambda_n < 0$	$d^T Q d \leq 0$ (degenerate concave)
Indefinite	$Q \succ \prec 0$	$\lambda_1 > 0 \wedge \lambda_n < 0$	Assumes both positive and negative values

Table 5.1: Definiteness taxonomy for real symmetric matrices.

5.9 Spectral Tomography Across Four Regimes

Selecting direction d along an orthonormal eigenvector H_i yields the pure form:

$$\phi_{H_i}(\alpha) = \frac{1}{2}\alpha^2(H_i^T Q H_i) = \frac{1}{2}\alpha^2(H_i^T(\lambda_i H_i)) = \frac{1}{2}\lambda_i \alpha^2 \quad (5.55)$$

We examine the four spectral regimes using classroom benchmarks:

5.9.1 Case I: Positive Definite Matrix ($Q \succ 0$)

Consider the positive definite configuration:

$$Q = \begin{bmatrix} 6 & -2 \\ -2 & 6 \end{bmatrix}, \quad H = \frac{\sqrt{2}}{2} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}, \quad \lambda = \begin{bmatrix} 8 \\ 4 \end{bmatrix} \quad (5.56)$$

- Along $d = \pm H_1$: $\phi(\alpha) = 4\alpha^2$ (maximum curvature).
- Along $d = \pm H_2$: $\phi(\alpha) = 2\alpha^2$ (minimum curvature).
- For intermediate directions, curvature varies continuously within $[4, 8] > 0$.
- **Optimization:** Strictly convex. Global unconstrained minimizer is **unique**:

$$x^* = 0, \quad f^* = 0, \quad \lim_{\|x\| \rightarrow \infty} f(x) = +\infty \quad (5.57)$$

5.9.2 Case II: Positive Semidefinite Matrix ($Q \succeq 0$)

Consider the singular case with a zero eigenvalue:

$$Q = \begin{bmatrix} 4 & -4 \\ -4 & 4 \end{bmatrix}, \quad H = \frac{\sqrt{2}}{2} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}, \quad \lambda = \begin{bmatrix} 8 \\ 0 \end{bmatrix} \quad (5.58)$$

- Eigenvalue $\lambda_2 = 0$ defines the nullspace: $\ker(Q) = \text{span}(H_2)$.
- Along $d = \pm H_1$: strictly convex parabola $\phi(\alpha) = 4\alpha^2$.
- Along $d = \pm H_2$: since $\lambda_2 = 0$:

$$\phi_{H_2}(\alpha) = 0, \quad \forall \alpha \in \mathbb{R} \quad (5.59)$$

The tomography is a horizontal flat line.

- **Optimization:** Optimal value is $f^* = 0$, but optimal minimizers are **non-unique**:

$$X^* = \ker(Q) = \{x \in \mathbb{R}^2 : x = \gamma H_2, \gamma \in \mathbb{R}\} \quad (5.60)$$

The entire H_2 axis consists of equivalent global minimizers.

5.9.3 Case III: Indefinite Matrix ($Q \succ \prec 0$)

Consider a matrix with mixed-sign eigenvalues:

$$Q = \begin{bmatrix} 3 & -5 \\ -5 & 3 \end{bmatrix}, \quad H = \frac{\sqrt{2}}{2} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}, \quad \lambda = \begin{bmatrix} 8 \\ -2 \end{bmatrix} \quad (5.61)$$

- Along $d = \pm H_1$: convex upward parabola ($\phi(\alpha) = 4\alpha^2 \rightarrow +\infty$).
- Along $d = \pm H_2$: concave downward parabola ($\phi(\alpha) = -\alpha^2 \rightarrow -\infty$).
- Along asymptotic directions with $d^T Q d = 0$, the profile is flat.
- **Optimization:** Unbounded both above and below:

$$\inf_{x \in \mathbb{R}^n} f(x) = -\infty, \quad \sup_{x \in \mathbb{R}^n} f(x) = +\infty \quad (5.62)$$

The origin $x = 0$ is a **saddle point**.

5.9.4 Case IV: Negative Definite Matrix ($Q \prec 0$)

Consider strictly negative eigenvalues:

$$Q = \begin{bmatrix} -6 & -2 \\ -2 & -6 \end{bmatrix}, \quad H = \frac{\sqrt{2}}{2} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}, \quad \lambda = \begin{bmatrix} -4 \\ -8 \end{bmatrix} \quad (5.63)$$

The function is strictly concave:

$$x_{\max}^* = 0, \quad f_{\max}^* = 0, \quad \inf_{x \in \mathbb{R}^n} f(x) = -\infty \quad (5.64)$$

5.10 Geometry of Level Sets

5.10.1 Diagonalization and Orthonormal Coordinate Rotation

Consider the unit level set:

$$L(f, 1) = \left\{ x \in \mathbb{R}^n : \frac{1}{2} x^T Q x = 1 \right\} \quad (5.65)$$

Applying spectral decomposition $Q = H \Lambda H^T$ and rotating coordinates via:

$$y = H^T x \iff x = H y \quad (5.66)$$

the equation assumes canonical decoupled form:

$$\frac{1}{2}(Hy)^T Q(Hy) = 1 \iff \frac{1}{2}y^T (H^T Q H)y = 1 \iff \frac{1}{2}y^T \Lambda y = 1 \quad (5.67)$$

Explicitly in coordinates y :

$$\sum_{i=1}^n \frac{1}{2} \lambda_i y_i^2 = 1 \iff \sum_{i=1}^n \frac{y_i^2}{\frac{2}{\lambda_i}} = 1 \quad (5.68)$$

5.10.2 Semi-Axis Theorem for Ellipsoids

Theorem 5.10.1 (Orientation and Length of Semi-Axes). *Let $Q \succ 0$ be positive definite. Level set $L(f, 1)$ is a compact n -dimensional ellipsoid satisfying:*

1. Principal **symmetry axes** align with the orthonormal eigenvectors H_i .
2. The **semi-axis length** along direction H_i equals:

$$\text{semi-axis}_i = \sqrt{\frac{2}{\lambda_i}} \quad (5.69)$$

Proof. Intersecting the ellipsoid with the line spanned by H_i ($x = \alpha H_i$):

$$\phi_{H_i}(\alpha) = 1 \iff \frac{1}{2} \lambda_i \alpha^2 = 1 \iff \alpha^2 = \frac{2}{\lambda_i} \iff \alpha = \sqrt{\frac{2}{\lambda_i}} \quad (5.70)$$

identifying the Euclidean distance from origin to vertex along the eigenvector. $\square$

Inverse Relation Between Curvature and Axis Length

Semi-axis length is inversely proportional to the square root of the eigenvalue:

- Large $\lambda_i \implies$ high curvature / rapid growth $\implies$ **short semi-axis**.
- Small $\lambda_i \implies$ low curvature / gentle growth $\implies$ **long semi-axis**.




Figure 5.3: Geometry of level set $L(f,1)$ for a positive definite matrix: alignment along orthonormal eigenvectors H_1, H_2 and inverse semi-axis proportionality.

5.10.3 Degeneration into Open Elliptic Cylinders ($\lambda_i \rightarrow 0$)

As an eigenvalue tends to zero ($\lambda_i \rightarrow 0^+$):

$$\lim_{\lambda_i \rightarrow 0^+} \text{semi-axis}_i = \lim_{\lambda_i \rightarrow 0^+} \sqrt{\frac{2}{\lambda_i}} = +\infty \quad (5.71)$$

In 3D space with $\lambda_1 > 0$, $\lambda_2 > 0$, and $\lambda_3 \rightarrow 0^+$, the compact ellipsoid elongates along H_3 . At the exact limit $\lambda_3 = 0$, the surface opens to infinity, forming a hollow **elliptic cylinder**.

5.10.4 Summary of Level Geometries

Spectral Regime	Graph Geometry $\text{gr}(f)$	Level Sets $L(f, v)$
$Q \succ 0$	Convex paraboloid	Compact ellipsoids ($v > 0$)
$Q \succeq 0$ ($\exists \lambda_i = 0$)	Parabolic trough / gutter	Open elliptic cylinders (∞)
$Q \succ \prec 0$	Hyperbolic paraboloid (Saddle)	One- or two-sheeted hyperboloids
$Q \prec 0$	Inverted concave paraboloid	Compact ellipsoids ($v < 0$)

Table 5.2: Summary of geometric profiles for homogeneous quadratic forms.

5.11 Quadratic Optimization: Unconstrained vs Box-Constrained

5.11.1 Unconstrained Optimization: Immediate Analytical Solution

For the unconstrained problem:

$$\min_{x \in \mathbb{R}^n} \frac{1}{2} x^T Q x \quad (5.72)$$

spectral analysis provides a complete characterization:

- If $Q \succ 0$: unique minimizer $x^* = 0$, $f^* = 0$.
- If $Q \succeq 0$: optimal subspace $X^* = \ker(Q)$, $f^* = 0$.
- If any eigenvalue is strictly negative ($\lambda_n < 0$): $f^* = -\infty$.

5.11.2 Box-Constrained Optimization: Complexity Shift and $\mathcal{NP}$ -Hardness

Consider the identical objective restricted to a bounded hyper-rectangle:

$$(P_{\text{box}}) \quad \min \left\{ \frac{1}{2} x^T Q x : x \in X = [x_-, x_+] \subset \mathbb{R}^n \right\} \quad (5.73)$$

where $[x_-, x_+] = \{x \in \mathbb{R}^n : x_i^- \leq x_i \leq x_i^+, i = 1, \dots, n\}$.

While univariate box-constrained optimization ($n = 1$) admitted an immediate $\mathcal{O}(1)$ solution, in general dimension a dramatic complexity transition occurs:

1. **If $Q \succeq 0$ (Convex Problem):** The problem possesses a unique global minimum (or a compact convex set of equivalent minima). It can be solved in polynomial time via interior-point methods or projected gradient algorithms [4, 5]. However, **no closed-form formula exists:** solutions require iterative numerical algorithms.
2. **If Q is Indefinite ($Q \succ \prec 0$) or Negative Definite ($Q \prec 0$) (Non-Convex Problem):** Negative curvature pushes minima to the outer boundary of the feasible domain. Local and global minima reside strictly on the **vertices of the box**:

$$x^* \in \text{vertices}(X) = \prod_{i=1}^n \{x_i^-, x_i^+\} \quad (5.74)$$

Since the hypercube possesses 2^n vertices, brute-force search is computationally intractable.

Theorem 5.11.1 (Complexity of Box-Constrained QP [8]). *Box-constrained quadratic programming (P_{box}) with a general non-positive semidefinite matrix Q is $\mathcal{NP}$ -hard.*

The Didactic Paradox of Optimization

Imposing seemingly trivial coordinate-wise bounds $x_i \in [x_i^-, x_i^+]$ on an elementary quadratic function elevates an instantly solvable problem into the class of $\mathcal{NP}$ -hard computational problems.

Bibliography

- [1] Gilbert Strang. *Linear Algebra and Learning from Data*. Wellesley-Cambridge Press, Wellesley, MA, 2019.
- [2] Lloyd N. Trefethen and David Bau III. *Numerical Linear Algebra*. SIAM, Philadelphia, 1997.
- [3] J. W. Demmel. *Applied Numerical Linear Algebra*. SIAM, Philadelphia, PA, 1997.
- [4] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, Cambridge, 2004.
- [5] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. Springer, New York, 2nd edition, 2006.
- [6] Mokhtar S. Bazaraa, Hanif D. Sherali, and Chitharanjan M. Shetty. *Nonlinear Programming: Theory and Algorithms*. John Wiley & Sons, Hoboken, NJ, 3rd edition, 2006.
- [7] Gene H. Golub and Charles F. Van Loan. *Matrix Computations*. Johns Hopkins University Press, Baltimore, 4th edition, 2013.
- [8] Etienne de Klerk. The complexity of a simple quadratic program on a box. *Operations Research Letters*, 36(4):468–472, 2008.